

# Mandated Transparency and the Production of Research: Evidence from a Data- and Code-Sharing Policy\*

John M. Barrios<sup>†</sup>      Roman Chychyla<sup>‡</sup>

*Preliminary draft — September 7, 2026*

## Abstract

We examine how research changes when a journal requires authors to document their data and provide the code used to build their samples. The *Journal of Accounting Research* (*JAR*) introduced this requirement for all empirical submissions in January 2015, while its closest rivals did not. We study roughly 4,600 articles from six journals around the policy change. Consistent with disclosure limiting authors' discretion, *JAR* papers published after the mandate are easier to read, rely less on hedging language, and report results more comprehensively. The reduction in hedging is largest for studies with the weakest headline results. Consistent with the policy changing who submits, *JAR* moves toward more novel work from scholars at top research universities; its articles receive fewer citations overall, but from more prominent outlets. The writing effects are absent from exempt theory articles and reappear when another journal adopts a comparable requirement. In sum, increasing the openness of the research process influences not only what can be checked and reproduced, but also what gets written and who ultimately publishes in which outlets.

**Keywords:** research transparency, open science, data-sharing policy, difference-in-differences, synthetic control, textual analysis, economics of science.

**JEL:** A14, C18, C52, L82, O33.

---

\*We thank *Thomas Bourveau, Menaka Hampole, Christian Leuz, Miguel Minutti-Meza, Jake Thomas, Edward Watts, Thomas Wollmann*, and conference participants at Yale SOM for helpful comments and discussions. Replication code and data are described in the Appendix and will be posted publicly upon publication. While we have reviewed for each of the six journals, and each of us has had work rejected at every one of them, neither of us has held an editorial or advisory role over the period studied. All measures are built from publicly available article text, metadata, and citation records; no journal provided data or cooperation. We study *JAR* because its January 2015 policy affords a clean start date, not to evaluate the journal, and we report costs and benefits together. John Barrios was on the faculty of the University of Chicago Booth School of Business, which publishes *JAR*, from 2015 to 2020. Comments welcome.

<sup>†</sup>Yale University and NBER. Email: [john.m.barrios@yale.edu](mailto:john.m.barrios@yale.edu).

<sup>‡</sup>University of Miami Business School. Email: [rchychyla@bus.miami.edu](mailto:rchychyla@bus.miami.edu)

# 1 Introduction

Academic research trades in a marketplace of ideas, and like any asset it carries a price. The results that the readers consider credible are adopted and extended, while results that are questioned are discounted (Barrios et al., 2025). Over the last ten years, readers have found increasing grounds for skepticism. Replication efforts in psychology, economics, and finance have shown high failure rates (Open Science Collaboration, 2015; Camerer et al., 2016, 2018; Chang and Li, 2022; Pérignon et al., 2024), and related work has highlighted how flexibility in research design decisions can produce statistically significant findings (Brodeur et al., 2016, 2020; Simonsohn et al., 2014; Andrews and Kasy, 2019). When a researcher can choose from many defensible sample cuts, controls, and specifications and only report those that meet the significance threshold, the published record moves away from the truth. This drift does not have to stem from academic fraud; it can stem from the research design discretion inherent in empirical work, exercised by researchers who face strong incentives to publish and who write for journals that prefer significant to null results.

The response to these issues has largely converged around a single idea: greater transparency. Funders and journals have demanded pre-analysis plans, registered reports, and, more specifically, the disclosure of the data and code that generate the findings of a paper (Christensen and Miguel, 2018; Miguel et al., 2014; Nosek et al., 2015). The rationale for these policies mirrors that of any mandated disclosure system. When others can review the materials that generate the results, the authors expect scrutiny and adjust what they study and how they write. Economists since Dasgupta and David (1994) and Stephan (1996) have examined how the distribution of scientific credit and the incentive system in open science influence the way research is carried out. Subsequent studies find that the composition of research teams, the influence of prominent scholars, and the access to data specifically affect the questions studied (Wuchty et al., 2007; Azoulay et al., 2010; Iaria et al., 2018; Hill and Stein, 2025).

Most of the research on these transparency mandates focuses on the replicability of published papers, while less is known about how authors' anticipation of disclosure affects what is written and who writes it. Just as the accounting literature has shown, disclosure is not just ex-post verification after publication. It also has an *ex-ante* cost that the author has to bear while writing the paper. This cost is also most likely borne by projects that are difficult to document or whose results depend on a few discretionary decisions. When an author chooses what to write, how to present it, and where to submit it, they should take this disclosure cost into account before any review. Thus, whether the mandate alters how authors write, who decides to submit, or neither is a supply-side issue rather than a

replication issue, and it is the central research question we study in this paper.

To study this question, we examine the *Journal of Accounting Research*'s (*JAR*) 2015 Data Policy. Beginning on January 1, 2015, authors of empirical papers had to submit a datasheet on their data sources, download dates, and who handled the data, and before final acceptance the authors had to provide the code that constructs the analysis sample from the raw files. These materials were then posted on the journal's website after publication. *JAR* did not require the actual final dataset nor do in-house replications. That is, the compliance was self-certified. The other top journals in the field – *The Accounting Review*, the *Journal of Accounting and Economics*, *Contemporary Accounting Research*, and the *Review of Accounting Studies* – imposed no similar change through the sample period of 2012–2017. We use this change as the source of identifying variation.

The accounting literature provides a natural setting to examine these issues. Empirical research in accounting is dominated by archival studies that use large datasets, frequently compiled manually from proprietary or licensed sources, where the process from raw data to a regression table is long and seldom transparent to readers. *JAR* is one of the three leading journals in accounting, so the mandate applies to work at the top of the academic research market, where the incentive to publish is strongest and the cost of submitting opaque or fragile research is largest. The policy was adopted unilaterally, on an unambiguous date, by a single journal without having similar policies adopted by its closest competitors. That gives us a sharply defined treatment and a set of closely comparable control journals that draw on the same authors, topics, and methods that we can use for comparison. A second journal, *Management Science*, adopted a similar transparency rule four years later, which gives us a second treated journal and variation in policy stringency that we use to check our inferences.<sup>1</sup>

The transparency mandate imposes a compliance cost that unevenly affects potential journal submissions. Authors can respond along two margins: they may change how they write and report, they may redirect their work to other outlets, or they can do both. We refer to the first as the *behavioral* margin and to the second as the *market* margin. Ex ante, the policy could lead to three outcomes: it might alter how authors who continue submitting to *JAR* write and conduct their work, it could also affect which papers and authors choose to submit to the journal in the first place, or it could have no effect. Under the *discipline* view, authors who keep submitting to *JAR* anticipate that the data trail and the code will be public, so they write plainly (i.e., in a more readable fashion), hedge their claims less, and report their results more completely, leaving insignificant results in the manuscript rather than omitting if they don't pass the significance threshold. Under *selection*, research projects

---

<sup>1</sup>While *Management Science* is not strictly an “accounting” area journal and publishes articles in other fields (such as management, marketing, and finance), it is considered a top outlet for accounting publications.

that are opaque or hard to document, where documenting the data trail exposes discretionary choices in sample construction or specification, are redirected to other outlets, and more transparent research stays. *JAR*’s author pool and portfolio then shifts toward more original contributions. Finally, under a *skepticism* view, self-certified disclosures have little effect, and neither margin moves. Importantly, both discipline and selection responses can coexist, and a journal-level comparison cannot tell them apart, because either one alone would change what *JAR* publishes in a similar way.

We use multiple empirical strategies to disentangle the possible responses. Our primary approach is a difference-in-differences design that compares papers submitted to *JAR* with papers submitted to five other accounting journals during 2012–2017, incorporating submission-year and topic fixed effects. Our coefficient of interest is  $JAR \times Post$ , which captures how the outcomes of *JAR* papers changed after the disclosure mandate changed relative to those of the papers in the set of control journals. Because this design alone cannot disentangle changes in author behavior from shifts in who submits to *JAR*, we complement it with additional designs intended to more cleanly isolate each margin. For the behavioral margin, we add author fixed effects to capture within-author shifts in writing, leverage a second treated journal (the 2019 policy at *Management Science*), and contrast empirical papers with exempt theory papers as another comparison group. For the market margin, we keep the baseline specification, but compare the results to finance journals as a never-treated control group and to a synthetic-control method built for a single treated unit.

Studying how research changes means examining what the papers actually communicate, rather than simply counting their citations. To this end, we use several text-based methods, including fine-tuned large language-model classifiers and text embeddings, on the full text of roughly 4,600 published papers in our sample around the policy change. We then combine these measures with bibliographic metadata and the citation history of each article (Gentzkow et al., 2019). With these inputs, we create four groups of outcome measures. The first captures the writing style, which we assess via section-level readability using the Gunning-Fog index and the fraction of sentences that hedge claims, following the disclosure literature in accounting and finance (Li, 2008; Loughran and McDonald, 2014; Bonsall et al., 2017; Bushee et al., 2018). The second measure reflects novelty, specifically, how far the abstract and introduction of a paper diverge from prior research. We operationalize this distance by, first, comparing the text in its own language using semantic text embeddings and, second, characterizing novelty through the portfolio of references cited (Kelly et al., 2021; Uzzi et al., 2013). The third measure focuses on the completeness of the empirical reporting. Specifically, we parse each paper to collect all the coefficients reported in the tables of the paper, counting the number of estimates that are used to support each main finding, and evaluate how the

test statistics reported cluster around conventional significance cutoff points, since a grouping of estimates just above these thresholds is consistent with p-hacking or specification search (Brodeur et al., 2016, 2020). Our fourth measure focuses on the paper’s reception and impact, which we measure as the overall citation counts and the prestige rankings of the journals that later cite the paper. Together, these measures allow us to evaluate whether the mandate altered research through the channels the policy was intended to affect.

We begin by analyzing the behavioral margin and document that authors publish more readable articles in *JAR*. Following the mandate, abstract readability improves by approximately one Gunning–Fog grade, corresponding to a decline of roughly one-third of a standard deviation. The improvement in readability is concentrated in sections of the paper that offer context and interpretation (the abstract and introduction) and does not appear in the more technical methods and results sections, a pattern consistent with a shift in how authors present and frame their contributions. However, *JAR* abstracts show a pre-trend toward increased readability that starts around two years before the rule is implemented. In addition, editor turnover or a stylistic shift in the journal in 2015 could also produce a similar change in writing. To rule this out, we conduct a series of tests designed to distinguish behavioral responses from other explanations, such as shifts in the composition of authors or concurrent changes in editorial policy. The estimated readability effect remains robust across all three tests, and none is easily attributable to a purely compositional mechanism or to an editorial change limited to a single journal. First, exploiting *Management Science*’s 2019 data and code policy, we find a within-author increase in readability: the same authors write more readable papers after the policy than before. Second, the magnitude of the effect increases with the stringency of the policy in the two journals, ordered by the degree of required compliance. Third, the effect is not observed for theory papers in *JAR*, which do not rely on data or code and are therefore exempt from the mandate. Collectively, these results are more consistent with a behavioral response to the mandate than simply changes in editorial direction or in the pool of authors who submit to *JAR*.

Next, we turn to how authors present evidence, documenting results consistent with a behavioral response. Post-mandate *JAR* papers make fewer headline claims and surround each result with more supporting coefficients, consistent with anticipating outside scrutiny of the analysis. Rather than selectively emphasizing the strongest outcomes, authors report results more completely: the number of main results per paper falls from 2.6 to 2.0, while the total number of reported coefficients rises from about 49 to 65. The coefficients authors add are disproportionately weak; the share of conventionally significant results decreases, but the average significance of the test statistics is essentially unchanged, and their distribution does not clearly shift away from the just-significant region. In other words, the mandate changes

how thoroughly authors report their results, but not the statistical significance of what they report.

Hedging also decreases, but the decline is not uniform: it is concentrated in the contribution claims of papers with weaker main results. That is, papers with weaker results stop “overselling” after the mandate. The overall drop in hedging also appears in the exempt theory papers, so part of it is journal style. Its concentration in papers with weaker empirical results, however, cannot be explained by style alone, since a change in style is unlikely to target papers by the statistical strength of their main results.

We then turn to the market margin. We find that following the mandate, *JAR* publishes papers that are more novel across all novelty metrics we construct. The estimate holds against a synthetic control and when compared to never-treated finance articles. Yet, we don’t observe a change in novelty when we control for authors’ fixed effects, suggesting that this result is driven by selection, rather than a behavior change. The mandate does not incentivize an author to do more original research, rather it sorts more original work into the journal. One interpretation is that authors with original, well-documented projects face a lower disclosure mandate cost and are less worried about scrutiny. We also observe a change in the author pool at *JAR* post-policy. Relative to controls, post-mandate *JAR* authors are about 0.3 standard deviations more likely to be at a top research department and carry a higher *h*-index; they are directionally more junior and write somewhat shorter papers, and they draw on a different mix of proprietary and public data.

Our citation analysis also supports a market-side (selection) channel. Post-mandate, *JAR* papers are cited less in raw counts, but the citations that remain come from higher ranked journals. These more prominent citations come both from the field’s leading accounting outlets and from outside the field: premier business journals account for a larger share of the smaller citation pool.

Taken together, our results show that both channels are at work: authors write differently, and the mix of authors and projects the journal attracts also changes. The novelty result holds against a synthetic control. The writing and citation effects are large and point in the same way throughout, though they are estimated less precisely with only one treated journal. Several caveats apply throughout, and we discuss them in more detail below. With only one journal treated, our tests have limited power, so some of our null results may simply reflect that, rather than showing that nothing changed. Our writing, novelty, and reported-inference measures come from text, built using classifiers and embeddings, so we treat them as informative proxies. Finally, *JAR* policy is applied based on the submission, not publication, date of the article. For some treated articles, the submission date information is not available, so we conservatively impute the treatment timing from publication dates; our

main results hold when we keep only the papers with an observed submission date.

Our study contributes to three strands of literature. First, it connects to the economics of science and to research transparency (Christensen and Miguel, 2018; Miguel et al., 2014; Nosek et al., 2015; Camerer et al., 2018). The work in the economics of science examines how institutions, incentives, and access to ideas shape knowledge production (Dasgupta and David, 1994; Stephan, 1996; Wuchty et al., 2007; Azoulay et al., 2010; Iaria et al., 2018; Hill and Stein, 2025), and a related line of research finds that greater openness can change research topics scientists decide to study (Murray et al., 2016). Accounting has faced similar concerns, Dyckman and Zeff (2014) examined one year of empirical articles in *The Accounting Review* and *JAR* and document unjustified sample-window choices and a reliance on statistical significance tests instead of economic significance. Salterio (2014) and Basu (2012) argue that field incentives discourage replication and a survey of attendees at the 2019 *JAR* conference found that most of the respondents saw irreproducible findings as common, attributing them to publication incentives and selective reporting, and reported encountering them in the work of others without following up (Hail et al., 2020). Journals have responded through a registered-reports pilot at the 2017 *JAR* conference (Bloomfield et al., 2018) and by implementing data and code requirements: *JAR* in 2015 and the American Accounting Association will now require a data description sheet and complete code, from data construction through analysis, for conditionally accepted papers to *The Accounting Review* on or after July 1, 2026, with its other journals beginning in July 2027.<sup>2</sup> This literature asks whether published results replicate and whether researchers believe they do. Instead, we explore a more upstream question: whether the possibility of disclosure of one’s code alters the supply of research, and find that the effect operates at least as much through the submissions market as it does through any individual author.

Second, we contribute to the economics-of-disclosure and the textual-analysis literature on disclosure (Leuz and Wysocki, 2016; Healy and Palepu, 2001; Beyer et al., 2010; Li, 2008; Loughran and McDonald, 2011; Dyer et al., 2017; Bushee et al., 2018; Gentzkow et al., 2019). This work has traditionally focused on whether mandatory reporting leads to behavioral changes or instead primarily reshapes the scope, content, and prominence of the information disclosed; how disclosure costs are allocated when information is proprietary or expensive to produce; and whether externalities fall on stakeholders not directly related to the disclosure requirement (Leuz and Wysocki, 2016; Verrecchia, 2001). Journal-level requirements to share data and code create a similar set of issues for authors and papers: these policies can

---

<sup>2</sup>American Accounting Association, “Data Integrity and Research Record,” *The Accounting Review* Guide for Authors, updated March 2026. The materials are kept in a private AAA repository for the editorial team with public posting only encouraged but not required.

influence how papers are designed and which findings are presented, reallocate which studies are publishable across journals, and alter the overall distribution of citations and readers' attention. We treat the research article as a tool for disclosure and measure these shifts using text-based proxies for writing style and the reporting of statistical inference, together with novelty measures and indicators capturing submission mix and citation outcomes. Bringing insights from accounting to the economics of open science.

Finally, we provide evidence on an issue that emerges whenever a single entity shifts its policies while its close substitutes do not. In our case, focusing only on published articles mixes changes in the authors' responses to the policy with changes in the pool of papers submitted to the journal. For authors who continue to publish in the outlet, the mandate can alter the way evidence is written and communicated. At the same time, it may cause papers and authors to reallocate across journals. To separate these two margins, we pair within-author comparisons with evidence from exempt theory papers, another mandate at another journal, and additional comparison groups, including never-treated finance journals. Where appropriate, we rely on design-based inference (Conley and Taber, 2011; Cameron et al., 2008; MacKinnon and Webb, 2017; Abadie, 2021; Arkhangelsky et al., 2021; Roth et al., 2023). This also has implications for market-level policy changes: as soon as close alternatives adopt comparable mandates, authors will find it harder to avoid disclosure by moving their submissions to other outlets. As a result, the sorting effect would weaken and that in turn could also affect the influence the mandate has on how authors frame and report their evidence.

The remainder of the paper proceeds as follows. Section 2 describes the policy and our conceptual framework. Section 3 describes the corpus, the text-based measures, and the sample. Section 4 outlines the empirical strategy, our inference approach, and a guide to which design supports each claim. Section 5 reports the behavioral and market results, along with parallel-trends, treatment-timing, and robustness checks of editorial-change on which identification is based. Section 6 examines the mechanisms, limitations, and general-equilibrium responses. Section 7 concludes.

## 2 Institutional Background and Conceptual Framework

### 2.1 The *JAR* data- and code-sharing policy

*JAR* is one of the three premier journals in accounting, along with *The Accounting Review* (*TAR*) and the *Journal of Accounting and Economics* (*JAE*). Its Data Policy took effect for all submissions received on or after January 1, 2015. The policy required, at the paper's



initial submission, a data-description sheet documenting how the raw data were obtained or generated, their sources, the dates on which the data were downloaded, the instruments used in surveys and experiments, and which authors handled the data and ran the analyses. Authors using proprietary data also had to provide the editors, privately, with a contact at the data provider and the terms of the data-sharing agreement. Before final acceptance, authors had to submit the computer code that converted the raw data into the dataset used in the analysis, or, where parts of that process were proprietary, a step-by-step description precise enough for another researcher to reconstruct the sample. The requirement covered archival, experimental, field, survey, and simulation work. Purely analytical papers were not subject to the requirement.

Two aspects of the 2015 policy warrant highlighting. First, it created a disclosure requirement rather than an obligation to reproduce results. The journal did not require submission of the final dataset or the code used to run the statistical analysis. Sharing was encouraged, but authors were only required to keep their data and programs for six years. *JAR* similarly did not have a data editor and did not conduct internal replications during our sample period. Adherence was self-attested; the datasheet and sample-construction code were made available on the journal’s website after publication. *JAR* did not expand the requirement to include complete analysis and table-production code, log files, or mandatory identifiers until an April 2024 revision, well beyond our sample window. A November 2016 revision renamed the policy the “Data and Code Sharing Policy” and added structural papers, but it did not alter the core requirements. Appendix A records the policy’s full revision history and the obligations under each version. *Management Science*, in contrast, implemented an active data editor in 2020; thus, the *JAR* rule functioned as a disclosure policy rather than a verification regime. Second, the other leading accounting journals had no comparable requirement during the 2012–2017 window. This contrast implies that *JAR* serves as the treated journal and the other accounting journals as the control group.

## 2.2 Descriptive framework

This disclosure requirement increases the cost of publishing empirical work in *JAR* that is opaque or not very robust. When authors must reveal data sources, download dates, and the code used to construct the sample, any finding that holds only for a specific sample restriction, a nonstandard winsorization choice, or an unreported data-processing step becomes much riskier to submit. In effect, it enables readers to inspect the code for such problems. The mandate operates by making work more visible rather than by having the journal replicate the analysis, and it bears most heavily on the very behaviors the transparency policy is

meant to curb (namely, specification searching and selective reporting) (Simonsohn et al., 2014; Brodeur et al., 2016; Andrews and Kasy, 2019).

We use a simple framework to organize these ideas before we turn to the data. An author with a completed project decides whether and where to submit. Each journal offers a prestige payoff, but imposes a cost of compliance with its standards. *JAR*'s mandate raises that compliance cost, and raises it most for the most *opaque* projects, those whose headline result leans on choices the author would rather not see disclosed. Authors compare the prestige-net-of-cost of *JAR* to that of the alternatives, and submit where the net payoff is highest. Two responses follow.

First, we consider how an author who keeps submitting to *JAR* writes and reports, knowing that the datasheet and code behind the paper will be public. We label this the *behavioral* margin, that is, the *discipline* response. The framework produces three predictions, the first two about how the paper is written and the third about how completely the results are reported. The first concerns how plainly a paper is written. A large literature in accounting and finance treats textual complexity as an economic choice rather than a matter of style. These studies document that firms write more opaquely when they have reason to obscure weak performance, that plain-language measures of readability predict how information is impounded, and that clearer disclosure carries consequences for valuation and attention (Li, 2008; Loughran and McDonald, 2014; Bonsall et al., 2017; Bushee et al., 2018; Dyer et al., 2017; Hwang and Kim, 2017). Authors tailor their language deliberately to the audience they expect to read it (Cao et al., 2023). The same logic applies to research writing. Once the pipeline behind a paper is public and any reader can retrace it, opaque exposition is less effective in masking a weak empirical result, and readable documentation complements the disclosed materials. Post-mandate *JAR* papers should therefore be easier to read in the sections that describe and interpret the analysis (prediction P1). A closely related prediction concerns the language of certainty. Because the reader can now test the author's interpretation against the documented analysis, disclosure rewards the precise claim and penalizes the hedged one. Therefore, we expect a lower linguistic hedging rate, particularly in cases where the author states the contribution (prediction P2).

The third prediction concerns how completely authors report their results. Once the full analysis can potentially be replicated, reporting only the specifications that clear the significance threshold becomes riskier, so authors should report more completely, leaving insignificant results in the manuscript rather than omitting them, and backing each headline claim with more supporting estimates rather than a few strong ones (Brodeur et al., 2016, 2020; McCloskey and Ziliak, 1996). The same logic implies that the distribution of reported test statistics should shift away from the just-significant thresholds toward the null (Simonsohn

et al., 2014; Andrews and Kasy, 2019; Elliott et al., 2022) (prediction P3).

The remaining responses are changes in *which* authors and projects *JAR* attracts and in how the resulting papers are received, not in how a given author works. This is the *market* margin, that is, the *selection* response, and it is the one a difference-in-differences alone cannot separate from behavior. We consider originality first. If marginal, threshold-driven results become harder to publish once the analysis is open to inspection, competition for acceptance at *JAR* tilts toward conceptual novelty. A growing literature measures scientific novelty through atypical recombination of prior ideas (Uzzi et al., 2013), through the disruptiveness of a work relative to its antecedents (Funk and Owen-Smith, 2017; Park et al., 2023), and through the distance of a work’s text from the existing literature (Kelly et al., 2021). This literature documents that novel work is both higher-impact and systematically disfavored by the acceptance process (Wang et al., 2017). A parallel literature shows that incentives and openness shape whether scientists pursue novel directions at all (Azoulay et al., 2011; Murray et al., 2016). We therefore expect measured novelty to rise (prediction P4). Whether novelty rises because a given author chooses more original projects or because the journal attracts more original work is the behavior-versus-selection question, and the within-author design answers it.

This composition response is not limited to novelty. Because submission to *JAR* is a choice, authors whose projects are opaque, or whose data are costly or impossible to share, are less likely to submit there, while authors with clean, well-documented projects are relatively more willing to pay the disclosure cost. *JAR*’s author pool and portfolio therefore shift at the cutoff. Part of any measured change in its output then reflects who now publishes there, not how a given set of authors behaves. This shift is the standard concern that an estimated effect reflects selection on unobservable characteristics (Oster, 2019), and the concern is especially salient in a market where authors sort across a few high-stakes outlets and compete for priority and placement (Card and DellaVigna, 2013; Heckman and Moktan, 2020; Hill and Stein, 2025) (prediction G1).

Once published, the paper is evaluated by the market through readership and citations. Work published under the disclosure regime may be read, trusted, and cited differently, with fewer of the citations coming from higher ranked journals, so that raw citation counts could fall even as the prominence of the citing journals increases (prediction P5). The control journals are part of the same market. Projects diverted from *JAR* could be submitted to the other accounting journals, so the control journals are themselves partially treated by the displaced submissions. Indeed, prior research documents that field-wide forces (the expansion of open-science policies, secular trends in how research is written and cited) can move treated and control journals together (Christensen and Miguel, 2018; Nosek et al., 2015). Either

channel violates the stable-unit-treatment-value assumption, which requires that a unit’s treatment leave the outcomes of other units unaffected, and biases the within-accounting comparison (Bertrand et al., 2004; Conley and Taber, 2011) (prediction G2).

A third possibility is that the mandate changes nothing; we call this the *skepticism* response. Compliance is self-certified, and the journal verifies nothing, so the rule may not be binding. In that case neither margin moves, and none of the predictions above should hold. Taking the predictions to the data is therefore also a test of whether a disclosure rule without verification has any effect.

## 3 Data and Measurement

### 3.1 Corpus

Our corpus consists of research articles published between 2008 and 2025 in six leading accounting journals: *JAR*, *TAR*, *JAE*, *Contemporary Accounting Research (CAR)*, the *Review of Accounting Studies (RAST)*, and accounting articles published in *Management Science (MS)*. For the general equilibrium analysis we also assemble the accounting articles in the three leading finance journals (the *Journal of Finance*, the *Journal of Financial Economics*, and the *Review of Financial Studies*), which serve as the never-treated benchmark.<sup>3</sup> For each article, we obtain the full text, parsed into sections (abstract, introduction, theory/hypotheses, methodology, results, conclusions), together with bibliographic metadata, submission and publication dates where available, author identifiers, and the forward citation record.

### 3.2 Measures

We construct four families of outcome measures: writing (readability and hedging), novelty, reported inference, and citations.

**Readability.** For each section of a paper we compute the Gunning–Fog index, a standard readability measure that increases in average sentence length and the share of multisyllabic words; higher values denote more complex, less readable text. The measure is widely adopted in the accounting and finance disclosure literature, where textual complexity is treated as a choice that conceals or reveals information (Li, 2008; Loughran and McDonald, 2014; Bonsall et al., 2017; Dyer et al., 2017).

---

<sup>3</sup>We classify a *Management Science* or finance-journal article as accounting if it either studies an accounting topic according to the OpenAlex database or at least one of its authors has previously published in an accounting journal, per the BYU accounting-research rankings list.

**Hedging.** We construct the hedging measure in two stages. In the first stage we classify *sentences* as hedged or not. We start from the corpus of Ningrum et al. (2025), 975 sentences drawn from academic journals and manually classified as hedged or not. Next, we use these data to fine-tune a transformer classifier, DistilRoBERTa, and label every sentence in every paper in our sample.<sup>4</sup> The fine-tuned classifier recognizes epistemic uncertainty and qualification (constructions such as “may suggest,” “appears to be consistent with,” “we cannot rule out,” “broadly in line with”), unlike a fixed keyword list, which misses the many ways authors qualify a claim and misclassifies in cases of negation and quotation. A sentence is classified as hedged when the model’s predicted probability exceeds 0.5. The paper- or section-level *hedging rate* is then the share of hedged sentences,  $h = (\text{hedged sentences})/(\text{total sentences})$ . A higher rate means more qualified, less committal text.

In the second stage we locate *where* in the paper the hedging occurs, because the framework predicts movement in the parts that frame and make a claim, not in the mechanical reporting of results. We parse each introduction into paragraphs and classify each paragraph by its rhetorical function, adapting the “create-a-research-space” (CARS) model of research-article introductions (Swales, 1990, 2004). CARS characterizes an introduction as a sequence of rhetorical moves in which the author establishes a territory, carves out a niche, and then occupies it. We map these moves onto three categories: *motivation* (Move 1), *gap* (Move 2), and *present research* (Move 3, Step 1), and we add *results* for the preview of principal findings (Move 3, Step 2). Because accounting introductions almost always devote at least one paragraph to stating what the paper adds to the literature, a rhetorical act that CARS subsumes within Move 3, we separate out *contribution* as a fifth category, following Hyland (2004) and Lim (2012). Paragraphs serving none of these functions, such as roadmaps, acknowledgments, and scope caveats, fall into a residual category that we exclude from the analysis. For each paragraph, we also check whether it claims novelty or identification. We classify the paragraphs with Qwen 3.6 27B, an open-weight large language model developed by Alibaba Group.<sup>5</sup> With sentences scored for hedging and paragraphs scored for their function, we compute a hedging rate *within* each rhetorical area.

---

<sup>4</sup>We fine-tune DistilRoBERTa on the 975 annotated sentences using stratified 10-fold cross-validation, the same evaluation protocol as Ningrum et al. (2025). Mean out-of-fold accuracy is 0.794 (standard deviation 0.030), comparable to the highest-accuracy fine-tuned models reported in Ningrum et al. (2025).

<sup>5</sup>We allow the model to pick up to two labels per paragraph, so that a paragraph can be both a gap and a contribution, for example. We validate the classifier against an independently produced benchmark: a separate large language model (Claude Opus) labeled 885 introduction paragraphs from 75 articles under the same taxonomy. The two models’ primary-or-secondary label sets overlap for 89.2% of paragraphs, suggesting high classification agreement.

**Novelty.** We measure novelty in three ways. Our main measure is textual: we embed each paper’s abstract and introduction using a pretrained sentence-embedding model (Qwen3-Embedding-4B, 2560 dimensions) and take the mean cosine distance between the paper and its twenty nearest neighbors among articles submitted in the prior three years ( $k$ -nearest-neighbor distance,  $k = 20$ ), pooling the prior-literature reference set across all journals so the measure is comparable across treated and control samples. A larger distance means the text is further from the recent literature. Because one embedding choice could drive our measure, we rebuild it across  $k \in \{1, 5, 10, 50\}$ , a five-year reference window, an abstract-only variant, a distance-to-field-centroid variant, and an “originality” variant, one minus the maximum similarity to any prior paper, in the spirit of Kelly et al. (2021).<sup>6</sup>

Our second measure is combinatorial, built from each paper’s reference list rather than its text. Following the recombination logic of Uzzi et al. (2013), we score the atypicality of the journal pairs a paper cites against a degree-preserving null, and we summarize how diverse the fields it draws on are (the entropy and the non-accounting share of its cited literature) and the average age of its references. These citation-based measures use no text at all, so their agreement with the embedding measure is an independent check rather than a mechanical one. Third, we use the introduction classifier to flag whether a paper makes an explicit rhetorical claim of novelty. We keep this separate from the distance measures throughout, because a claim of novelty is a statement authors make, not a property of the text, and the two are weakly correlated.

**Reported inference.** We use the Qwen 3.6 27B large language model to extract the test statistics reported in each paper’s tables and recover, for each reported coefficient, an approximate two-sided  $p$ -value from the reported  $t$ - or  $z$ -statistic or from the coefficient and its standard error. Where a table reports only significance stars, the  $p$ -value can be located only within a coarse band, so we exclude star-imputed results from the test-statistic outcomes. Because authors showcase some results and relegate others, we also classify each extracted result by salience: the full universe of extracted coefficients, the primary hypothesis tests, the results discussed in the introduction, and the paper’s single most-important (headline) result. We then summarize, at the paper level and within each salience tier, the share of results falling in conventional  $p$ -value bands. Following Brodeur et al. (2016), we also report the mass of  $z$ -statistics in a caliper just above versus just below the 1%, 5%, and 10% thresholds, an indicator of specification search.

---

<sup>6</sup>The effect is stable across all of them (Appendix H).

**Citations.** We build our citation outcomes from the OpenAlex citation graph. We match each paper to OpenAlex by DOI (97.5% of the sample) and retrieve every indexed work that cites it, together with the citing work’s journal and publication year. From this record we construct three types of outcome. The first is quantity: the log of one plus the paper’s total accumulated citations (a count restricted to the four years after publication behaves similarly). The second is field-adjusted impact: the log of one plus the paper’s field-weighted citation index, which scales the citations a paper receives by the count expected for works of the same field, type, and publication year. The third is the status composition of the citing journals, which asks where the citations come from rather than how many: the share of a paper’s citations originating in the top six accounting outlets, the share originating in seven top finance and economics journals, and the mean prestige of the citing outlets. We compute the mean prestige by matching each citing journal to its Scopus prestige scores, SJR (Scimago Journal Rank) and SNIP (source-normalized impact per paper), in the year the citation appears, and averaging across the paper’s citations.<sup>7</sup> Because SJR and SNIP exist only for Scopus-rated outlets, the prestige means are computed over the Scopus-rated subset of a paper’s citers. We also report variants that exclude the six journals of our sample, which isolate the prestige earned outside the field’s leading journals, in Table OA2.

### 3.3 Sample and treatment timing

For our main sample we keep archival and experimental research papers with an abstract and a classifiable introduction, yielding roughly 4,600 papers in the six accounting journals. We assign treatment timing at the submission date, since the policy binds at submission. To determine the submission date we use the manuscript’s received date where observed, and otherwise impute it from the publication date net of a calibrated review lag, dropping papers whose imputed timing straddles the January 2015 cutoff ambiguously. We classify a paper as *Post* if its assignment date falls on or after January 1, 2015, and as *JAR* if it is published there. Our main sample window is the three years on either side of the policy (submission years 2012–2017); in robustness checks, we widen this window. Because *JAR* reports a received date for every article in the window, the treated papers do not need their submission dates imputed: their treatment timing is observed directly, and all imputation falls on the control journals. Imputation is concentrated in *RAST* (no received dates reported) and *CAR*, and

---

<sup>7</sup>The accounting set comprises the six journals of our sample universe: *JAR*, *TAR*, *JAE*, *CAR*, *RAST*, and *Management Science* (on the citing side, all *Management Science* articles count, not only its accounting section). The finance–economics set comprises the *Journal of Finance*, the *Journal of Financial Economics*, the *Review of Financial Studies*, the *Quarterly Journal of Economics*, the *American Economic Review*, the *Journal of Political Economy*, and *Econometrica*.

the overall imputed share is stable across the pre and post periods (28.6% versus 30.0%).<sup>8</sup>

Table 1 reports descriptive statistics for the main outcomes and controls. The controls used (paper length, the number of authors, mean author *h*-index and career age, and an indicator for authors at top research departments) are the paper- and author-level characteristics available for both treated and control journals. Because authorship is key to the selection story, Table OA3 reports team and author characteristics by journal, and Table OA4 indicates that post-mandate *JAR* tilts toward more top-20-department and, directionally, more junior authors relative to controls. Author teams average roughly two to three authors, and *JAR* sits at the top of the field on author seniority and top-20 affiliation, alongside *JAE*. The author panel underlying the within-author estimates covers 9,770 author-paper observations among 3,774 unique authors in the accounting sample, of whom 1,794 appear more than once and 603 publish in both *JAR* and a control journal. The within-author estimates in Section 5.2.1 use only the part of the panel observed inside the 2012–2017 estimation window (2,500 author-paper observations, 810 repeat authors), which is why the repeat-author counts reported there are smaller than the 1,794 above.

Figure 1 plots the hedging rate, readability, and novelty in *JAR* over time, relative to the pooled accounting controls. The figure shows the yearly difference between the *JAR* mean and the pooled accounting-control mean, normalized so that the average 2012–2014 difference is zero. Before the policy, the difference varies from year to year but shows no trend. After 2015 it shifts, toward lower hedging and complexity and higher novelty. The pre-to-post movement in each panel is close to the regression difference-in-differences reported alongside it (readability  $-1.30$  in the data,  $-1.25$  estimated). The confidence bands are wide because the *JAR* side of the difference is based on only a few dozen papers per year. That is why the analyses below examine parallel trends directly and run design-based inference.

## 4 Empirical Strategy

### 4.1 Difference-in-differences


Our baseline compares papers in *JAR* to papers in the control journals, before and after the January 2015 policy. For paper  $i$  in journal  $j$  submitted in year  $t$ , we estimate

$$y_{ijt} = \beta (JAR_j \times Post_t) + \gamma JAR_j + \delta_t + \theta_{k(i)} + X'_{ijt} \lambda + \varepsilon_{ijt}, \quad (1)$$

---

<sup>8</sup>Section 5.4.2 confirms that re-estimating on the observed-date-only subsample does not change our results.





**Figure 1: The raw *JAR*–control gap around the 2015 policy.** Each panel plots the yearly difference between the *JAR* mean and the pooled accounting-control mean, demeaned so the 2012–2014 average gap is zero, with a 95% confidence band from the yearly difference in means. The grey and red bars mark the mean gap over the 2012–2014 pre-window and the 2015–2017 post-window (the estimation sample); the arrow is the raw level shift  $\Delta$ , shown next to the corresponding regression difference-in-differences. Years after 2017 are displayed for context but fall outside the estimation window. Series are unadjusted; no controls or fixed effects enter this figure.

where  $y_{ijt}$  is one of the writing, novelty, reported-inference, or citation measures from Section 3;  $JAR_j$  marks the treated journal;  $Post_t$  equals one for submission years 2015 and later;  $\delta_t$  are submission-year fixed effects; and  $\theta_{k(i)}$  are research-topic fixed effects. The coefficient  $\beta$  on the  $JAR \times Post$  interaction is the difference-in-differences estimand: the change in *JAR* papers from before to after the mandate, net of the contemporaneous change in the control journals.

The fixed effects and controls are chosen to absorb the obvious alternatives. Submission-year effects  $\delta_t$  account for anything common to the field in a given year, including the broader open-science movement and secular trends in how accounting research is written and cited; identification comes only from *JAR*’s movement *relative* to the controls in the same year. Topic effects  $\theta_{k(i)}$  ensure we are not simply comparing a different mix of subjects, whose writing and citation conventions differ. The control vector  $X_{ijt}$  holds fixed paper length, the number of authors, mean author *h*-index and career age, and an indicator for authors at top-20 departments, so that  $\beta$  does not absorb a shift in the size, seniority, or institutional composition of the work *JAR* attracts; Section 5.3.2 returns to that composition as an object of interest in its own right. When a citation or *p*-value measure is itself the outcome we drop it from the controls.

We report the estimates for two comparison sets. The “Top-3” design pits *JAR* against the other two leading accounting journals, *TAR* and *JAE*, the closest match in prestige and audience. The “All-Accounting” design adds *CAR*, *RAST*, and the accounting papers in *Management Science*. The two deliver a similar picture throughout. We estimate (1) by OLS

and cluster standard errors by journal  $\times$  submission-year.<sup>9</sup>

We trace the response over time to examine the assumption of parallel trends that underlies our specification (1). Specifically, we estimate an event-study analogue that replaces  $Post_t$  with a full set of indicators for submission year relative to the 2015 cutoff, each interacted with  $JAR$ :

$$y_{ijt} = \sum_{s \neq -1} \beta_s (JAR_j \times \mathbf{1}[t - 2015 = s]) + \gamma JAR_j + \delta_t + \theta_{k(i)} + X'_{ijt} \lambda + \varepsilon_{ijt}. \quad (2)$$

The year before the policy ( $s = -1$ ) is the omitted category, so each  $\beta_s$  measures the  $JAR$ -control gap in event-year  $s$  relative to that baseline. A flat, near-zero pre-period sequence supports the identifying assumption, the post-period coefficients trace the dynamic effect.

## 5 Results

We structure our results along two dimensions. The first is *behavioral*: holding the author constant, we examine whether the rule changes how a given researcher writes and reports. The second is the *market* dimension: we assess whether the rule alters *which* authors and *which* projects the journal attracts, as well as how the resulting papers are subsequently cited.

### 5.1 The behavioral margin: what a given author changes

This subsection documents the shifts in writing and reporting found in  $JAR$  articles after the disclosure rule takes effect. We label these shifts as potential behavioral responses. We first analyze writing, then how statistical evidence is reported, and finally two sources of heterogeneity—the opacity of the underlying project and the strength of the paper’s headline finding—that shed light on possible mechanisms.

---

<sup>9</sup>The treatment is assigned to a single journal, which warrants a discussion on the standard errors. With one treated cluster, analytic cluster-robust standard errors can be anti-conservative (Bertrand et al., 2004; Cameron et al., 2008; Conley and Taber, 2011; MacKinnon and Webb, 2017). We therefore interpret the clustered  $t$ -statistics as descriptive and confirm every result with procedures that do not lean on large-cluster asymptotics: the restricted wild cluster bootstrap, an in-time placebo, and synthetic-control and synthetic difference-in-differences estimators built for a single treated unit (Abadie, 2021; Arkhangelsky et al., 2021; Roth et al., 2023). We also relabel each control journal in turn as the treated unit; because the policy was not randomly assigned and there are only six journals, we report that exercise as a placebo ranking rather than as a  $p$ -value. Appendix E reports these checks as well as multiple-testing-adjusted  $p$ -values across the family of outcomes (Anderson, 2008; Romano and Wolf, 2005).

### 5.1.1 Writing: papers become easier to read and rely less on hedging language

Table 2, Panel A presents estimates from specification (1) for the writing outcomes. After the mandate, *JAR* articles become more readable: abstract complexity decreases by 1.25 Gunning–Fog points and introduction complexity falls by 0.96, each about one-third of a standard deviation. In economic terms, a 1.25-point reduction in Fog corresponds to a little over one U.S. grade level in reading difficulty. In line with the framework’s first prediction, the biggest shifts occur in the sections readers rely on to assess the paper’s claims.<sup>10</sup>


We then break the readability improvement into its two mechanical components, since doing so lets us assess whether the result is merely cosmetic. The Gunning–Fog index equals average sentence length plus the proportion of complex, multisyllabic words, so the  $JAR \times Post$  effect can be cleanly partitioned into the contribution of each component. Figure 2 shows this breakdown by section. Changes in complex-word usage explain about two-thirds of the introduction effect and almost the entire whole-paper effect, and they generate essentially all of the statistical significance that remains after the split. By contrast, the sentence-length channel is modest and is never significant on its own. Thus, post-mandate *JAR* authors are not merely turning long sentences into shorter ones—the simplest way to improve a readability metric without altering meaning—but are instead opting for simpler vocabulary. In addition, a separate vocabulary-based measure (SMOG) tracks Fog closely (Appendix J), further supporting our conclusions.

Next we examine hedging in authors’ writing, finding that after the mandate the introduction hedging rate falls by 0.028, around 0.3 standard deviations. Read together, the writing results say that when the data construction of a paper is open to inspection, authors hedge their results less and state more directly what they did and what they found. Two results ahead qualify this reading, and we flag them now. The average hedging decline also appears in the exempt theory papers, so part of it reflects journal style (Section 5.2.3). Its concentration in the papers with the weakest headline evidence does not, since a style change cannot target the papers whose test statistics are thin (Section 5.1.4).

The writing response varies across the different parts of the paper, and this variation serves as an internal check. Figure 3 plots the readability effect section by section and the hedging effect part by part. Both line up the same way. Readability improves by 0.7 to 1.3 Fog points in the abstract, introduction, theory, and conclusions, the sections that motivate and interpret the result, while the methodology and results sections, where notation and reported numbers fix the prose, do not move. Hedging falls almost entirely in the two rhetorical areas where authors make a claim, the statement of the research gap (a drop

---

<sup>10</sup>Section 5.4.1 examines a pre-period decline in abstract readability; consequently, we use the designs in Section 5.2 to further evaluate the behavioral interpretation of the drop.




**Figure 2: The readability gain is simpler words, not shorter sentences.** The  $JAR \times Post$  effect on Gunning–Fog, decomposed by section into the contribution of its two components: average sentence length (gold) and the share of complex, multisyllabic words (crimson). The two sum exactly to the total effect (labeled at right). The word-complexity component dominates in every section and accounts for the significance; the sentence-length component is small.

of 0.087, about an eighth of its mean) and in claims of contribution ( $-0.091$ ), while the motivation, present-research, and results claims are flat. A change in editorial or copyediting standards, or a shift in *JAR*’s topic mix, would not explain why only the persuasive claims moved in both panels. The underlying coefficients and  $t$ -statistics, section by section and claim by claim, are reported in Appendix Tables OA13 and OA14.


The timing of the hedging decline also lines up with the policy. Figure 4 reports event-study estimates of (2) for the gap and contribution claims. In the gap move the *JAR*–control gap is near zero through the pre-period and turns negative once the mandate binds, the contribution claims show the same pattern more noisily, as one would expect from a sub-paragraph measure in a single selective journal. Overall, the dynamics are consistent with the framework’s prediction that authors become less guarded about their gap and contribution claims once the paper’s data trail and code will be public.

### 5.1.2 Reported inference: more nulls on the page, no robust retreat from the thresholds

We next examine how authors report the statistical evidence beneath the prose. If the mandate discouraged tuning a specification until it just cleared a significance cutoff, the test



**Figure 3: The mandate moves persuasive language, not mechanical reporting.**  $JAR \times Post$  estimates (All-Accounting; 95% confidence intervals). Left: readability (Gunning-Fog) by section. Right: hedging rate by introduction rhetorical move. Filled markers are framing/claim-making components; hollow markers are mechanical reporting. The effect concentrates in the framing components and is absent in the mechanical ones.



**Figure 4: Hedging in the claim-making moves, by event time.**  $JAR \times$  event-time coefficients from (2) for the research-gap and contribution rhetorical moves (reference year -1; 95% confidence intervals clustered by journal×submission year).


statistics authors report should pull away from the conventional thresholds and toward the null. We find that half of this prediction holds: post-mandate *JAR* papers leave more of their insignificant results on the page, but the apparent retreat from the one-percent threshold is a migration to the five-percent threshold rather than a move into insignificance.

Table OA5 establishes the baseline. Reported significance is the norm in this literature: across every extracted coefficient the median reported  $p$ -value is 0.13 and 62.6% clear the five-percent cutoff, and among the single most-important result per paper the median  $p$ -value is 0.002 and 92.5% clear five percent. The introduction- and abstract-discussed results fall between these two extremes. Star-imputed rows move these levels trivially—the with- and without-imputation columns never differ by more than 0.001—so the threshold aggregates that follow are built on genuine test statistics only.

Table OA6 estimates the specification (1) on the share of the reported results of a paper falling in each  $|z|$  band, for four nested sets of results: every coefficient extracted (Panel A), the results discussed in the introduction (Panel B), the primary hypothesis tests (Panel C) and the single headline result (Panel D). Across the universe of extracted coefficients, the share of statistically insignificant results ( $|z| < 1.65$ ) rises by 4.3 percentage points against *TAR+JAE* ( $t = 2.53$ ) and by 4.4 points against all accounting controls ( $t = 2.23$ ), from a pre-mandate *JAR* mean of 27.7%. To put this magnitude in context, this is a relative increase of about one-sixth in the share of reported nulls. No single significant band offsets the gain, and mean  $|z|$  is essentially unchanged, so the rise reflects added mass at the null end of the distribution rather than a reshuffling across thresholds. Section 5.1.3 identifies the source as authors making fewer headline claims and surrounding each with more supporting estimates, disproportionately weaker ones, so the null share rises because reporting is fuller.

The evidence at the thresholds themselves is weaker. Reading down the salience hierarchy, the share of showcased results bunched just above the one-percent cutoff falls, and it falls most where the incentive to push a statistic across a line is largest: by 0.077 for the primary results that carry each paper, against 0.015 across the universe of incidental coefficients (Brodeur et al., 2016, 2020). Taken alone, this pattern would suggest deterred specification search. Three additional pieces of evidence argue against that reading. First, the mass that leaves the just-above-one-percent caliper does not move into the null region; it reappears just above the *five*-percent cutoff, where the showcased share rises by about 0.10. Second, the overall share of significant showcased results, if anything, increases, and the null share among showcased results does not rise. Figure 5 shows the corresponding raw shift. The pre-policy spike in the density of absolute  $z$ -statistics in the  $[2.0, 2.6)$  region, just above the five- and one-percent cutoffs, thins after the policy, but the mass does not collapse into the body below  $z = 2$ ; much of it settles just above the five-percent cutoff, consistent with statistics

relocating across thresholds rather than leaving significance.



**Figure 5: Reported test statistics in *JAR*, before vs. after the policy.** Density of the absolute  $z$ -statistic of reported results in *JAR* papers, pre- and post-mandate. Dotted lines mark the ten-, five-, and one-percent two-sided thresholds. Post-mandate, the spike just above the one-percent threshold thins while mass appears just above the five-percent threshold, the pattern of a relocation across cutoffs rather than a retreat from significance (Section 5.1.2).

Third, with a single treated journal, the one-percent decrease rests on a placebo comparison across six journals: *JAR* ranks first when each control journal is relabeled as treated, which is suggestive but cannot establish the effect on its own. We therefore treat the test-statistic channel as an unresolved null, consistent with reported statistics sorting one threshold down rather than with less specification search. Appendix Figure OA3 reports the complete caliper picture for the various thresholds.

### 5.1.3 Reporting discipline: fewer headline claims, more supporting evidence

The reported-inference results of Section 5.1.2 raise a further question: do post-mandate *JAR* papers test more thoroughly, or do they report fewer, stronger headline numbers? Following the test-multiplicity logic of Brodeur et al. (2020), we decompose each paper’s reported results into its headline (“primary”) results and the full census of reported coefficients, and estimate equation (1) on the counts.

Table 3 reports the results. First, *JAR* papers report *fewer* headline results after the mandate with the log count of primary results falling by 0.11, which translates to about 10% fewer headline claims in proportional terms, with the raw mean dropping from 2.6 to 2.0 per paper. Second, the total number of reported coefficients does not fall. If anything it rises,

from roughly 49 to 65 per paper, with the number of tables unchanged, so the ratio of total to headline coefficients climbs sharply (a difference-in-differences of about +14). Put differently, each headline claim in a post-mandate *JAR* paper is surrounded by many more supporting estimates. This is the pattern one would expect when the pipeline behind a result is open to inspection, authors make fewer headline claims and back up each with more robustness. This pattern also has a selection flavor, since reporting fewer headline results is consistent with trimming marginal findings. On this margin the discipline and selection channels seem to go in the same direction. Ellison (2002) documents the same margin in economics publishing more broadly: over three decades referees came to demand more robustness and generality around each claim, papers grew longer, and review times lengthened. A disclosure mandate raises the return to that investment for one journal at a known date.

To disentangle the behavioral from the compositional margin in this case we next focus on within-author variation. Absorbing author identity with author fixed effects (the design Section 5.2.1 formalizes), we find that the same authors report about 15 more coefficients in their post-mandate *JAR* papers than in their other work ( $p = 0.055$ ; Table 4, Panel B). That is roughly five-sixths of the cross-sectional gap of 18. The rise is not driven by the more junior or top-department entrants as triple-interactions are insignificant. To isolate what these additional coefficients are, we turn to their significance levels and magnitudes. We find that the share of reported results that have clear conventional significance falls by seven to eight percentage points, while the mean magnitude of the test statistics is unchanged. This combination suggests that the coefficients that authors are adding are disproportionately weaker, non-significant ones. The added mass is in the supporting coefficients, not the main ones, which is how the falling overall significant share coexists with Section 5.1.2’s finding that the significant share among showcased results, if anything, rises. In other words, post-mandate *JAR* papers seem to show more of the supporting evidence and lean less on a few headline numbers, with no inflation in the strength of the statistics they report (see Section 5.1.2 and Appendix Figure OA3).<sup>11</sup>

Next, if the response is driven by disclosure cost, it should be larger where that cost is larger, such as more opaque projects whose data and pipeline are harder to document and share. We test this prediction directly. From the disclosed data descriptions in papers, extracted with the same Qwen 3.6 27B model used for reported inference, we construct a continuous project-opacity index (private versus public ownership, restricted versus open access, proprietary-private status, with a Compustat-like adjustment that lowers the opacity

---

<sup>11</sup>We note that this is a *JAR*-specific response: *Management Science*’s 2019 mandate (the second treated journal of Section 5.2.2) does not reproduce it, so we read it as discipline particular to the *JAR* setting rather than a general property of disclosure mandates.



score for canned public panels that can simply be acquired for nominal fees) and add a triple interaction  $JAR \times Post \times Opacity$  to the baseline.

Table OA7 reports the estimates, and the evidence is mixed. The two most language-elastic outcomes move in the predicted direction: abstract readability improves more, and introduction hedging falls more, for higher-opacity projects (about  $-0.40$  Fog and  $-0.021$  hedging per standard deviation of opacity).<sup>12</sup> The same relation does not appear in main-text readability, reporting discipline, or novelty. When opacity is proxied by the single proprietary-private flag, the readability gradient remains ( $-0.52$  Fog per standard deviation) but the hedging gradient does not, so we regard the channel as suggestive. Whether the mandate also changed the opacity composition of the pool, exit rather than effort, is a sorting question, which we take up with the other composition results in Section 5.3.2. Overall our results above indicate that the disclosure cost is paid through behavior, the same authors write in simpler language on their harder-to-document projects.

#### 5.1.4 Heterogeneity by the strength of the headline result

Next we focus on the incentive side. If the mandate deters authors from overstating thin evidence, the writing effects should be concentrated where overstatement pays most, in papers whose headline result is weak. To examine this, we split papers at the sample-median headline test statistic ( $|t| = 3.12$ , taken from each paper’s self-identified most important result) and add  $JAR \times Post \times Weak$  to the baseline specification.

Figure 6 graphs the net  $JAR \times Post$  effect for weak (filled) and strong (hollow) papers, and Appendix Table OA8 reports the underlying triple-interaction estimates. The decline in hedging is carried almost entirely by the weak papers. Introduction hedging falls by  $0.053$  among weak papers against an insignificant  $-0.011$  among strong ones, and contribution hedging falls by  $0.185$  against an insignificant  $-0.029$ . The difference between weak and strong papers is significant in both cases (triple interactions of  $-0.042$  and  $-0.156$ ). In terms of magnitude, the weak-paper effects are five to six times larger than the strong-paper coefficients. In terms of pre-period standard-deviation units, these are declines of roughly two-thirds to a full standard deviation. These results suggest that the full-sample hedging declines documented above are not a broad shift in the journal style but rather they come from the papers whose headline evidence is thin.

Three additional checks discipline our reading of the result. First, the concentration is not mechanical. Pre-mandate,  $JAR$  authors already hedged their contribution claims more when the headline evidence was weaker (a slope of  $-0.058$  per log point of headline strength),

---

<sup>12</sup>Both estimates are similar under a provenance-only index that drops the methods-length and experimental components (abstract readability  $-0.44$ ,  $t = -1.9$ ; introduction hedging  $-0.018$ ,  $t = -2.1$ ).

and the raw within-*JAR* correlation between contribution hedging and headline strength falls from  $-0.33$  before the mandate to  $-0.12$  after. The mandate does not shift all weak papers by a constant, rather it compresses precisely the strategic margin on which prose tracked evidence. Second, the concentration is specific to claim-making language. Gap hedging is the exception in Figure 6, falling for strong papers ( $-0.152$ ,  $t = -4.24$ ) rather than weak ones (a net of  $-0.065$ ), and the readability improvements are shared by both groups, the weak-paper abstract-Fog net difference of  $-1.34$  is at most marginal, and neither readability triple interaction is distinguishable from zero. The mandate changes how strongly authors claim a contribution, not the complexity of their sentences.


Finally, the concentration does not just reflect the selection of papers. If the mandate pushed the weak papers with the most inflated claims to *JAR*'s neighbors, the weak papers still at *JAR* would be the ones that never overstated their evidence, and measured hedging would fall with no author changing a word. For that story to be true, the weak share of *JAR*'s pool must fall, and it does not. Relative to controls, the share of *JAR* papers whose headline statistic falls below the sample median rises by about eleven percentage points after the mandate (*JAR* ranks first among the six accounting journals under cross-journal relabeling). Within *JAR*, the raw share moves from 46 to 54 percent of the measurable pool, and the count of weak papers is flat at roughly eleven per year (33 pre, 32 post).<sup>13</sup> If anything, the mandate leaves *JAR* carrying more thin-evidence papers, not fewer, so the hedging decline among them reads as authors writing differently, not different authors writing. Put differently, once readers can check the evidence, the writing response is concentrated exactly where the incentive to overstate thin evidence is strongest.

## 5.2 Behavior or composition? Three tests that distinguish the margins

Any of the behavioral findings so far could, in principle, be produced by composition alone, the same reduced-form change would appear if the mandate merely re-sorted who publishes in *JAR*. We use three designs to separate these two margins. First, we exploit within-author variation, which holds the researcher fixed, so any remaining effect is the same person writing differently. Second, we reproduce each test on an independent mandate at a second treated journal. Finally, for the text outcomes, we use theory papers, which carry no data or code and

---

<sup>13</sup>This split requires a numeric headline statistic, and the share of *JAR* papers carrying one falls by about 15 percentage points relative to controls as designs whose headline is not a  $t$ -statistic enter alongside the rare-topic broadening of Section 5.3.1, so we read the weak-share rise as a statement about the measurable pool. Selection could still operate through entry (i.e., weak papers arriving after the mandate may be written by authors who never hedged), but that margin is what the within-author evidence of Section 5.2.1 tackles.



**Figure 6:** Writing effects by headline-result strength. Net  $JAR \times Post$  effects for papers whose headline test statistic is below (filled) versus above (hollow) the sample median, in pre-period standard-deviation units with 95% confidence intervals (analytic, clustered journal  $\times$  year). Estimates from the triple-interaction specification of Appendix Table OA8; weak-paper net effects are delta-method sums of  $JAR \times Post$  and the triple interaction.

are therefore exempt from the mandate, as a within-field placebo. Findings that survive all three designs we read as behavioral, while findings that do not, we attribute to the mandate re-sorting the pool of papers.

### 5.2.1 Within-author estimates and robustness to selection on observables

The first and most demanding test holds the researcher fixed. Asking, how much of the measured change in  $JAR$ 's output is a change in behavior, and how much is it a change in the types of authors that publish in  $JAR$ ? Two complementary exercises bound the answer, and together they localize the mechanism: first we reweight the controls to look like  $JAR$  on observables, then we utilize author fixed effects.

We first ask whether the reduced-form effects survive by making the control group look like  $JAR$  on observables. Using entropy balancing (Hainmueller, 2012), we reweight the control journals to have the same average author  $h$ -index, seniority, team size, paper length, and top-department affiliation with those of  $JAR$ , and re-estimate specification (1) on the reweighted sample. The estimates are essentially unchanged (Table 4, Panel A), with hedging in the introduction of the papers falling by 0.029, the abstract complexity by roughly 1.19 Fog points, and novelty rising by 0.013, each with the same sign and magnitude as the baseline. This stability suggests that the behavioral results are robust to selection on the observable characteristics of authors that we can measure.

Next, we go on to absorb the author identity directly. Expanding the data to author-paper observations and adding author fixed effects, we identify  $\beta$  only from within-author variation, which means the same researcher publishing in  $JAR$  after the mandate relative to their other

work. The within-author estimates collapse to zero (Table 4, Panel B). For example, the introduction-hedging coefficient is  $-0.002$ , abstract complexity  $+0.08$ , and novelty  $+0.003$ , and the same holds in the subsample of “switcher” authors who publish in both treated and control journals.

Together, the two tests localize the margin of selection. The cross-sectional effects survive balancing on observables but vanish once author identity is held fixed, which means that the relevant selection is not on the handful of covariates a referee would think to control for, but on the identity, and the largely unobservable type, of the authors who choose *JAR* after the mandate. This is the selection channel (prediction G1) and is the most important qualification for a purely behavioral reading of the reduced form estimates. One caveat cuts the other way, the within-author contrast nets out any response an author carries to all their outlets, so if the mandate changed writing habits portfolio-wide, this design would miss it, the null bounds the *JAR*-specific behavioral component, not the total. Additionally, we caution that the within-author design is demanding in a single small treated journal, only about 810 authors contribute repeat observations and 227 are switchers. The minimum effect the within-author estimator can detect at conventional (80 percent) power ranges from 27 to 37 percent of an outcome’s standard deviation across the four columns. For abstract readability and novelty, that detection threshold lies below the cross-sectional estimate, so the within-author null is informative in the sense that the behavioral component is genuinely smaller than the cross-section. For hedging and citations, the threshold is at or above the cross-sectional estimate, so there the within-author null reflects limited power as much as composition.

Finally, the cross-sectional effects appear widespread rather than concentrated. When  $JAR \times Post$  is interacted with an indicator for experimental (as opposed to archival) studies, and separately with an indicator for junior authors, we find no statistically significant heterogeneity across any channel (Table 4, Panel C), consistent with a broad-based shift that is not limited to a particular methodology or career stage.

### 5.2.2 A second mandate and the dose-response design

The main constraint of our design is that it includes only one treated journal. To address this, we also study *Management Science*’s move in 2019 to a more stringent data- and code-sharing policy (requiring complete analysis and table code, log files, and, starting in 2020, an active data editor who re-runs submissions), which provides a second treated journal. When we treat its accounting articles as an additional treated unit in a two-way fixed-effects framework (journal and submission-year effects), readability rises by 0.63 Fog points after 2019, as shown in Table 5. This corresponds to roughly three-fifths of a U.S. grade level, the same direction

and about half the size of the *JAR* effect. As we explain below, the appropriate benchmark is *JAR*’s behavioral component rather than the entire cross-section, and on that margin the estimates align closely.

The fact that an independent journal with a much stronger rule moves the same outcome the same way is a strong external check to our readability inferences. The second journal also supplies the within-author test that *JAR*’s thin panel cannot power. Absorbing author fixed effects alongside the journal’s own style, the same authors write post-2019 *Management Science* abstracts that are 0.64 Fog points more readable than their other work, with the main text moving by  $-0.46$ , these within-author magnitudes match the journal-level estimate. However, their hedging, as at *JAR*, does not fall.<sup>14</sup> Our behavioral interpretation of readability therefore does not rest on cross-author variation at either treated journal. The hedging, novelty, and citation channels are *JAR*-specific. Next, we use the full cross-journal variation in policy stringency. We code every journal-period on an eight-item index that counts what each policy actually requires—sample-construction code, analysis code, table-generation code, a comprehensive log file, a prohibition on manual data steps, mandatory sample identifiers, active data-editor verification, and a disclosure document at submission—read directly from the policy texts (Appendix A defines the index and documents the coding). Within our window the index runs from 0 at the never-mandating journals, to 2 at *JAR* from 2015 (and at *CAR*, which adopted a comparable disclosure rule in mid-2020), to 5 and then 6 at *Management Science*, whose 2019 policy requires the full code chain and log files and whose 2020 revision adds the data editor. The eight items are a formal data-disclosure document at submission, sample-construction code, statistical-analysis code, table-generation code, a comprehensive log file, mandatory sample identifiers, a prohibition on manual data manipulation, and active editorial verification; *JAR*’s 2015 policy scores the first two, *Management Science* scores five in 2019 and six once its data editor began work in 2020, and *JAR*’s 2024 revision, outside our window, scores seven. Regressing each outcome on the index with journal and submission-year fixed effects, readability scales with the strength of the rule: a one-unit increase in the index lowers Fog by 0.15 points, statistically significant whether we cluster on journal $\times$ year or on the six journals, and similar under an ordinal recoding of the same record. Because only

---

<sup>14</sup>Author–paper panel across the six accounting journals, 2008–2021 submissions: 7,747 observations on 1,791 repeat authors, with author, submission-year, and topic fixed effects, the *Management Science* journal effect, and the controls of equation (1); standard errors clustered by author. The abstract estimate is  $-0.75$  restricting to the 2016–2021 window around the policy, and is unchanged controlling for *JAR*  $\times$  *Post*. The journal-level hedging coefficient in Table 5 is positive ( $+0.026$ ) and significant but does not read as a policy response: an event study shows *MS*–control deviations of the same size well before the rule (2012, 2013, and 2017), the coefficient falls to  $+0.015$  over the full 2008–2021 window and to  $+0.010$  without the 2020–2021 submission years, no weak-paper gradient accompanies it (the triple interaction is insignificant), and relabeling a control journal with the same 2019 date produces a drift of comparable size. We therefore read the *MS* evidence as the absence of a hedging response, not a perverse one.

two journals mandate within the window, the index’s variation is concentrated: the estimate keeps its sign when any single journal is dropped but loses precision when either treated journal is excluded, so we read the dose-response as corroboration for the second-journal design rather than stand-alone evidence. No other outcome shows a similar dose-response.


Across both extensions, readability is the robust, generalizable result, while the remaining channels are particular to *JAR*. Additionally, *JAR*’s informative within-author null caps the behavioral component below the  $-1.25$  cross-section, and *MS*’s within-author estimates put that component near  $-0.6$ . We interpret this as the remainder reflecting the sorting that Section 5.3.2 documents directly, which suggests that the headline readability effect is roughly equal parts behavior and sorting.

### 5.2.3 Within-journal placebo: theory papers

Given that the mandate requires authors to file datasheets and post the code that builds their samples, it can only bind on empirical work. Theory papers in *JAR* are therefore exposed to the same editors, journal style, and time trends as their empirical counterparts but are not treated by the disclosure rule. If the readability effect is driven by the mandate rather than a journal-wide drift, it should appear for empirical papers and be absent for theory papers.

To test this, we reintroduce the analytical papers we otherwise exclude and estimate the difference-in-differences separately by paper type, and a triple difference that interacts  $JAR \times Post$  with an indicator for empirical work. The separation is clean, as shown in Figure 7. For empirical papers, abstract readability improves by 1.25 Fog points, about one and a quarter U.S. grade levels, the introduction by 0.96, and the main text by 0.53. For theory papers, the same coefficients are near zero and statistically indistinguishable from it. The triple-difference coefficients carry the right sign and a magnitude close to the empirical-only effect, and *JAR* is the most extreme of the nine journals when each is relabeled as treated.

Despite the results, we are cautious in our interpretation here since *JAR* publishes only about two dozen theory papers per period, so the triple difference is estimated imprecisely, its standard error is several times that of the empirical-only estimate, and the design could not reject even an empirical-sized effect most of the time. We therefore read the theory papers as corroborating the disclosure interpretation of the readability result, alongside the within-author and second-journal evidence, rather than as an independent test. The contrast does not extend to hedging: empirical and theory papers hedge less by nearly the same amount, which points to a journal-wide change in style, not a response to the mandate, and reinforces our reading of readability, not hedging, as the behavioral disclosure channel.



**Figure 7: The mandate moves empirical papers, not theory papers.** Left:  $JAR \times Post$  effect on readability (Gunning–Fog) estimated separately for empirical papers (the mandate binds) and analytical/theory papers (it does not); 95% confidence intervals. Right: the abstract-readability triple-difference coefficient with each of the nine journals relabeled as treated; *JAR* (diamond) is the most extreme.

### 5.3 The market margin: compositional shift in authors and works

We next turn to the market margin asking whether the mandate changed which authors and which projects *JAR* attracts, and how the resulting papers are cited. These are selection effects rather than changes in any one author’s conduct, given our results above. We specifically examine the novelty and topic composition of the work first, then the composition of the author pool, then citations, and we close by testing the assumption that the control journals are themselves untreated, which matters for our inferences.

#### 5.3.1 Novelty: papers move further from the prior literature

We first ask whether the work *JAR* publishes moves away from the prior literature. The framework predicts it should, if a marginal empirical result is harder to defend once the data trail and code are open, the authors who continue to choose *JAR* compete on a different margin, and the natural one is the novelty of the idea. Table 2, Panel B reports the estimates. The embedding distance of the introduction from its twenty nearest predecessors rises by 0.013, roughly a third of a standard deviation, and abstract novelty by 0.016. The share of papers that explicitly assert novelty rises by about five percentage points but is not statistically significant. We interpret that rhetorical margin apart from the distance measures, since the two turn out to be close to uncorrelated.

The distance result holds against a synthetic control<sup>15</sup> (Appendix E), and it is the one outcome on which *JAR* separates from the placebo journals. Holding the author fixed, however, the novelty change is essentially zero (Section 5.2.1), the mandate does not seem to make a given author write more original papers, it seems to sort more original work into the journal. Moreover, the novelty result does not hinge on how we measure novelty. It holds across every embedding variant we construct, nearest-neighbor distances at  $k$  from 1 to 50, a five-year reference window, and an originality measure, each landing between one-third and one-half of a standard deviation, and it reappears in measures built entirely from the citation network, with no text input. Post-mandate *JAR* papers draw on a wider range of fields, and the non-accounting share of their references rises by about half a standard deviation ( $t \approx 7$ ), robust to how many references they carry. The fact that these two different measures, constructed from different data, move together (Figure 8; Appendix H), makes it less likely that our inferences are driven by some data issue.


Two features of the change in novelty help characterize the kind of novelty the mandate sorts in. First, it is local rather than global, the distance from a *JAR* paper to its nearest predecessors rises sharply, but its distance to the field’s center barely moves, so the mandate pushes *JAR* into the sparser corners of the field, not into a new part of it. Second, it is concentrated in the tail. Figure 9 reports quantile treatment effects showing how the effect roughly rises from about +0.007 at the median to +0.021 at the ninetieth percentile ( $t = 2.4$ ), and the share of *JAR* papers as novel as the pre-mandate top decile rises by about ten percentage points, directionally consistent. The change is concentrated in the upper tail of the novelty distribution; the typical paper moves little. Both patterns are sharpest in financial accounting and auditing, the archival areas where the disclosure cost binds the most, and the most novel post-mandate papers, on subjects such as analyst facial structure and the information in product-launch tweets, sit further from the field’s staple questions than the papers they displace.

A natural alternative interpretation is that *JAR* simply began publishing different topics, with the embeddings picking up the relabeling. We test this directly and find that the core topic mix does not move. Across coarse topic classes, no share shifts survive a multiple-testing correction (smallest Benjamini–Hochberg  $q = 0.20$ ). The overall distance between *JAR*’s topic distribution and the controls’ (the Jensen–Shannon divergence) barely changes, and a multinomial test of the post-mandate mix against the controls’ rejects nothing. Despite these results, we further remove the issue of topical changes by re-estimating the headline


---

<sup>15</sup>Because we examine many outcomes across the four channels, some individually statistically significant coefficients are expected by chance; adjusting for multiple testing within each channel (Anderson, 2008; Romano and Wolf, 2005) leaves novelty as the channel that survives.





**Figure 8: Novelty across measures.**  $JAR \times Post$  difference-in-differences on each novelty measure, in standard-deviation units, with 95% confidence intervals (analytic clustering by journal  $\times$  year; All-Accounting sample). Crimson markers are text-embedding measures; blue markers are built from the reference network and use no text. The effect is stable across every embedding choice and reappears in the reference-diversity measures. The figure shows the distance and diversity measures; Table OA15 reports the full set, including distance to the field centroid, mean reference age, and Uzzi journal-pair atypicality.



**Figure 9: Quantile treatment effects on novelty.**  $JAR \times Post$  quantile treatment effects on introduction novelty at the 25th, 50th, 75th, and 90th percentiles (95% confidence intervals), for both comparison sets. The effect roughly doubles from the median to the 90th percentile: the mandate produces more highly novel papers rather than shifting the typical one.


novelty regression inside financial accounting alone, the dominant area holding three-quarters of the sample, with the coefficient continuing to be positive and significant on the mandate (+0.011). The novelty gain is therefore not a relabeled topic mix, but rather papers move further from their neighbors within the field’s staple areas.<sup>16</sup>

What does change is *JAR*’s reach into small topics. We define a paper’s topic as rare if its granular OpenAlex topic holds less than one percent of the control-journal corpus (the threshold is computed omitting *JAR*, so the definition cannot respond to *JAR*’s own behavior). Figure 10 reports the result. The share of *JAR* papers in rare topics jumps from 4.8 to 19.2 percent at the mandate while the controls drift from 8.3 to 10.9, a difference-in-differences of +12.9 percentage points. The change starts exactly at the 2015 submission cohort and persists through 2021. Under cross-journal relabeling, *JAR* ranks first among all six journals, the strongest statement the design permits, with a placebo coefficient 1.65 times the largest control’s in absolute value.<sup>17</sup> On a nearly constant paper count, *JAR* goes from publishing in 11 distinct granular topics to 20. These papers read as genuine boundary-crossing rather than

<sup>16</sup>If we exclude the rare-topic papers described next, this roughly halves the novelty coefficient. Tail-opening into new topics and within-topic distance are two faces of the same phenomenon, embedding novelty is higher in thin neighborhoods precisely because entering them is novel, so we describe the effect as both-and, not as purely within-topic.

<sup>17</sup>The estimate survives rarity thresholds from 0.5 to 2 percent, leave-one-year-out estimation, and defining rarity from pre-period control shares alone.

simply classifier noise, touching on topics such as machine-learning fraud detection, deception in CEO narratives, judicial ideology as litigation risk, integrity-hotline field work, and the politics of accounting regulation.<sup>18</sup> Read together with the stable core, the composition evidence suggests that the mandate did not change what *JAR* is about, rather it widened how far afield *JAR* reaches. This is the market-margin counterpart to the thicker novelty tail above (Appendix H reports the full topic-share and per-topic difference-in-differences evidence).



**Figure 10: Rare-topic share around the 2015 mandate.** A topic is rare if it holds less than one percent of the control-journal corpus (defined out of *JAR*). Left: yearly rare-topic shares, 2008–2021; bars mark the pooled 2012–14 and 2015–17 shares in the estimation sample (4.8% → 19.2%). Right: the difference-in-differences with each journal relabeled as treated; *JAR* ranks first among the six.


### 5.3.2 Selection: the composition of *JAR*’s author and project pool

Prediction G1 of our framework says that the mandate re-sorts those who publish in *JAR*. Table 6 and Figure 11 report estimates of difference differences with the author and the characteristics of the paper as outcomes. Relative to controls, post-mandate *JAR* authors are about 0.32 standard deviations more likely to be affiliated with a top-20 research department and carry a higher *h*-index (about 0.25 standard deviations). They are also directionally more junior and write somewhat shorter papers, though neither margin is statistically significant ( $p \approx 0.20$ ). The change in what *JAR* publishes therefore reflects, in part, a change in who publishes there. Oster-style bounds confirm what the balancing exercise in Section 5.2.1

<sup>18</sup>*JAR*’s rare-topic share is volatile further back (15 percent in 2008, zero in 2009–10), so the 2012–14 base is on the low side of its history; the post-mandate level is, however, the first stretch of seven consecutive years above 16 percent. And the rare bucket is a residual of small topics, so we lean on the out-of-*JAR* definition, the threshold sweep, and a continuous rarity version (log inverse control share, +0.56), all concordant.

showed, the coefficients are stable to observables, so the relevant selection is on author type (Oster, 2019).

The sorting extends to the data used in the projects. The share of *JAR* papers built on proprietary data rises by about eleven percentage points relative to controls after the mandate (a difference-in-differences of 0.11), consistent with proprietary data authors, who can substitute a documented data description for the data itself, facing a smaller compliance cost than authors of public-data work. The rise is a fact about provenance, not opacity, given that when measured with the continuous opacity index of Section 5.1.3, the composition of the pool does not move (+0.05 standard deviations).<sup>19</sup> Read together, the two results indicate that the compliance cost fell on documentation, not on proprietariness, and we find no evidence of substitution between proprietary or opaque work over our window.




**Figure 11: Selection: the composition of *JAR*'s author and project pool.**  $JAR \times Post$  difference-in-differences on each author/paper characteristic, expressed in standard-deviation units, with 95% confidence intervals (analytic clustering by journal×year). Markers are coloured by significance (red  $p < 0.05$ , orange  $0.05 \leq p < 0.10$ , grey  $p \geq 0.10$ ) and annotated with the two-sided  $p$ -value. After the mandate *JAR* tilts toward more top-20-department, higher- $h$ -index authors, with directionally (but no longer significantly) more junior authors and shorter papers.

The composition shift has two distinct sources, and separating them characterizes how the market re-sorts the papers. We divided the post-mandate *JAR* authors into those returning

<sup>19</sup>The topic-composition evidence of Section 5.3.1 corroborates this null from the other direction. The one composition margin that does move, entry into rare topics, tilts toward more opaque, more proprietary subjects (pre-period opacity 0.46 versus 0.31, proprietary-data share 46 versus 21 percent), and the opacity shift it mechanically implies (+0.019 on the index) is almost exactly the insignificant sorting estimate here (+0.017, estimated on the wider  $\pm 7$ -year window). If the mandate were deflecting opaque work away from *JAR*, the entering topics would tilt the other way.

to the journal and those publishing in it for the first time (Figure 12), and the two groups pull the composition in different directions. The new juniors come from the entrants, authors new to *JAR* after the mandate are about 2.8 years more junior than the returning pool ( $t = -6.4$ ), carry a lower  $h$ -index (14 versus 25), and are cited less than a third as often. The top-department affiliation comes from the incumbents, with returning authors sitting at a top-20 department 66% of the time, against 47% for the new entrants. Established authors anchor the journal while a broad set of junior researchers enters alongside them, suggesting a widening of the author pool at the bottom of the seniority ladder rather than a wholesale turnover, drawn from many institutions, not a few (an institutional Herfindahl of 0.010, against 0.025 among returning authors). We looked for, but do not find, a *JAR*-specific change in team size or in the within-team spread of seniority, post-mandate *JAR* teams are no more mixed in seniority than the controls', and the raw rise in mixed-seniority teams is a field-wide shift, not a response to the mandate. The one team-level margin that does move is geographic, post-mandate *JAR* papers draw on fewer distinct institutions per team (a difference-in-differences of  $-0.19$  standard deviations, which is robust to design-based inference), so the disclosed projects are somewhat more co-located, consistent with the added coordination that assembling and sharing a documented data package requires.



**Figure 12: Characteristics of post-mandate entrants and returning authors.** Mean characteristics of post-mandate *JAR* authors, split into those returning to the journal and those appearing for the first time. New entrants are more junior and less often at a top-20 department, on slightly smaller teams; the top-20 affiliation comes from the returning authors.

### 5.3.3 Citations: fewer, but from higher-ranked journals

We next examine how the market receives the post-mandate papers. Table 7 reports the estimates. Raw citations fall (Panel A), with log total citations dropping by 0.29 against


the top-3 controls, roughly a quarter fewer citations, and the field-weighted index moves with it.<sup>20</sup> The decline is not just reflecting the mechanical price of the journal’s rare-topic reach, excluding the rare-topic entrants leaves it slightly larger, and the rare-topic papers themselves are, if anything, cited more than the rest of the post-mandate pool (Section 6.1 reports this decomposition).

However, the citations that remain come from higher-status outlets. Panel B traces the citing journals across nested tiers, the share of a paper’s citations from *TAR* and *JAE* rises by 1.8–1.9 percentage points, the share from the top-six accounting journals by 2.5–2.8, the share from the top finance and economics journals by 0.2–0.3, and the share from the top-thirteen business journals by 2.8–3.0. The accounting and business tiers are each statistically significant at the one-percent level, and the small finance-and-economics tier is marginal. Panel C confirms the pattern in journal-prestige scores, the mean SJR of citing outlets rises by 0.38–0.43, which is roughly a sixth of the pre-mandate mean citing-journal SJR of 2.428, and the mean SNIP by 0.15–0.16. The estimates are consistent across the two comparison sets and sharper, if anything, in the broader All-Accounting design. Table OA2 bounds how far the prominence extends: recomputing the mean SJR and SNIP of citing journals after dropping the top-6 accounting journals entirely as citers, the increases shrink to about 0.18 SJR and 0.09 SNIP and are no longer statistically significant. The rising prestige of citing outlets is therefore concentrated among the top accounting and business journals.

Two further features describe the shift. First, the shift is of the whole distribution of citing-journal prestige, not of its tails, the *JAR*  $\times$  *Post* effect on the mean citing SJR is essentially the same, between +0.28 and +0.39, at every decile from the tenth to the ninetieth, so the papers are cited up and down the prestige ladder by more prominent outlets rather than picking up a few marquee citations. Figure 13 shows the corresponding movement over time, the *JAR*–control gap in citing prestige is stable before the mandate and widens steadily after it. Second, the attention comes from a wider set of fields, the share of *JAR*’s citations originating inside business and accounting falls by about eight percentage points after the mandate, while the share from economics and finance, and from the social and decision sciences, rises correspondingly (Appendix Figure OA9). The shift does not come from *JAR* papers citing themselves, whose share does not rise, while the prestige gain is concentrated in the top accounting and business outlets (Appendix Table OA2). The complete set of prestige and tier measures appears in Appendix Figure OA8.

---

<sup>20</sup>Restricting the count to citations received within three years of publication gives the same sign and a similar share of the mean:  $-0.19$  against the top-3 controls and  $-0.20$  against all accounting controls.



**Figure 13: Post-mandate *JAR* papers are cited by more prominent journals.** Mean SJR of citing journals by submission year, for *JAR* and the pooled accounting controls; the dashed line marks January 2015. The *JAR*–control gap is stable before the policy and widens after it.

#### 5.3.4 General equilibrium: is the control group really untreated?

Until now, all our estimates compare *JAR* with other accounting journals and interpret the resulting gap as the effect of the mandate. That reading relies on an assumption that we are cautious off (prediction G2): that the control journals constitute an untreated counterfactual, unaffected from both *JAR*’s policy and from wider forces influencing the field. In a research marketplace, neither part of that assumption is secure. Because submission is a choice, papers diverted from *JAR* would be submitted to *TAR*, *JAE*, and others, effectively “treating” the controls via the displaced inflow of manuscripts. And because accounting is a small, highly interconnected field, secular influences (the wider open-science movement, evolving journal styles, shifting citation norms) tend to move *JAR* and the comparison journals in tandem. Either mechanism violates the stable-unit-treatment-value assumption, but they distort different contrasts. Spillovers onto the controls—such as redirected submissions that pull *TAR* and *JAE* toward *JAR*’s outcomes—would compress the treated–control difference, causing within-accounting difference-in-differences estimates to understate the true effect. By contrast, a field-wide trend that raises all accounting journals together is removed by that within-field differencing; instead, it contaminates estimates that benchmark accounting against journals outside the field, which would misattribute the shared accounting trend to *JAR* and thereby *overstate* the effect. Since the direction of overall bias cannot be determined *a priori*, we therefore try to quantify these channels.


The first test leverages a market that *JAR*'s policy is unlikely to influence. A data- and code-sharing mandate at an accounting journal should not reroute submissions to finance, and the three top finance journals (*Journal of Finance*, *Journal of Financial Economics*, *Review of Financial Studies*) introduced no similar requirement during our sample period. We therefore treat their accounting-relevant articles as a *never-treated outer benchmark* and estimate two comparisons for each writing outcome (Table 8; Figure 14). We limit this benchmark exercise to writing outcomes because these are the dimensions on which an external market can be evaluated on the same footing: readability and hedging are measured the same way in finance and accounting text, while citation counts and citation conventions differ mechanically across fields, and our novelty metric is defined relative to the accounting literature's own local neighborhood. The first comparison examines whether the controls are themselves affected: a difference-in-differences that contrasts the accounting controls (excluding *JAR*) with finance around 2015, where a nonzero estimate indicates that the allegedly untreated controls are moving, either because submissions were displaced or because of a shared trend across fields. The second comparison re-estimates the main effect using finance, rather than the accounting journals, as *JAR*'s comparison group, yielding a control that is not subject to within-field spillovers. Taken together, the two comparisons form a triple-difference: the finance-benchmarked *JAR* effect, purged of whatever movement occurred in the accounting controls regardless.

Together, the two contrasts paint a straightforward picture. In terms of readability, the finance-benchmarked estimate is comparable in magnitude to the within-accounting estimate ( $-1.19$  Fog points, slightly more than one U.S. grade level—versus  $-1.25$  relative to the accounting controls), and the contamination contrast is small and statistically insignificant. This implies the readability finding is robust whether the reference group is other accounting journals or an external benchmark. For hedging, the contamination contrast is effectively nil (coefficient  $\approx 0$ ), indicating that any hedging difference is not an artifact of the chosen control group. The broader takeaway extends beyond this specific policy: if treated and control units operate in the same market, a simple two-group difference-in-differences can mistake a market-wide shift for a treatment effect, and only an outside, never-treated benchmark can separate the two. In this case, that benchmark corroborates the writing result rather than reversing it.

## 5.4 Identification and inference

Once the two margins are disentangled, four design issues remain: whether *JAR* and the control journals followed parallel trends prior to the policy; whether the results are driven





**Figure 14: Benchmarking against the never-treated finance journals.** One panel per writing outcome (each in its own units). Within each panel: the headline  $JAR \times Post$  effect against the accounting controls (top), the same effect against the never-treated finance benchmark (middle), and the contamination contrast—the accounting controls themselves versus finance (bottom). Bars are 95% confidence intervals, coloured by significance (red  $p < 0.05$ , orange  $0.05 \leq p < 0.10$ , grey  $p \geq 0.10$ ). For readability the finance-benchmarked effect is of the same order as the within-accounting one and the contamination contrast is small and insignificant, so the result is robust to the choice of control group; for hedging the contamination contrast sits on zero, so the controls are clean.

by imputed treatment timing; whether an editorial change that happened to line up with the mandate could generate the same findings; and whether a contemporaneous shift in the journal’s preferences toward economics-style work could also account for them. We address each in turn.<sup>21</sup>

#### 5.4.1 Parallel trends


The headline difference-in-differences assume *JAR* and its controls were moving together before the policy. Figure 15 reports event-study versions of Equation (1), replacing *Post* with submission-year-relative-to-2015 indicators interacted with *JAR* (reference year  $-1$ ). For embedding novelty, the pre-policy coefficients are flat and statistically indistinguishable from zero, and the effect appears only after 2015. For abstract readability the pre-trend is less clean, a negative pre-period coefficient two years before the policy indicates that *JAR* abstracts were already trending toward lower complexity. Total citations display a pre-existing downward drift that is partly mechanical, since older papers have had longer to accumulate citations. Of the three, novelty is the only outcome whose pre-period coefficients are flat. For readability, the support that does not lean on the pre-period comparison is the within-author and placebo evidence of Section 5.2. For citations, the compositional (share-based) facts of Section 5.3.3 and the timing check below.

#### 5.4.2 Robustness to imputed treatment timing

Because we assign treatment at submission and submission dates are imputed for some control papers, a natural concern is that the imputation drives the results. Table OA9 shows that it does not. First, the treated unit needs no imputation as *JAR* reports a received date for every article in the window, so all imputation falls on the controls, where only around nine percent need to be imputed for *TAR+JAE* and they are balanced across the pre and post periods (33.1 versus 33.4 percent in the broad set). *JAR* loses exactly one paper to the ambiguous-window screen. Second, re-estimating the baseline on the observed-date-only subsample, dropping every imputed observation, leaves the abstract-readability, introduction-novelty, and introduction-hedging effects unchanged in sign and significance and, if anything, slightly larger (abstract readability moves from  $-1.26$  to  $-1.32$ , novelty from  $0.014$  to  $0.015$ ), while the citation-count effect, the least precisely estimated channel, keeps its sign. This evidence suggests a mild non-differential attenuation in the full sample, not an effect manufactured by imputation, which would shrink when the imputed observations are removed. Third, the realized review lag lengthened modestly after 2015, but the *JAR*-versus-control difference in

---

<sup>21</sup>Design-based inference for the single-treated-cluster problem is done in Appendix E.



**Figure 15: Event-study (dynamic DiD) estimates.** Each panel plots the  $JAR \times \text{event-time}$  coefficients (95% confidence intervals, analytic clustering by journal $\times$ year), reference year  $-1$  (2014). The grey and red bars mark the mean of the pre-period and post-period coefficients, and  $\Delta$  is their difference—the level shift at the mandate. Flat, near-zero pre-period coefficients support parallel trends; the readability pre-period dip at  $-2$  and the citation drift discussed in the text are visible as departures of the pre-period points from the grey bar.

that lengthening is statistically insignificant (a difference-in-differences of roughly +72 days) and, by construction, moves only imputed control dates, never *JAR*'s. Since non-differential misclassification of the post indicator attenuates the estimate toward zero, our full-sample coefficients are, if anything, conservative.

### 5.4.3 Editorial change?

A final concern is that our writing and selection results may simply reflect an editor or house-style shock that happened to land at the same time as the mandate rather than the mandate itself. With a single treated journal, any *JAR*-specific shock that coincided with January 2015, an editor transition, a change in desk-reject norms, a house-style revision, is collinear with the mandate, and the field-wide year fixed effects cannot absorb it. We address the concern in two ways. First we examine changes in who edits the journal, which is directly testable. Next we examine changes in the stylistic regime imposed from the masthead, which cannot be tested head-on.

If *JAR*'s post-2015 shifts merely track a change in editors, they should disappear once the editor is held fixed. Table OA10 and Figure OA1 show that they do not. Three editors—Berger, Leuz, and Skinner—each handled at least five papers on both sides of the mandate and together account for the majority of the journal's decisions. Within this subsample, with editor fixed effects, the same before/after pattern reappears: abstracts get 1.1 Fog points more readable, introductions hedge 1.7 percentage points less, novelty rises by 0.007, and review durations lengthen by roughly 20%. The readability and hedging estimates essentially match the whole-journal figures, and all four keep their sign and significance.<sup>22</sup> The same people, editing at the same desk, changed what they accepted once the rule bound.

A falsification test sharpens the point, *JAR*'s actual masthead turnover came in 2019, not 2015, and nothing breaks there, among post-mandate papers the 2019 break on readability, hedging, and novelty is small and insignificant (readability  $-0.35$ ,  $p = 0.22$ ), and the editors who joined in 2019 look like the continuing ones on every outcome. We are careful about the reach of this test, while it rules out that the 2015 shifts are an artifact of the editors, it does not separate the demand margin from supply, since the same editor still accepts from a changed submission pool.

The diffuse version of a change in the stylistic regime imposed from the masthead leaves no editor to hold fixed, so here we rely on the separating designs. The theory-paper placebo holds

---

<sup>22</sup>The one outcome that does not move within editor is data opacity, which is consistent with the paper's null result on project sorting (Section 5.3.2). The lone exception is Skinner, whose auditing share rises from 6% to 18% over the window, his opacity change is more than twice the journal-wide figure and survives topic fixed effects on his desk alone, so it reflects a shift in the mix of papers routed to his desk rather than a change in what a given editor accepts.

the editorial regime fixed: theory papers pass through the same editors and the same desk in the same years, yet, carrying no data or code to disclose, they do not move (Section 5.2.3), which a stylistic change imposed by an editor would not respect. The second treated journal (Section 5.2.2) identifies the readability response from variation that an idiosyncratic *JAR* editorial change cannot generate, a different journal, in a different field, under different editors, moves the same way after its own mandate and does so within author, the same researchers writing in simpler language once a rule binds. The weak-paper relationship is a second pattern no editor-imposed style regime should produce, for the reason Section 5.1.4 gives. A coincident editorial shift could affect the two-group reduced form, but it cannot produce the convergence of these independent designs, which is why we rely on their inference for the behavioral claim rather than on the headline difference-in-differences alone.

#### 5.4.4 Is it simply a shift toward economics?

A sharper version of the editorial concern is that *JAR* changed tastes rather than editors. Accounting researchers often describe the journal as having turned, over the 2010s, toward economics—toward quasi-experimental designs, shock-based settings, and identification-first rhetoric—and a taste of that kind could produce more readable, less-hedged prose even without the mandate. The accounting controls cannot difference out a *JAR*-specific drift, and our topic fixed effects absorb it only where it operates through the topic label rather than within it. Unlike a diffuse house style, however, this drift is measurable in the corpus, so we test it directly. We build two paper-level proxies: an indicator for a quasi-experimental identification strategy (difference-in-differences, instrumental variables, regression discontinuity, event study, or matching) and the share of introduction paragraphs making an identification claim. Both confirm the shift: the quasi-experimental share of *JAR* papers rises from 30 to 46 percent across the window and the identification-claim share from 6.5 to 11 percent, so the question is not whether the journal moved toward economics but whether that move explains the writing results.

If the econ shift were itself a response to the mandate, the mandate sorting econ-adjacent work into the journal, then conditioning on it would control away part of the treatment effect rather than a confound. So before any horse race we ask when the shift happened. Figure 16 puts the two proxies on the left-hand side of the event-study design of Section 5.4.1. Both trend smoothly through January 2015: the estimated 2015 step is 0.04 standard deviations for the quasi-experimental share and 0.09 for the identification-claim share, each a fraction of the median step the mandate’s outcomes take at the same date (about a sixth and two-fifths, respectively), and neither is unusual among the six journals (neither ranks first among the six journals when each is relabeled as treated). The best-fitting break date for the identification-

claim share is 2018, three years after the mandate. The writing outcomes, by contrast, step at 2015 and rank first among the six journals. Put differently, the econ shift is a ramp and the writing change is a step: the shift cannot account for the timing of the writing results.

Because the shift predates the mandate and does not break at it, conditioning on it tests the confound rather than controlling away the treatment. Adding both proxies, in levels and interacted with *Post*, leaves the writing coefficients essentially unchanged: abstract readability retains 87 percent of its baseline effect ( $-1.25$  to  $-1.09$  Fog points) and introduction hedging strengthens ( $-0.028$  to  $-0.031$ ). In the bounding terms of Oster (2019), selection on unobserved culture would have to be more than four times as important as everything we already control to drive the readability effect to zero. Appendix Table OA17 reports the full set. The weak-paper hedging concentration of Section 5.1.4 is fully retained after conditioning, 100 to 105 percent of the baseline effect, still ranking first among the six journals, and this is a demanding test because quasi-experimental papers are, if anything, mildly more likely to be statistically weak in our sample. A broad taste for economics gives the hedging decline no reason to concentrate in the introductions of statistically weak papers; a rule that raises the cost of overclaiming binds precisely there.


The editor evidence reaches the same verdict as the previous subsection: the readability effect survives editor fixed effects and holds, indeed grows, within the three editors who span the mandate, and an in-time placebo that moves the break to 2012 finds nothing.<sup>23</sup> We are careful about the reach of this battery. It defends the behavioral writing results, not the novelty result: because the mandate itself plausibly sorts more novel, econ-adjacent work into the journal, “the policy sorted in economics-style work” and “the journal developed a taste for economics” are nearly the same object on that margin, and we do not use this test to separate them. For the writing claim the four tests converge—no 2015 break in the shift, coefficients stable under conditioning, the weak-paper concentration intact, and the same editors moving at the mandate but not at the placebo—and our inferences remain unchanged.

## 6 Discussion

Our results document how a transparency mandate adopted by a single journal changes writing and papers, when only that journal, and not its competitors, adopts. Field-wide adoption is a different experiment, and our framework distinguishes these two. The market-margin results are relative-price effects in the sense that novelty rises, top-department authors arrive, and attention concentrates at *JAR* because authors can still circumvent the disclosure

---

<sup>23</sup>With editor fixed effects the readability coefficient is  $-1.19$  (95 percent of baseline); within the spanning editors it is  $-1.78$ ; the fake-2012 break yields  $-0.04$  for readability and  $+0.01$  for hedging.



**Figure 16: Event-study timing of the econ-shift proxies versus abstract readability.** Standardized  $JAR \times \text{year}$  event-study coefficients (dots) with a fitted pre-trend+step+post-trend overlay (lines); submission years 2008–2021, six accounting journals, reference year 2014. Solid lines: the two econ-shift proxies (quasi-experimental design share and introduction identification-claim share) show no 2015 step—the estimated step is 0.04 and 0.09 standard deviations, and neither proxy ranks first among the six journals when each is relabeled as treated—and the identification-claim share’s best-fitting break date is 2018. Dashed line: abstract readability shifts discretely at 2015 and ranks first among the six journals under the same relabeling. The absence of a 2015 break in the proxies is what makes conditioning on them in Table OA17 a valid test.

cost by submitting elsewhere. If every journal adopted the rule, that escape valve closes, which would flatten the re-sorting across journals, and the composition gains we document, which came partly at the other journals' expense, would not aggregate to the whole field.

The behavioral results are the opposite case, a given author's simpler writing and fuller reporting do not depend on where the marginal paper goes, they reappear under an independent mandate at *Management Science*, and they scale with policy stringency, so these are the effects we would expect to survive field-wide adoption. What other things would universal adoption change becomes the more important question, with no untreated journal of similar quality to absorb high-documentation-cost projects, the adjustment must run through what gets researched and how it is documented rather than where it is sent. Our reduced form is therefore not a forecast of a field-wide mandate. It is the market's response while an outside option still exists. That outside option is narrowing. *Management Science* has required a replication package since 2019 and *Contemporary Accounting Research* code since 2020, and the American Accounting Association will require code at *The Accounting Review* from July 2026 and at its other journals from 2027, which leaves the *Journal of Accounting and Economics* and the *Review of Accounting Studies* as the only control journals in our sample without such a rule. The market-margin estimates describe how authors and projects re-sorted across outlets while alternatives existed and should not be read as the effect of the rules now being adopted; the behavioral estimates are the ones that carry over.<sup>24</sup>

## 6.1 Is the research better?

Next comes the question of whether research quality improved. A mandate can change research without improving it, so we close by asking whether the shifts we document leave the field better off, and we try to answer it to the extent the evidence allows. The question of research quality has both a private and a social side, which may not agree. Privately, *JAR* appears to come out ahead, its post-mandate papers are more readable, more novel, and cited by more prominent outlets, the portfolio an academic journal competes to assemble. Socially, the results are more mixed, and part of the private gain is a transfer rather than a creation of value.

Start with the specification fishing, the aspect that the transparency movement cares about, where the case is strongest. Post-mandate *JAR* papers leave more of their insignificant

---

<sup>24</sup>The general-equilibrium results provide a methodological lesson that extends beyond this setting. The natural control group for a single-journal mandate, the other journals in the field, is endogenous to the very behavior under study, both because submissions are reallocated and because the field moves together. The robustness of our estimates to a never-treated outside benchmark had to be verified, not assumed (Section 5.3.4). Studies that evaluate journal- or institution-level policies should benchmark against a plausibly untreated outside market wherever one exists.



results on the page, the share of reported nulls rises by about a sixth, and they back each headline claim with more supporting coefficients rather than a few strong ones. Whatever else the mandate did, it made the evidentiary record behind a *JAR* paper more complete and more scrutinizable, which is the stated objective of the policy and of the broader movement (Christensen and Miguel, 2018; Miguel et al., 2014). Plain prose points the same way, as a reader pays less to extract what a paper did and found.

However, these benefits come with costs, which we do not net out. The first is reach. Post-mandate *JAR* papers are cited about a quarter less often, although the citations that remain come from more prominent journals. To the extent that citations proxy for use, some use was lost, and lost on comparable work, not only on the papers that reached into sparser topics. Whether the field prefers a narrower, higher-status audience to a broader one is a value judgment we do not take a stand on. The second margin is the work the rule prices, and here the evidence tempers the standard worry rather than confirming it. The proprietary cost of disclosure theory that the accounting literature has studied for four decades (Verrecchia, 2001; Beyer et al., 2010) predicts that a sharing mandate deflects proprietary-data work to the journal’s neighbors. Over our window it did not, because the 2015 rule asked for a datasheet and sample-construction code rather than the data themselves, proprietary-data authors could comply at low marginal cost, as proxied by the rise in the share of proprietary-data papers in *JAR* while its measured opacity does not move (Sections 5.1.3 and 5.3.2). The redistribution we can document runs through authors instead: the top-department, highly cited researchers who sort into *JAR* would otherwise have published elsewhere in the field, so part of *JAR*’s gain is a within-field transfer, not new value. The deflection cost remains the right worry for stronger policies, here *Management Science*’s verification regime and the full analysis code required by *JAR*’s own 2024 revision impose the cost on the work the 2015 rule spared. Moreover, a single-journal design cannot measure the net effect in any case, because we never observe the projects of the authors who chose not to come.

Two forces skew even the reporting gains as pure welfare. The cleanest behavioral gain, more complete reporting, is *JAR*-specific and does not reappear at *Management Science*, so we cannot claim it as a result of the disclosure mandate. And a rule strong enough to push proprietary-data work elsewhere would trade one kind of credibility for another. For example, results that anyone can re-run, in exchange for questions only proprietary data can answer. The 2015 mandate largely avoided that trade-off, the datasheet substituted for the data, but its stricter successors do not, and whether the trade is worth making is itself the welfare question, which we are not able to answer with this data.

## 7 Conclusion

We studied how a data- and code-sharing mandate requiring authors to disclose the provenance of their data and hand over the code that builds their samples changes the research a journal publishes, using *JAR*'s January 2015 policy and a corpus of roughly 4,600 accounting articles. The reduced-form patterns are economically meaningful. Post-mandate *JAR* papers are about one Gunning–Fog grade easier to read, hedge less in the parts of the paper that state the contribution, especially where the headline evidence is weakest, score higher on every distance- and breadth-based novelty measure, leave more of their insignificant results on the page, and draw a two- to three-percentage-point larger share of their citations from leading journals. Because treatment is assigned to a single journal, we held each result to design-based inference and to explicit tests for the general-equilibrium responses a single-journal mandate provokes. The novelty effect holds against a synthetic control and is corroborated by reference-network measures that use no text. Our readability effect is corroborated by a second treated journal and a stringency dose-response, while the hedging and citation effects are large and consistently signed. We also document direct evidence that the mandate re-sorted *JAR*'s author pool, toward top-department, higher-*h*-index authors, and directionally toward more junior authors.

Our work shows how transparency requirements do not merely make existing research checkable; they change what research is written, by whom, and where it is sent. For editors and scholarly associations, the takeaway is that a disclosure rule is a market intervention whose effects run through the composition of submissions as much as the conduct of any one author, and that once every journal adopts the rule, the relative shifts we document, in who submits where and what they submit, would no longer appear, because there would be no untreated outlet to absorb them. We hope our work encourages future research to ask whether the same forces operate under registries and pre-analysis plans, and in settings with enough treated journals to identify the market response directly.

## References

- Alberto Abadie. Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature*, 59(2):391–425, 2021. doi: 10.1257/jel.20191450.
- Michael Anderson. Multiple inference and gender differences in the effects of early intervention: A reevaluation of the abecedarian, perry preschool, and early training projects. *Journal of the American Statistical Association*, 103(484):1481–1495, 2008. doi: 10.1198/016214508000000841.
- Isaiah Andrews and Maximilian Kasy. Identification of and correction for publication bias. *American Economic Review*, 109(8):2766–2794, 2019. doi: 10.1257/aer.20180310.
- Dmitry Arkhangelsky, Susan Athey, David A. Hirshberg, Guido W. Imbens, and Stefan Wager. Synthetic difference-in-differences. *American Economic Review*, 111(12):4088–4118, 2021. doi: 10.1257/aer.20190159.
- Pierre Azoulay, Joshua Graff Zivin, and Jialan Wang. Superstar extinction. *The Quarterly Journal of Economics*, 125(2):549–589, 2010. doi: 10.1162/qjec.2010.125.2.549.
- Pierre Azoulay, Joshua Graff Zivin, and Gustavo Manso. Incentives and creativity: Evidence from the academic life sciences. *The RAND Journal of Economics*, 42(3):527–554, 2011. doi: 10.1111/j.1756-2171.2011.00140.x.
- John Manuel Barrios, Filippo Lancieri, Joshua Levy, Shashank Singh, Tommaso M. Valletti, and Luigi Zingales. The conflict-of-interest discount in the marketplace of ideas. Working Paper 33645, National Bureau of Economic Research, 2025.
- Sudipta Basu. How can accounting researchers become more innovative? *Accounting Horizons*, 26(4):851–870, 2012. doi: 10.2308/acch-10311.
- M. Bertrand, Esther Duflo, and Sendhil Mullainathan. How much should we trust differences-in-differences estimates? *The Quarterly Journal of Economics*, 119(1):249–275, 2004. doi: 10.1162/003355304772839588.
- Anne Beyer, Daniel A. Cohen, Thomas Z. Lys, and Beverly R. Walther. The financial reporting environment: Review of the recent literature. *Journal of Accounting and Economics*, 50(2–3):296–343, 2010. doi: 10.1016/j.jacceco.2010.10.003.
- Robert J. Bloomfield, Kristina M. Rennekamp, and Blake A. Steenhoven. No system is perfect: Understanding how registration-based editorial processes affect reproducibility

- and investment in research quality. *Journal of Accounting Research*, 56(2):313–362, 2018. doi: 10.1111/1475-679X.12208.
- Samuel B. Bonsall, Andrew J. Leone, Brian P. Miller, and Kristina M. Rennekamp. A plain english measure of financial reporting readability. *Journal of Accounting and Economics*, 63(2-3):329–357, 2017. doi: 10.1016/j.jacceco.2017.03.002.
- Abel Brodeur, Mathias Lé, Marc Sangnier, and Yanos Zylberberg. Star wars: The empirics strike back. *American Economic Journal Applied Economics*, 8(1):1–32, 2016. doi: 10.1257/app.20150044.
- Abel Brodeur, Nikolai Cook, and Anthony Heyes. Methods matter: p-hacking and publication bias in causal analysis in economics. *American Economic Review*, 110(11):3634–3660, 2020. doi: 10.1257/aer.20190687.
- Brian J. Bushee, Ian D. Gow, and Daniel J. Taylor. Linguistic complexity in firm disclosures: Obfuscation or information? *Journal of Accounting Research*, 56(1):85–121, 2018. doi: 10.1111/1475-679x.12179.
- Colin F. Camerer, Anna Dreber, Eskil Forsell, Teck-Hua Ho, Jürgen Huber, Magnus Johannesson, Michael Kirchler, Johan Almenberg, Adam Altmejd, Taizan Chan, Emma Heikensten, Felix Holzmeister, Taisuke Imai, Siri Isaksson, Gideon Nave, Thomas Pfeiffer, Michael Razen, and Hang Wu. Evaluating replicability of laboratory experiments in economics. *Science*, 351(6280):1433–1436, 2016. doi: 10.1126/science.aaf0918.
- Colin F. Camerer, Anna Dreber, Felix Holzmeister, Teck-Hua Ho, Jürgen Huber, Magnus Johannesson, Michael Kirchler, Gideon Nave, Brian A. Nosek, Thomas Pfeiffer, Adam Altmejd, Nick Buttrick, Taizan Chan, Yiling Chen, Eskil Forsell, Anup Gampa, Emma Heikensten, Lily Hummer, Taisuke Imai, Siri Isaksson, Dylan Manfredi, Julia Rose, Eric-Jan Wagenmakers, and Hang Wu. Evaluating the replicability of social science experiments in nature and science between 2010 and 2015. *Nature Human Behaviour*, 2(9):637–644, 2018. doi: 10.1038/s41562-018-0399-z.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.
- Sean Cao, Wei Jiang, Baozhong Yang, and Alan L. Zhang. How to talk when a machine is listening: Corporate disclosure in the age of ai. *Review of Financial Studies*, 36(9): 3603–3642, 2023. doi: 10.1093/rfs/hhad021.

- David Card and Stefano DellaVigna. Nine facts about top journals in economics. *Journal of Economic Literature*, 51(1):144–161, 2013. doi: 10.1257/jel.51.1.144.
- Andrew C. Chang and Phillip Li. Is economics research replicable? sixty published papers from thirteen journals say “often not”. *Critical Finance Review*, 11(1):185–206, 2022. doi: 10.1561/104.000000053.
- Garret Christensen and Edward Miguel. Transparency, reproducibility, and the credibility of economics research. *Journal of Economic Literature*, 56(3):920–980, 2018. doi: 10.1257/jel.20171350.
- Timothy G. Conley and Christopher Taber. Inference with “difference in differences” with a small number of policy changes. *The Review of Economics and Statistics*, 93(1):113–125, 2011. doi: 10.1162/rest\_a.00049.
- Partha Dasgupta and Paul A. David. Toward a new economics of science. *Research Policy*, 23(5):487–521, 1994. doi: 10.1016/0048-7333(94)01002-1.
- Thomas R. Dyckman and Stephen A. Zeff. Some methodological deficiencies in empirical research articles in accounting. *Accounting Horizons*, 28(3):695–712, 2014. doi: 10.2308/acch-50818.
- Travis Dyer, Mark H. Lang, and Lorien Stice-Lawrence. The evolution of 10-k textual disclosure: Evidence from latent dirichlet allocation. *Journal of Accounting and Economics*, 64(2-3):221–245, 2017. doi: 10.1016/j.jacceco.2017.07.002.
- Graham Elliott, Nikolay Kudrin, and Kaspar Wüthrich. Detecting  $p$ -hacking. *Econometrica*, 90(2):887–906, 2022. doi: 10.3982/ecta18583.
- Glenn Ellison. The slowdown of the economics publishing process. *Journal of Political Economy*, 110(5):947–993, 2002. doi: 10.1086/341868.
- Russell J. Funk and Jason Owen-Smith. A dynamic network measure of technological change. *Management Science*, 63(3):791–817, 2017. doi: 10.1287/mnsc.2015.2366.
- Matthew Gentzkow, Bryan Kelly, and Matt Taddy. Text as data. *Journal of Economic Literature*, 57(3):535–574, 2019. doi: 10.1257/jel.20181020.
- Luzi Hail, Mark Lang, and Christian Leuz. Reproducibility in accounting research: Views of the research community. *Journal of Accounting Research*, 58(2):519–543, 2020. doi: 10.1111/1475-679X.12305.

- Jens Hainmueller. Entropy balancing for causal effects: A multivariate reweighting method to produce balanced samples in observational studies. *Political Analysis*, 20(1):25–46, 2012. doi: 10.1093/pan/mpr025.
- Paul M. Healy and Krishna G. Palepu. Information asymmetry, corporate disclosure, and the capital markets: A review of the empirical disclosure literature. *Journal of Accounting and Economics*, 31(1–3):405–440, 2001. doi: 10.1016/S0165-4101(01)00018-0.
- James J. Heckman and Sidharth Moktan. Publishing and promotion in economics: The tyranny of the top five. *Journal of Economic Literature*, 58(2):419–470, 2020. doi: 10.1257/jel.20191574.
- Ryan Hill and Carolyn Stein. Scooped! estimating rewards for priority in science. *Journal of Political Economy*, 133(3):793–845, 2025. doi: 10.1086/733398.
- Byoung-Hyoun Hwang and Hugh Hoikwang Kim. It pays to write well. *Journal of Financial Economics*, 124(2):373–394, 2017. doi: 10.1016/j.jfineco.2017.01.006.
- Ken Hyland. *Disciplinary Discourses: Social Interactions in Academic Writing*. University of Michigan Press, Ann Arbor, MI, michigan classics edition, 2004. ISBN 9780472030248. doi: 10.3998/mpub.6719.
- Alessandro Iaria, Carlo Schwarz, and Fabian Waldinger. Frontier knowledge and scientific production: Evidence from the collapse of international science. *The Quarterly Journal of Economics*, 133(2):927–991, 2018. doi: 10.1093/qje/qjx046.
- Bryan Kelly, Dimitris Papanikolaou, Amit Seru, and Matt Taddy. Measuring technological innovation over the long run. *American Economic Review Insights*, 3(3):303–320, 2021. doi: 10.1257/aeri.20190499.
- Christian Leuz and Peter D. Wysocki. The economics of disclosure and financial reporting regulation: Evidence and suggestions for future research. *Journal of Accounting Research*, 54(2):525–622, 2016. doi: 10.1111/1475-679x.12115.
- Feng Li. Annual report readability, current earnings, and earnings persistence. *Journal of Accounting and Economics*, 45(2-3):221–247, 2008. doi: 10.1016/j.jacceco.2008.02.003.
- Jason Miin-Hwa Lim. How do writers establish research niches? A genre-based investigation into management researchers’ rhetorical steps and linguistic mechanisms. *Journal of English for Academic Purposes*, 11(3):229–245, 2012. ISSN 1475-1585. doi: 10.1016/j.jeap.2012.05.002.

- Tim Loughran and Bill McDonald. When is a liability not a liability? textual analysis, dictionaries, and 10-ks. *The Journal of Finance*, 66(1):35–65, 2011. doi: 10.1111/j.1540-6261.2010.01625.x.
- Tim Loughran and Bill McDonald. Measuring readability in financial disclosures. *The Journal of Finance*, 69(4):1643–1671, 2014. doi: 10.1111/jofi.12162.
- James G. MacKinnon and Matthew D. Webb. Wild bootstrap inference for wildly different cluster sizes. *Journal of Applied Econometrics*, 32(2):233–254, 2017. doi: 10.1002/jae.2508.
- Deirdre N. McCloskey and Stephen T. Ziliak. The standard error of regressions. *Journal of Economic Literature*, 34(1):97–114, 1996. JSTOR 2729411.
- Edward Miguel, Colin F. Camerer, Katherine Casey, Jérémie F. Cohen, Kevin Esterling, Alan S. Gerber, Rachel Glennerster, Donald P. Green, Macartan Humphreys, Guido W. Imbens, David D. Laitin, Temina Madon, Lewis S. Nelson, Brian A. Nosek, Morten Aagaard Petersen, Richard Sedlmayr, Joseph P. Simmons, Uri Simonsohn, and Mark van der Laan. Promoting transparency in social science research. *Science*, 343(6166):30–31, 2014. doi: 10.1126/science.1245317.
- Fiona Murray, Philippe Aghion, Mathias Dewatripont, Julian Kolev, and Scott Stern. Of mice and academics: Examining the effect of openness on innovation. *American Economic Journal: Economic Policy*, 8(1):212–252, 2016. doi: 10.1257/pol.20140062.
- Panggih Kusuma Ningrum, Philipp Mayr, Nina Smirnova, and Iana Atanassova. Annotating scientific uncertainty: A comprehensive model using linguistic patterns and comparison with existing approaches. *Journal of Informetrics*, 19(2):101661, 2025. doi: 10.1016/j.joi.2025.101661.
- Brian A. Nosek, George Alter, George C. Banks, Denny Borsboom, Sara Bowman, Steven J. Breckler, Stuart Buck, Chris Chambers, G. J. Chin, Garret Christensen, Mónica Contestabile, Allan Dafoe, Eric Eich, Jeremy Freese, Rachel Glennerster, Daniel L. Goroff, D. P. Green, Bradford W. Hesse, Macartan Humphreys, John Ishiyama, Dean Karlan, Alan Kraut, Arthur Lupia, Patricia L. Mabry, Temina Madon, Neil Malhotra, Evan Mayo-Wilson, Marcia McNutt, Edward Miguel, Elizabeth Levy Paluck, Uri Simonsohn, Courtney K. Soderberg, Barbara A. Spellman, James Turitto, Gary R. VandenBos, Simine Vazire, Eric-Jan Wagenmakers, Rory Wilson, and Tal Yarkoni. Promoting an open research culture. *Science*, 348(6242):1422–1425, 2015. doi: 10.1126/science.aab2374.

- Open Science Collaboration. Estimating the reproducibility of psychological science. *Science*, 349(6251):aac4716–aac4716, 2015. doi: 10.1126/science.aac4716.
- Emily Oster. Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics*, 37(2):187–204, 2019. doi: 10.1080/07350015.2016.1227711.
- Michael Park, Erin Leahey, and Russell J. Funk. Papers and patents are becoming less disruptive over time. *Nature*, 613(7942):138–144, 2023. doi: 10.1038/s41586-022-05543-x.
- Christophe Pérignon, Olivier Akmansoy, Christophe Hurlin, Anna Dreber, Felix Holzmeister, Jürgen Huber, Magnus Johannesson, Michael Kirchler, Albert J. Menkveld, Michael Razen, and Utz Weitzel. Computational reproducibility in finance: Evidence from 1,000 tests. *Review of Financial Studies*, 37(11):3558–3593, 2024. doi: 10.1093/rfs/hhae029.
- Joseph P. Romano and Michael Wolf. Stepwise multiple testing as formalized data snooping. *Econometrica*, 73(4):1237–1282, 2005. doi: 10.1111/j.1468-0262.2005.00615.x.
- Jonathan Roth, Pedro H. C. Sant’Anna, Alyssa Bilinski, and John Poe. What’s trending in difference-in-differences? a synthesis of the recent econometrics literature. *Journal of Econometrics*, 235(2):2218–2244, 2023. doi: 10.1016/j.jeconom.2023.03.008.
- Steven E. Salterio. We don’t replicate accounting research—or do we? *Contemporary Accounting Research*, 31(4):1134–1142, 2014. doi: 10.1111/1911-3846.12102.
- Uri Simonsohn, Leif D. Nelson, and Joseph P. Simmons. P-curve: A key to the file-drawer. *Journal of Experimental Psychology General*, 143(2):534–547, 2014. doi: 10.1037/a0033242.
- Paula E. Stephan. The economics of science. *Journal of Economic Literature*, 34(3):1199–1235, 1996. JSTOR 2729500.
- John M. Swales. *Genre Analysis: English in Academic and Research Settings*. Cambridge University Press, Cambridge, UK, 1990.
- John M. Swales. *Research Genres: Explorations and Applications*. Cambridge University Press, Cambridge, UK, 2004.
- Brian Uzzi, Satyam Mukherjee, Michael J. Stringer, and Benjamin F. Jones. Atypical combinations and scientific impact. *Science*, 342(6157):468–472, 2013. doi: 10.1126/science.1240474.
- Robert E. Verrecchia. Essays on disclosure. *Journal of Accounting and Economics*, 32(1–3): 97–180, 2001. doi: 10.1016/S0165-4101(01)00025-8.



Jian Wang, Reinhilde Veugelers, and Paula E. Stephan. Bias against novelty in science: A cautionary tale for users of bibliometric indicators. *Research Policy*, 46(8):1416–1436, 2017. doi: 10.1016/j.respol.2017.06.006.

Stefan Wuchty, Benjamin F. Jones, and Brian Uzzi. The increasing dominance of teams in production of knowledge. *Science*, 316(5827):1036–1039, 2007. doi: 10.1126/science.1136099.

# Tables

**Table 1:** Descriptive statistics

|                                    | N     | Mean  | SD    | Median |
|------------------------------------|-------|-------|-------|--------|
| <i>Outcomes</i>                    |       |       |       |        |
| Abstract readability (Fog)         | 1,370 | 19.29 | 2.98  | 19.21  |
| Introduction readability (Fog)     | 1,370 | 17.82 | 1.73  | 17.76  |
| Introduction hedging rate          | 1,370 | 0.506 | 0.097 | 0.511  |
| Intro-document novelty (kNN-20)    | 1,370 | 0.532 | 0.035 | 0.531  |
| Log total citations, $\log(1 + c)$ | 1,370 | 4.16  | 1.13  | 4.20   |
| Cite share from top-6 acct.        | 1,356 | 0.118 | 0.083 | 0.106  |
| <i>Controls</i>                    |       |       |       |        |
| Paper length (words, 000s)         | 1,370 | 10.7  | 2.0   | 10.5   |
| Number of authors                  | 1,369 | 2.75  | 0.99  | 3.00   |
| Mean author $h$ -index             | 1,369 | 9.5   | 5.2   | 8.8    |
| Mean career length (years)         | 1,356 | 7.8   | 5.2   | 7.0    |
| Top-20 department share            | 1,370 | 0.384 | 0.487 | 0.000  |

*Notes.* Sample of archival and experimental research articles in the six accounting journals, submission years 2012–2017 (the headline DiD window). Readability is the Gunning Fog index (higher = more complex); novelty is mean cosine distance to the 20 nearest prior-three-year documents; hedging is the share of uncertainty-marked sentences.

**Table 2:** Main difference-in-differences estimates, by channel

| Outcome ( $JAR \times Post$ )                       | Top-3<br>coef. | All-Acct<br>coef. | Pred. |
|-----------------------------------------------------|----------------|-------------------|-------|
| <i>Panel A. Writing (Gunning–Fog; hedging rate)</i> |                |                   |       |
| Abstract readability (Fog)                          | −1.352***      | −1.250**          | P1    |
| Introduction readability (Fog)                      | −1.079***      | −0.959**          | P1    |
| Conclusions readability (Fog)                       | −1.006***      | −0.823**          | P1    |
| Introduction hedging rate                           | −0.027**       | −0.028**          | P2    |
| Contribution hedging rate                           | −0.094***      | −0.091***         | P2    |
| <i>Panel B. Novelty</i>                             |                |                   |       |
| Intro-document novelty (kNN-20)                     | +0.014***      | +0.013**          | P4    |
| Abstract novelty (kNN-20)                           | +0.020***      | +0.016***         | P4    |
| Novelty claim (indicator)                           | +0.045         | +0.049            | P4    |
| <i>Panel C. Reported inference</i>                  |                |                   |       |
| Share insignificant ( $ z  < 1.65$ ), all results   | +0.039**       | +0.039*           | P3    |
| <i>Panel D. Citations</i>                           |                |                   |       |
| Log total citations                                 | −0.287**       | −0.277*           | P5    |
| Log field-weighted citations                        | −0.213*        | −0.221            | P5    |
| Cite share from top-6 accounting                    | +0.025***      | +0.028***         | P5    |
| Mean prestige of citing outlets                     | +0.425***      | +0.382*           | P5    |

*Notes.* Each cell is the  $JAR \times Post$  coefficient from specification (1) with submission-year and topic fixed effects and paper/author controls. “Top-3” compares  $JAR$  to  $TAR+JAE$ ; “All-Acct” compares  $JAR$  to all five accounting control journals. Panel C uses the full empirical sample and the baseline control set; Appendix Table OA6, which restricts to papers whose headline result carries a  $p$ -value and adds the number of tables as a control, reports +0.043 and +0.044 for the same outcome. Stars denote analytic clustered significance and are descriptive only: Appendix E reports design-based inference appropriate for a single treated cluster. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table 3:** Mechanism: test multiplicity (discipline vs. selection)

| Outcome ( $JAR \times Post$ )        | Top-3     | All-Acct  |
|--------------------------------------|-----------|-----------|
| Log # headline (primary) results     | −0.098*   | −0.104**  |
| Log # reported coefficients (total)  | +0.153**  | +0.101*   |
| Log # tables                         | +0.006    | −0.019    |
| Multiplicity ratio (total / primary) | +15.54*** | +14.87*** |

*Notes.* Each cell is the  $JAR \times Post$  coefficient from specification (1) with the listed count as the outcome. Headline (“primary”) results are coefficients flagged as a paper’s main findings; the total is the full census of reported coefficients across all tables. For  $JAR$ , the mean number of primary results falls from 2.6 (pre) to 2.0 (post) while the mean total rises from  $\sim 49$  to  $\sim 65$ , so each headline claim is backed by more supporting estimates after the mandate—the discipline signature.  $p$ -values are analytic and subject to the single-treated-cluster caveat of Appendix E. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table 4:** Author fixed effects, entropy balancing, and heterogeneity

|                                                                   | Hedging<br>(intro) | Readability<br>(abstract) | Novelty<br>(intro-doc) | Log<br>citations |
|-------------------------------------------------------------------|--------------------|---------------------------|------------------------|------------------|
| <i>Panel A. Entropy-balanced DiD (controls reweighted to JAR)</i> |                    |                           |                        |                  |
| $JAR \times Post$                                                 | -0.029**           | -1.187*                   | +0.013***              | -0.324*          |
| $p$ -value                                                        | 0.016              | 0.058                     | 0.004                  | 0.058            |
| <i>Panel B. Author fixed effects (within-author / switchers)</i>  |                    |                           |                        |                  |
| $JAR \times Post$ , all repeat authors                            | -0.002             | +0.080                    | +0.003                 | -0.005           |
| $p$ -value                                                        | 0.889              | 0.825                     | 0.535                  | 0.960            |
| $JAR \times Post$ , switchers only                                | +0.007             | +0.214                    | +0.003                 | +0.073           |
| $p$ -value                                                        | 0.619              | 0.564                     | 0.449                  | 0.530            |
| $JAR \times Post$ , no. of reported coefficients                  |                    | +15.0*                    | $(p = 0.055)$          |                  |
| cross-sectional, paper level                                      |                    | +18.0***                  | $(p < 0.001)$          |                  |
| <i>Panel C. Heterogeneity (triple-interaction coefficient)</i>    |                    |                           |                        |                  |
| $JAR \times Post \times$ experimental                             | -0.021             | -0.467                    | +0.016                 | +0.016           |
| $p$ -value                                                        | 0.508              | 0.582                     | 0.276                  | 0.963            |
| $JAR \times Post \times$ junior author                            | -0.025             | +1.165**                  | +0.006                 | +0.246           |
| $p$ -value                                                        | 0.236              | 0.041                     | 0.352                  | 0.374            |

*Notes.* All columns use the All-Accounting sample. *Panel A* reweights control papers by entropy-balancing weights that equate the means of author *h*-index, career length, team size, paper length, and top-department share with *JAR*'s (Hainmueller, 2012); post-balancing standardized covariate imbalance is zero. *Panel B* expands to author-paper observations and absorbs author fixed effects, so the coefficient is identified from within-author variation; standard errors are clustered by author ( $N \approx 2,518$  author-paper observations, 810 repeat authors, 227 switchers). The final two *Panel B* rows report the same within-author design on the number of reported coefficients from the comprehensive census of Section 5.1.3 (2,493 author-paper rows, 801 repeat authors), with the paper-level cross-sectional  $JAR \times Post$  for comparison ( $N = 1,340$ ). *Panel C* reports the triple-interaction of  $JAR \times Post$  with the listed indicator. *p*-values are analytic and, for *Panels A* and *C*, remain subject to the single-treated-cluster caveat of Appendix E. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table 5:** A second treated journal and a dose-response design

|                                                 | Abstract<br>readability | Introduction<br>hedging | Novelty<br>(intro-doc) | Log<br>citations |
|-------------------------------------------------|-------------------------|-------------------------|------------------------|------------------|
| <i>Management Science</i> $\times$ Post (2019)  | -0.625**                | +0.026***               | -0.001                 | +0.337**         |
| Policy index (eight-item count), acct. journals | -0.151**                | +0.002                  | -0.001                 | +0.013           |

*Notes.* Row 1 treats *Management Science*'s accounting papers as a second treated unit (post-2019), in a two-way fixed-effects design with journal, submission-year, and topic effects plus controls ( $N = 1,295$ ; control journals *TAR*, *JAE*, *RAST*). Row 2 regresses each outcome on the journal $\times$ period policy index—an eight-item count of what each policy requires, coded from the journals' policy documents (defined in Appendix A, Table OA1); in our window the index runs from 0 (never-mandating journals) through 2 (*JAR* from 2015; *CAR* from mid-2020) to 5–6 (*MS* from 2019, adding a data editor in 2020)—with journal, submission-year, and topic effects across the six accounting journals ( $N = 3,480$ ). Readability is the only channel that generalizes across journals in the predicted direction and scales with policy stringency. *Management Science*'s positive hedging coefficient does not read as a policy response—deviations of the same size predate the 2019 rule and the estimate is fragile to the estimation window (see text)—and novelty and citations do not track the policy index, so the behavioral margins beyond readability are journal-specific. Standard errors clustered on journal $\times$ year. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table 6:** Selection: change in *JAR*’s author and paper composition

| Characteristic ( $JAR \times Post$ ) | DiD (SD units) | $p$ -value |
|--------------------------------------|----------------|------------|
| Top-20 department share              | +0.32          | 0.002      |
| Mean author $h$ -index               | +0.25          | 0.044      |
| Mean author career length            | −0.24          | 0.193      |
| Paper length                         | −0.23          | 0.197      |
| Number of authors                    | +0.11          | 0.595      |
| Prior top-3 pubs. per author         | +0.07          | 0.619      |

*Notes.* Each row is a DiD with the listed *characteristic* as the outcome, expressed in standard-deviation units of that characteristic. Post-mandate *JAR* authors are more likely to be at top-20 departments and carry a higher  $h$ -index; the shifts toward more junior authors and shorter papers are directionally similar but not statistically significant—direct evidence of the selection channel (G1).

**Table 7:** Citations: count, citing-journal status, and prestige

| Outcome ( $JAR \times Post$ )               | Top-3     | All-Acct  |
|---------------------------------------------|-----------|-----------|
| <i>Panel A. Citation count</i>              |           |           |
| Log total citations                         | −0.287**  | −0.277*   |
| Log field-weighted citations                | −0.213*   | −0.221    |
| <i>Panel B. Share of citations from...</i>  |           |           |
| $TAR + JAE$                                 | +0.018*** | +0.019*** |
| top-6 accounting journals                   | +0.025*** | +0.028*** |
| top-7 finance/economics journals            | +0.003*   | +0.002    |
| top-13 business journals                    | +0.028*** | +0.030*** |
| <i>Panel C. Prestige of citing journals</i> |           |           |
| Mean SJR of citing outlets                  | +0.425*** | +0.382*   |
| Mean SNIP of citing outlets                 | +0.160*** | +0.146**  |

*Notes.* Each cell is the  $JAR \times Post$  coefficient from (1). Panel A uses logged total and field-weighted citation counts. Panel B reports the share of a paper’s citations originating in nested journal tiers (SJR is the Scimago Journal Rank; SNIP, source-normalized impact per paper). Raw citations fall while the status of citing journals rises across every tier. Stars are analytic clustered significance and descriptive only. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table 8:** General equilibrium: accounting controls vs. a never-treated finance benchmark

| Outcome              | Contamination<br>(acct. controls vs. finance) |      | <i>JAR</i> vs. finance<br>(never-treated) |       |
|----------------------|-----------------------------------------------|------|-------------------------------------------|-------|
|                      | coef.                                         | $p$  | coef.                                     | $p$   |
| Abstract readability | +0.179                                        | 0.67 | −1.192                                    | 0.009 |
| Introduction hedging | −0.006                                        | 0.47 | −0.025                                    | 0.004 |

*Notes.* “Contamination” is a DiD comparing the accounting *control* journals (excluding *JAR*) to the finance benchmark around the policy: a nonzero value means the controls are themselves moved by field-wide forces. “*JAR* vs. finance” uses finance as the (cleaner) control for *JAR*. The readability effect against the never-treated finance benchmark (−1.19) is of the same order as the within-accounting estimate (−1.25), and the contamination contrast is small and insignificant. The result is robust to the choice of control group and hedging shows no contamination.

# Online Appendix

## *Mandated Transparency and the Production of Research: Evidence from a Data- and Code-Sharing Policy*

### A The History of *JAR*'s code and data policy

This appendix documents the full history of the *JAR* data policy. Table OA1 summarizes the four versions of the policy, the coding is taken directly from the policy documents posted by the journal and the accompanying FAQ.

The original “Data Policy for the *Journal of Accounting Research*,” announced in December 2014, took effect for all submissions received on or after January 1, 2015. It imposed two obligations. At initial submission, authors had to file a data-description sheet disclosing how the raw data were obtained or generated—sources, download dates, the instrument for surveys and experiments—and which authors handled the data and ran the analyses; authors using proprietary data had also to give the editors, privately, a contact at the providing organization and the terms of the data-sharing agreement. Before final acceptance, authors had to submit the computer code that converts the raw data into the analysis dataset (or, for genuinely proprietary portions, a step-by-step description precise enough for another researcher to reconstruct the sample), and the datasheet and code were posted on the journal’s website after publication. The policy covered archival, experimental, field, survey, and simulation work; purely analytical papers were exempt. It did not require the final dataset, the statistical analysis code, or the table-generation code, and compliance was self-certified—the journal appointed no data editor and ran no in-house replication.

The two subsequent updates inside and just beyond our sample window left these requirements unchanged. The November 2016 revision renamed the policy the “Data and Code Sharing Policy,” added an FAQ, and extended coverage to structural papers, without altering what authors had to submit. The November 2022 revision expanded the proprietary-data disclosure: authors must now indicate whether the data provider restricts publication, requires pre-approval of results, or retains control of the data. On the feature-count index in Table OA1, all three versions score 2 of 8.

The substantive escalation came only in April 2024, well outside our window. Version 4 requires the full replication package short of the data themselves: statistical analysis code, table-generation code, a comprehensive log file showing the execution of the entire codebase, observation identifiers for the final sample (previously “whenever feasible,” now mandatory), and an explicit prohibition on manual data manipulation outside the submitted

code. Even under Version 4, compliance remains self-certified: unlike *Management Science*, which appointed a data editor in 2020 to verify replication packages before publication, *JAR* has no active verification. Two implications follow for our design. First, within our sample the policy in force is the 2015 disclosure rule described in Section 2, a single, sharp adoption date with no confounding mid-sample tightening. Second, the 2024 increase implies that our estimates speak to a comparatively light mandate, the behavioral responses we document arise from disclosure and exposure alone, not from verification.

## **B Additional descriptive and robustness tables**

This appendix collects the descriptive and robustness tables referenced in the main text: the citing-journal prestige result excluding the top-six accounting journals (Table OA2), author-pool descriptives (Tables OA3 and OA4), the reported- $p$ -value descriptives (Table OA5), the significance-distribution difference-in-differences (Table OA6), the opacity-heterogeneity estimates (Table OA7), the writing effects split by headline-result strength (Table OA8), and the treatment-timing robustness (Table OA9).



**Table OA1:** The *JAR* data- and code-sharing policy across its four versions

| Effective for submissions from                                              | Version 1<br>Jan. 2015 | Version 2<br>Nov. 2016 | Version 3<br>Nov. 2022 | Version 4<br>Apr. 2024 |
|-----------------------------------------------------------------------------|------------------------|------------------------|------------------------|------------------------|
| <i>Panel A: Required at initial submission</i>                              |                        |                        |                        |                        |
| Data-description sheet (sources, download dates, who handled the data)      | ✓                      | ✓                      | ✓                      | ✓                      |
| Provider contact and terms for proprietary data (privately, to the editors) | ✓                      | ✓                      | ✓ <sup>a</sup>         | ✓                      |
| <i>Panel B: Required prior to final acceptance</i>                          |                        |                        |                        |                        |
| Sample-construction code (raw data → analysis dataset) <sup>b</sup>         | ✓                      | ✓                      | ✓                      | ✓                      |
| Statistical/econometric analysis code                                       | –                      | –                      | –                      | ✓                      |
| Table-generation code                                                       | –                      | –                      | –                      | ✓                      |
| Comprehensive log file of the full code execution                           | –                      | –                      | –                      | ✓                      |
| Prohibition on manual (out-of-code) data manipulation                       | –                      | –                      | –                      | ✓                      |
| Observation identifiers for the final sample                                | Encouraged             | Encouraged             | Encouraged             | ✓                      |
| Final analysis dataset                                                      | –                      | –                      | –                      | –                      |
| <i>Panel C: Enforcement</i>                                                 |                        |                        |                        |                        |
| Compliance self-certified by authors                                        | ✓                      | ✓                      | ✓                      | ✓                      |
| Data editor / in-house verification of packages                             | –                      | –                      | –                      | –                      |
| Feature-count stringency index (0–8) <sup>c</sup>                           | 2                      | 2                      | 2                      | 7                      |

*Notes.* Requirements are coded from the four policy documents posted by the journal (“Data Policy for the *Journal of Accounting Research*,” December 2014, and the “Data and Code Sharing Policy” as updated November 2016, November 2022, and April 2024) and the journal’s FAQ dated April 1, 2024. Version 1 took effect for all submissions received on or after January 1, 2015; the November 2016 and November 2022 documents state only the month of the update, so those boundaries are approximate. A checkmark denotes a mandatory requirement; “–” denotes absence; “Encouraged” denotes language of the “whenever feasible” form. <sup>a</sup>Version 3’s substantive change is expanded proprietary-data disclosure: authors must indicate whether the provider restricts publication, requires pre-approval of results, or retains control of the data. <sup>b</sup>All versions allow a detailed step-by-step description in place of code for genuinely proprietary portions of the sample-construction process. <sup>c</sup>Sum of eight binary indicators: the six code/data requirements in Panel B (counting identifiers only when mandatory), the submission-stage disclosure sheet, and active editorial verification. Version 2 renamed the policy the “Data and Code Sharing Policy” and extended coverage to structural papers; Version 3 left the code requirements unchanged; neither moves the index.

**Table OA2:** Citing-journal prestige: robustness to excluding the top-6 accounting journals

| Outcome                      | <i>JAR</i> mean |       | <i>JAR</i> $\times$ <i>Post</i> difference-in-differences |                             |                     |
|------------------------------|-----------------|-------|-----------------------------------------------------------|-----------------------------|---------------------|
|                              | Pre             | Post  | Within- <i>JAR</i>                                        | vs. <i>TAR</i> + <i>JAE</i> | vs. All-Acct        |
| Mean SJR of citing journals  | 2.428           | 3.059 | +0.704***<br>(+4.40)                                      | +0.399***<br>(+4.03)        | +0.383*<br>(+1.78)  |
| excl. top-6 accounting       | 1.485           | 1.873 | +0.388***<br>(+3.04)                                      | +0.203<br>(+1.63)           | +0.184<br>(+0.83)   |
| Mean SNIP of citing journals | 1.864           | 2.124 | +0.283***<br>(+5.26)                                      | +0.149***<br>(+4.78)        | +0.145**<br>(+2.04) |
| excl. top-6 accounting       | 1.501           | 1.693 | +0.196***<br>(+5.34)                                      | +0.093**<br>(+2.25)         | +0.088<br>(+1.13)   |

*Notes.* Each cell is a difference-in-differences estimate of the *JAR*  $\times$  *Post* effect on the mean Scopus prestige (SJR or SNIP) of the journals citing a paper in its first three years, with *t*-statistics in parentheses. *Within-JAR* regresses the outcome on *Post* in a *JAR*-only sample (topic FE, SE clustered by year); the two DiD columns add *JAR*  $\times$  *Post* with topic and submission-year FE and SEs clustered by journal  $\times$  year, over the 2012–2017 window. The *excl. top-6* rows recompute the mean citing prestige after dropping citers that are themselves top-6 accounting journals (*JAR*, *JAE*, *TAR*, *CAR*, *RAST*, accounting *Management Science*). Excluding them, the DiD increases attenuate and, against the accounting controls, are no longer significant, indicating the citing-prestige gain is concentrated among the top-6 accounting and adjacent business journals whose citation share rises. Controls *X* as in equation (1). \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .

**Table OA3:** Author and team characteristics by journal

|                              | JAR  | TAR  | JAE  | CAR  | RAST | MS   |
|------------------------------|------|------|------|------|------|------|
| Number of authors            | 2.62 | 2.70 | 2.66 | 2.94 | 2.86 | 2.76 |
| Mean author <i>h</i> -index  | 8.8  | 9.0  | 9.4  | 9.8  | 8.7  | 11.6 |
| Mean career length (years)   | 7.4  | 7.4  | 7.0  | 9.0  | 7.5  | 8.6  |
| Top-20 department share      | 0.50 | 0.37 | 0.51 | 0.29 | 0.36 | 0.33 |
| Prior top-3 pubs. per author | 2.98 | 3.19 | 3.09 | 2.26 | 2.41 | 1.45 |
| Papers                       | 162  | 436  | 195  | 269  | 135  | 173  |

*Notes.* Means of paper- and author-level characteristics for the empirical sample (archival and experimental research articles, submission years 2012–2017). *h*-index, career length, and top-3 publication counts are averaged across a paper’s authors; the top-20 share is the fraction of papers with at least one author at a top-20 research department.

**Table OA4:** Author characteristics: *JAR* versus controls, before and after the policy

|                              | <i>JAR</i> |      | Controls |      | DiD  |
|------------------------------|------------|------|----------|------|------|
|                              | Pre        | Post | Pre      | Post |      |
| Number of authors            | 2.48       | 2.77 | 2.66     | 2.86 | 0.10 |
| Mean author <i>h</i> -index  | 7.5        | 10.3 | 8.7      | 10.3 | 1.3  |
| Mean career length (years)   | 7.2        | 7.6  | 7.2      | 8.4  | -0.9 |
| Top-20 department share      | 0.42       | 0.59 | 0.38     | 0.36 | 0.19 |
| Prior top-3 pubs. per author | 2.60       | 3.40 | 2.52     | 2.72 | 0.60 |

*Notes.* Group means by treatment status and period (submission before vs. on/after January 2015). The final column is the difference-in-differences (*JAR* post–pre minus Controls post–pre). Post-mandate *JAR* shifts toward more top-20-department and more junior authors; see Section 5.3.2 for inference.

**Table OA5:** Statistical results extracted from papers in the empirical sample

|                                                  | N     | Mean    | SD        | Q1    | Median | Q3     |
|--------------------------------------------------|-------|---------|-----------|-------|--------|--------|
| <b>All results</b>                               |       |         |           |       |        |        |
| <i>Mean p-value (reported)</i>                   | 4,347 | 0.145   | 0.100     | 0.068 | 0.133  | 0.204  |
| <i>Mean p-value (with imputation)</i>            | 4,448 | 0.146   | 0.100     | 0.068 | 0.133  | 0.205  |
| <i>% significant at <math>p \leq 0.10</math></i> | 4,448 | 72.5    | 18.6      | 60.0  | 74.3   | 87.0   |
| <i>% significant at <math>p \leq 0.05</math></i> | 4,448 | 62.6    | 21.3      | 48.9  | 63.3   | 78.6   |
| <i>% significant at <math>p \leq 0.01</math></i> | 4,448 | 42.6    | 24.2      | 24.4  | 40.6   | 59.3   |
| <b>Intro-discussed results</b>                   |       |         |           |       |        |        |
| <i>Mean p-value (reported)</i>                   | 4,372 | 0.052   | 0.083     | 0.007 | 0.020  | 0.052  |
| <i>Mean p-value (with imputation)</i>            | 4,445 | 0.053   | 0.082     | 0.007 | 0.020  | 0.054  |
| <i>% significant at <math>p \leq 0.10</math></i> | 4,445 | 92.4    | 15.1      | 88.9  | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.05</math></i> | 4,445 | 84.1    | 20.7      | 75.0  | 91.7   | 100.0  |
| <i>% significant at <math>p \leq 0.01</math></i> | 4,445 | 57.4    | 30.6      | 33.3  | 57.1   | 83.3   |
| <b>Abstract-discussed results</b>                |       |         |           |       |        |        |
| <i>Mean p-value (reported)</i>                   | 4,330 | 0.046   | 0.089     | 0.003 | 0.015  | 0.037  |
| <i>Mean p-value (with imputation)</i>            | 4,414 | 0.046   | 0.088     | 0.003 | 0.015  | 0.038  |
| <i>% significant at <math>p \leq 0.10</math></i> | 4,414 | 93.4    | 16.2      | 100.0 | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.05</math></i> | 4,414 | 86.1    | 22.6      | 75.0  | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.01</math></i> | 4,414 | 60.1    | 34.7      | 33.3  | 62.5   | 100.0  |
| <b>Primary results</b>                           |       |         |           |       |        |        |
| <i>Mean p-value (reported)</i>                   | 4,354 | 0.028   | 0.077     | 0.000 | 0.007  | 0.023  |
| <i>Mean p-value (with imputation)</i>            | 4,448 | 0.029   | 0.079     | 0.000 | 0.007  | 0.023  |
| <i>% significant at <math>p \leq 0.10</math></i> | 4,448 | 96.5    | 15.0      | 100.0 | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.05</math></i> | 4,448 | 91.1    | 23.3      | 100.0 | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.01</math></i> | 4,448 | 66.0    | 40.2      | 33.3  | 100.0  | 100.0  |
| <b>Most important result</b>                     |       |         |           |       |        |        |
| <i>Mean p-value (reported)</i>                   | 4,333 | 0.023   | 0.084     | 0.000 | 0.002  | 0.016  |
| <i>Mean p-value (with imputation)</i>            | 4,448 | 0.024   | 0.086     | 0.000 | 0.003  | 0.017  |
| <i>% significant at <math>p \leq 0.10</math></i> | 4,448 | 97.1    | 16.7      | 100.0 | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.05</math></i> | 4,448 | 92.5    | 26.4      | 100.0 | 100.0  | 100.0  |
| <i>% significant at <math>p \leq 0.01</math></i> | 4,448 | 68.5    | 46.5      | 0.0   | 100.0  | 100.0  |
| <i>Mean sample size</i>                          | 4,398 | 223,241 | 2,843,096 | 921   | 7,504  | 33,776 |

*Notes.* Sample is the Table 1 empirical sample of archival and experimental research articles, limited to the 4,448 papers whose *most-important* (headline) result carries a  $p$ -value (reported,  $|t|$ -derived, or imputed from significance stars). Statistical results are coefficient estimates extracted from each paper's tables at the regression-row level; the unit in each row is the paper, and each subset row is conditional on the paper having a usable value for that statistic. *All results* uses the *comprehensive* extraction — every coefficient of every column of every table/panel — restricted to empirical papers also present in the main extraction; the *intro-discussed*, *abstract-discussed*, *primary*, and *most-important* subsets use the main extraction (the main result per table/panel, which carries those flags). *Mean p-value (reported)* uses an explicit  $p$ -value when reported, otherwise a  $p$ -value derived from a numeric test statistic via the 2-sided normal approximation. *Mean p-value (with imputation)* additionally fills a midpoint  $p$ -value for rows that report only stars (3→0.005, 2→0.030, 1→0.075, 0→0.500). *Intro-/abstract-discussed* are results whose linked discussion paragraphs appear in the introduction/abstract. *Primary* results are flagged as the paper's main hypothesis-test coefficients; *most important* is the single headline result (one per paper). Because each row counts only papers with a usable value, subsets built on a single result (e.g. *most important*) cover marginally fewer papers than subsets that pool several results per paper.

**Table OA6:** Inference: difference-in-differences on the statistical-significance distribution of reported empirical results around the *JAR* January 2015 data- and code-sharing policy

**Panel A: Aggregates across all extracted results**

|                                                     | Pre-Policy<br>( $N = 83$ ) | Post-Policy<br>( $N = 73$ ) | Difference                  |           | Difference-in-Differences          |           |                                    |           |
|-----------------------------------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|------------------------------------|-----------|
|                                                     |                            |                             | Within-JAR<br>( $N = 156$ ) |           | JAR vs. TAR + JAE<br>( $N = 760$ ) |           | JAR vs. All Top<br>( $N = 1,291$ ) |           |
|                                                     | Mean                       | Mean                        | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                              | $t$ -stat |
| $ z  \geq 3.29$<br>( $p < 0.001$ )                  | 0.246                      | 0.232                       | -0.026                      | -1.15     | -0.010                             | -0.39     | -0.015                             | -0.60     |
| $ z  \in [2.58, 3.29)$<br>( $p \in [0.001, 0.01)$ ) | 0.156                      | 0.142                       | -0.014                      | -0.58     | -0.014                             | -0.70     | -0.009                             | -0.45     |
| $ z  \in [1.96, 2.58)$<br>( $p \in [0.01, 0.05)$ )  | 0.221                      | 0.202                       | -0.023                      | -1.19     | -0.026                             | -1.34     | -0.016                             | -0.73     |
| $ z  \in [1.65, 1.96)$<br>( $p \in [0.05, 0.10)$ )  | 0.100                      | 0.096                       | +0.001                      | +0.11     | +0.006                             | +0.58     | -0.003                             | -0.32     |
| $ z  < 1.65$<br>( $p \geq 0.10$ )                   | 0.277                      | 0.328                       | +0.061**                    | +2.43     | +0.043**                           | +2.53     | +0.044**                           | +2.23     |
| <i>Mean</i> $ z $                                   | 2.62                       | 2.54                        | -0.160*                     | -1.73     | -0.019                             | -0.36     | -0.057                             | -0.77     |

**Panel B: Results discussed in the introduction**

|                                                     | Pre-Policy<br>( $N = 81$ ) | Post-Policy<br>( $N = 74$ ) | Difference                  |           | Difference-in-Differences          |           |                                    |           |
|-----------------------------------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|------------------------------------|-----------|
|                                                     |                            |                             | Within-JAR<br>( $N = 155$ ) |           | JAR vs. TAR + JAE<br>( $N = 759$ ) |           | JAR vs. All Top<br>( $N = 1,296$ ) |           |
|                                                     | Mean                       | Mean                        | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                              | $t$ -stat |
| $ z  \geq 3.29$<br>( $p < 0.001$ )                  | 0.320                      | 0.334                       | -0.023                      | -0.81     | +0.021                             | +1.05     | +0.014                             | +0.52     |
| $ z  \in [2.58, 3.29)$<br>( $p \in [0.001, 0.01)$ ) | 0.226                      | 0.214                       | -0.024                      | -0.74     | +0.000                             | +0.01     | -0.003                             | -0.12     |
| $ z  \in [1.96, 2.58)$<br>( $p \in [0.01, 0.05)$ )  | 0.294                      | 0.267                       | -0.007                      | -0.25     | -0.069**                           | -2.23     | -0.051                             | -1.48     |
| $ z  \in [1.65, 1.96)$<br>( $p \in [0.05, 0.10)$ )  | 0.079                      | 0.102                       | +0.036**                    | +2.33     | +0.022                             | +1.21     | +0.012                             | +0.61     |
| $ z  < 1.65$<br>( $p \geq 0.10$ )                   | 0.081                      | 0.084                       | +0.018                      | +0.53     | +0.025                             | +0.75     | +0.028                             | +0.83     |
| <i>Mean</i> $ z $                                   | 3.16                       | 3.26                        | -0.084                      | -0.78     | +0.176**                           | +2.07     | +0.092                             | +0.90     |

**Table OA6 (continued).** Primary results and the paper’s most-important (headline) result.

**Panel C: Aggregates across primary results**

|                                                     | Pre-Policy<br>( $N = 82$ ) | Post-Policy<br>( $N = 73$ ) | Difference                  |           | Difference-in-Differences          |           |                                    |           |
|-----------------------------------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|------------------------------------|-----------|
|                                                     |                            |                             | Within-JAR<br>( $N = 155$ ) |           | JAR vs. TAR + JAE<br>( $N = 754$ ) |           | JAR vs. All Top<br>( $N = 1,288$ ) |           |
|                                                     | Mean                       | Mean                        | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                              | $t$ -stat |
| $ z  \geq 3.29$<br>( $p < 0.001$ )                  | 0.387                      | 0.388                       | -0.026                      | -0.79     | +0.015                             | +0.33     | +0.006                             | +0.12     |
| $ z  \in [2.58, 3.29)$<br>( $p \in [0.001, 0.01)$ ) | 0.259                      | 0.216                       | -0.046                      | -1.08     | -0.052                             | -1.19     | -0.051                             | -1.08     |
| $ z  \in [1.96, 2.58)$<br>( $p \in [0.01, 0.05)$ )  | 0.256                      | 0.323                       | +0.088***                   | +3.96     | +0.060                             | +1.45     | +0.068*                            | +1.71     |
| $ z  \in [1.65, 1.96)$<br>( $p \in [0.05, 0.10)$ )  | 0.044                      | 0.035                       | -0.003                      | -0.31     | -0.003                             | -0.17     | -0.008                             | -0.51     |
| $ z  < 1.65$<br>( $p \geq 0.10$ )                   | 0.054                      | 0.038                       | -0.013                      | -0.36     | -0.019                             | -0.58     | -0.015                             | -0.45     |
| <i>Mean</i> $ z $                                   | 3.38                       | 3.48                        | -0.023                      | -0.19     | +0.281*                            | +1.76     | +0.145                             | +0.90     |

**Panel D: Paper’s most-important result**

|                                                     | Pre-Policy<br>( $N = 82$ ) | Post-Policy<br>( $N = 72$ ) | Difference                  |           | Difference-in-Differences          |           |                                    |           |
|-----------------------------------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|------------------------------------|-----------|
|                                                     |                            |                             | Within-JAR<br>( $N = 154$ ) |           | JAR vs. TAR + JAE<br>( $N = 752$ ) |           | JAR vs. All Top<br>( $N = 1,283$ ) |           |
|                                                     | Mean                       | Mean                        | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                              | $t$ -stat |
| $ z  \geq 3.29$<br>( $p < 0.001$ )                  | 0.439                      | 0.375                       | -0.089                      | -1.40     | -0.080                             | -1.46     | -0.095                             | -1.61     |
| $ z  \in [2.58, 3.29)$<br>( $p \in [0.001, 0.01)$ ) | 0.244                      | 0.222                       | -0.015                      | -0.30     | +0.014                             | +0.24     | +0.015                             | +0.26     |
| $ z  \in [1.96, 2.58)$<br>( $p \in [0.01, 0.05)$ )  | 0.232                      | 0.361                       | +0.142***                   | +2.77     | +0.100                             | +1.63     | +0.113*                            | +1.69     |
| $ z  \in [1.65, 1.96)$<br>( $p \in [0.05, 0.10)$ )  | 0.049                      | 0.014                       | -0.026                      | -0.89     | -0.014                             | -0.45     | -0.016                             | -0.59     |
| $ z  < 1.65$<br>( $p \geq 0.10$ )                   | 0.037                      | 0.028                       | -0.011                      | -0.30     | -0.020                             | -0.54     | -0.018                             | -0.50     |
| <i>Mean</i> $ z $                                   | 3.48                       | 3.46                        | -0.184                      | -0.63     | +0.083                             | +0.30     | -0.075                             | -0.27     |

*Notes.* Each row reports JAR’s mean before/after January 1, 2015. *Within-JAR* is the OLS  $\beta$  on  $\text{Post}_t$  in a JAR-only specification with topic FE and submission-year-clustered SEs; *Difference-in-Differences* columns report  $\beta_3$  on  $\text{JAR} \times \text{Post}$  from a two-way-FE specification (topic + submission-year FE; SEs clustered at journal  $\times$  submission year). The DiD sample is empirical (archival + experimental) papers with policy assign-date in January 2012–December 2017, excluding those whose  $p_{10}/p_{90}$ -lag bounds straddle the cutoff. The sample is limited to papers whose *most-important* (headline) result carries a  $p$ -value — reported,  $|t|$ -derived, or imputed from significance stars — which drops only papers whose headline result is non-numeric (e.g. a machine-learning AUC, a structural confidence region, or a magnitude read from text). Panel A reports the indicators across *all* extracted results (the *comprehensive* extraction: every coefficient of every column of every table/panel). Panel B restricts to regression rows the extraction LLM flagged as discussed in the paper’s introduction (the results an author chose to highlight; roughly six per paper); Panel C restricts to the paper’s *primary* hypothesis-test coefficients (roughly two per paper); Panel D reports the same indicators for the paper’s *most-important* (headline) result. Results whose  $p$ -value is reported only via significance stars (roughly 4% of extracted rows) are dropped because the imputation rule places stars-only rows into specific  $|z|$  bands mechanically. Controls  $X = \{\text{PAPERLENGTH}, \text{MEANHINDEX}, \text{MEANYEARSCAREERSCOPUS}, \text{NUMAUTHORS}, \text{NUMTABLES}, \text{TOP20UTD}\}$ .  $\text{PAPERPRIMARYMEANP}$  is omitted because the outcome is itself the  $p$ -value distribution. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (two-sided tests).

**Table OA7:** Does the *JAR* mandate response scale with project opacity?

| Outcome                                              | Triple interaction $JAR \times Post \times Opacity$ |                                                    |
|------------------------------------------------------|-----------------------------------------------------|----------------------------------------------------|
|                                                      | Opacity index (z)                                   | Proprietary-private (z)                            |
| Abstract readability (Fog)                           | -0.425*<br>(0.253) [t=-1.68]<br>Rank 2/6, N=3,480   | -0.516**<br>(0.217) [t=-2.37]<br>Rank 3/6, N=3,480 |
| Main-text readability (Fog)                          | -0.035<br>(0.150) [t=-0.24]<br>Rank 6/6, N=3,480    | 0.003<br>(0.138) [t=0.02]<br>Rank 5/6, N=3,480     |
| Introduction hedging                                 | -0.021**<br>(0.009) [t=-2.43]<br>Rank 1/6, N=3,480  | -0.009<br>(0.008) [t=-1.09]<br>Rank 3/6, N=3,480   |
| Reporting discipline (avg stars)                     | -0.063<br>(0.074) [t=-0.86]<br>Rank 4/6, N=3,363    | 0.056<br>(0.063) [t=0.90]<br>Rank 3/6, N=3,363     |
| Novelty (intro-doc kNN-20)                           | 0.004<br>(0.003) [t=1.24]<br>Rank 1/6, N=3,480      | 0.005*<br>(0.003) [t=1.88]<br>Rank 2/6, N=3,480    |
| Topic FE, Pub-year FE                                | Yes                                                 | Yes                                                |
| Controls (length, h-index, career, #authors, top-20) | Yes                                                 | Yes                                                |
| SE clustered journal $\times$ year                   | Yes                                                 | Yes                                                |

*Notes.* Accounting-only empirical sample ( $JAR + TAR, JAE, CAR, RAST, MS$ ),  $\pm 7$ -year window. Each cell is the coefficient on the triple interaction  $JAR \times Post \times Opacity$  from a DiD that includes all lower-order interactions, topic and submission-year fixed effects, and the paper-level controls. Opacity is standardized (per 1 SD). The additive index combines data ownership (private>mixed>public), access (restricted>mixed>public), a proprietary-private flag, methods-section length, an experimental bump, and Compustat-like relief. Expected signs: readability and reporting-discipline coefficients negative (Fog/stars fall more where opacity is higher), hedging negative, novelty positive. “Rank” is *JAR*’s placebo position when each of the six accounting journals is relabeled as treated in turn (1/6 means first); it is a placebo ranking, not a significance level. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (analytic, clustered).

**Table OA8:** Writing effects by the strength of the headline result

| Outcome                                              | JAR $\times$ Post split by headline strength |                               |                                                     |
|------------------------------------------------------|----------------------------------------------|-------------------------------|-----------------------------------------------------|
|                                                      | Strong papers                                | $\times$ Weak (triple)        | Weak papers (net)                                   |
| Introduction hedging                                 | -0.011<br>(0.023) [t=-0.48]                  | -0.042*<br>(0.022) [t=-1.92]  | -0.053***<br>(0.019) [t=-2.78]<br>Rank 1/6, N=1,010 |
| Contribution hedging (CARS)                          | -0.029<br>(0.039) [t=-0.76]                  | -0.156**<br>(0.076) [t=-2.05] | -0.185***<br>(0.049) [t=-3.78]<br>Rank 1/6, N=982   |
| Gap hedging (CARS)                                   | -0.152***<br>(0.036) [t=-4.24]               | 0.087*<br>(0.049) [t=1.76]    | -0.065<br>(0.056) [t=-1.17]<br>Rank 2/6, N=847      |
| Abstract readability (Fog)                           | -0.80<br>(1.16) [t=-0.69]                    | -0.55<br>(1.04) [t=-0.53]     | -1.34*<br>(0.73) [t=-1.83]<br>Rank 2/6, N=1,010     |
| Introduction readability (Fog)                       | -0.67**<br>(0.34) [t=-1.99]                  | -0.01<br>(0.67) [t=-0.02]     | -0.69<br>(0.64) [t=-1.08]<br>Rank 2/6, N=1,010      |
| Topic FE, Submission-year FE                         | Yes                                          | Yes                           | Yes                                                 |
| Controls (length, h-index, career, #authors, top-20) | Yes                                          | Yes                           | Yes                                                 |
| SE clustered journal $\times$ year                   | Yes                                          | Yes                           | Yes                                                 |

*Notes.* Accounting-only sample, submission years 2012–2017. “Weak” papers have a headline test statistic (the paper’s most-important result,  $|t|$  or  $|z|$ ) below the sample median (3.12); the split leaves roughly 32 treated post-mandate papers on the weak side. Each row is one regression of the writing outcome on JAR  $\times$  Post  $\times$  Weak with all lower-order terms, topic and submission-year fixed effects, and the paper-level controls. “Strong papers” is the JAR  $\times$  Post coefficient; “Weak papers (net)” adds the triple interaction (delta-method standard error). Hedging outcomes are hedge rates (higher = more hedged); readability is the Gunning Fog index (higher = more complex). The signed prediction is negative throughout; “Rank” is the one-sided placebo position of the net weak-paper effect when each of the six accounting journals is relabeled as treated in turn (1/6 means JAR ranks first). With six journals and no random assignment this is a placebo ranking, not a significance level. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (analytic, clustered).

**Table OA9:** Timing-imputation robustness: received-date-only re-estimation

| Outcome                         | Full sample |       |       | Received-date only |       |     |
|---------------------------------|-------------|-------|-------|--------------------|-------|-----|
|                                 | Coef.       | $t$   | $N$   | Coef.              | $t$   | $N$ |
| Abstract readability (Fog)      | -1.250**    | -2.03 | 1,356 | -1.398**           | -2.48 | 959 |
| Intro-document novelty (kNN-20) | 0.013**     | 2.53  | 1,356 | 0.014***           | 2.84  | 959 |
| Introduction hedging rate       | -0.028**    | -2.41 | 1,356 | -0.025**           | -2.11 | 959 |
| Log citations                   | -0.277*     | -1.67 | 1,356 | -0.202             | -1.28 | 959 |

*Cross-journal placebo relabelings (received-only, each control journal relabeled as treated):*

Abstract readability (Fog)  $\hat{\beta} = -1.398$ ; exceeded in magnitude by 2 of 5 placebo relabelings  
Intro-document novelty (kNN-20)  $\hat{\beta} = 0.014$ ; exceeded in magnitude by 2 of 5 placebo relabelings

*Notes.* DiD estimates of JAR $\times$ Post from  $y \sim \text{JAR} + \text{JAR} \times \text{Post} + \text{Post} + C(\text{pub. year}) + C(\text{topic}) + \text{controls}$ , empirical (archival+experimental) papers, assign-date Jan 2012–Dec 2017, all-accounting control set (TAR, JAE, CAR, RAST, MS). Controls: paper length, mean author  $h$ -index, mean career years, number of authors, top-20 UTD. Standard errors clustered by journal $\times$ year. “Full sample” uses received plus imputed-certain submission dates; “Received-date only” keeps papers with an observed submission date. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ .




## C Within-editor stability at *JAR* (Design D)

This appendix supports the within-editor test in Section 5.4.3. Table OA10 reports the post-versus-pre-2015 change on each outcome estimated within the three *JAR* editors who span the mandate—editor by editor and pooled with editor fixed effects—alongside the whole-journal benchmark and a falsification at the journal’s actual 2019 masthead turnover. Figure OA1 plots the same coefficients. Editor assignments and the received-to-accepted duration measure are constructed for published papers only, so every estimate here conditions on acceptance and speaks to the composition of what a given editor published, not to submissions or rejections; the design accordingly rules out that the 2015 shifts are an artifact of masthead turnover, but does not separate demand from supply.

**Table OA10:** The same editors show the same shifts: within-editor before/after changes at *JAR*

| Outcome                     | Post – pre (2015), editor fixed effects |                      |                   |                     |                     | Falsification     |
|-----------------------------|-----------------------------------------|----------------------|-------------------|---------------------|---------------------|-------------------|
|                             | JAR-wide                                | Pooled (3 ed.)       | Berger            | Leuz                | Skinner             | Break at 2019     |
| Abstract readability (Fog)  | -1.166***<br>(0.295)                    | -1.125***<br>(0.343) | -0.723<br>(0.571) | -1.190**<br>(0.555) | -1.532**<br>(0.728) | -0.350<br>(0.284) |
| Introduction hedging rate   | -0.016*<br>(0.009)                      | -0.017**<br>(0.008)  | -0.021<br>(0.021) | -0.029*<br>(0.016)  | +0.002<br>(0.023)   | -0.013<br>(0.008) |
| Intro-doc novelty (kNN-20)  | +0.008**<br>(0.003)                     | +0.007**<br>(0.003)  | +0.010<br>(0.008) | +0.003<br>(0.007)   | +0.006<br>(0.007)   | +0.002<br>(0.007) |
| Log review duration         | +0.335***<br>(0.073)                    | +0.198***<br>(0.051) | +0.113<br>(0.166) | +0.250**<br>(0.117) | +0.228**<br>(0.116) | —                 |
| Data opacity (prop. either) | +0.112**<br>(0.045)                     | +0.067<br>(0.068)    | +0.012<br>(0.096) | -0.033<br>(0.102)   | +0.245**<br>(0.095) | +0.019<br>(0.062) |
| Editor FE                   | Yes                                     | Yes                  | —                 | —                   | —                   | Yes               |
| SE clustered assign-year    | Yes                                     | Yes                  | Yes               | Yes                 | Yes                 | Yes               |

*Notes.* Design D. The sample is the three *JAR* editors who handled at least five papers on each side of the January-2015 mandate—Berger (45/43), Leuz (34/55), and Skinner (35/39); Sapra is excluded (pre-mandate  $n = 5$ ). Each cell is the coefficient on a post-2015 indicator from a regression of the outcome on that indicator, estimated within the pooled spanning subsample with editor fixed effects (columns 2–3) or editor by editor (columns 4–6), with standard errors clustered on the assignment year; the *JAR*-wide column is the whole-journal benchmark. Adding topic fixed effects to the pooled specification leaves the estimates essentially unchanged (readability  $-1.172$ , hedging  $-0.016$ , novelty  $+0.009$ , log duration  $+0.203$ ). The final column is a falsification at the journal’s *actual* masthead turnover: among post-2015 *JAR* papers, the break at the 2019 editor transition (Hail, Wittenberg-Moerman, Verdi) on each outcome; log duration is untestable there because the received-by-2017 censoring rule leaves no computable post-2019 durations. Data opacity does not move within editor except for Skinner, whose desk’s auditing share rises from 6% to 18% over the window—a composition shift, not a within-editor change (his opacity coefficient is more than twice the journal-wide change and survives topic fixed effects on his desk alone). This design rules out that *JAR*’s 2015 shifts are an artifact of who sits in the editor’s chair; it does not separate the demand margin from supply, because the same editor accepts from a changed submission pool. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (analytic, clustered).



**Figure OA1: The same editors show the same shifts.** Post-versus-pre-2015 change at *JAR* on each outcome, estimated editor by editor for the three editors who span the mandate (Berger, Leuz, Skinner; navy) and for the whole journal (grey), with 95% confidence intervals. Points left of the dashed line mark declines (more readable, less hedged); points to the right mark increases (more novel, slower reviews, more opaque data). Data opacity moves only for Skinner, a desk-composition artifact discussed in Section 5.4.3.

## D Data construction and reproducibility

All estimates are computed from an aggregate corpus of accounting (and, for the general-equilibrium tests, finance) research articles. Text measures (section-level Gunning–Fog, hedging rates, and embedding-based novelty) and the extracted reported-results data are constructed at the paper level and merged to bibliographic and citation metadata. The difference-in-differences frame is restricted to archival and experimental research articles with an abstract and a classifiable introduction, with treatment timing assigned at the submission date. The design-based inference (the restricted wild cluster bootstrap, in-time placebos, cross-journal placebo relabelings), the synthetic-control and synthetic difference-in-differences estimators, the Oster bounds, and the finance-benchmark general-equilibrium tests, the second-treated-journal (*Management Science*, 2019) and continuous policy-index dose-response designs, and the proprietary-data and weak-result heterogeneity are implemented in Python. The replication package (`code/`, documented in its README) is organized as a numbered pipeline: a shared data loader (`_engine.py`) followed by scripts 01–09 that build each channel’s estimates, the design-based inference, the extensions, and every figure and table in the paper, with an `extensions/` folder holding the exploratory analyses and their result files. Every figure and table here is reproduced by that pipeline from the bundled aggregate data. The full package—code, the bundled aggregate data, and the README—will be posted publicly upon publication; until then it is available from the authors on request.

## E Design-based inference

Before the inference battery, Table OA11 collects in one place the design each headline result’s interpretation rests on.

Because treatment is assigned to a single journal, we confirm the main estimates with inference that does not rely on large-cluster asymptotics. Table OA12 reports, for one primary outcome per channel, the  $JAR \times Post$  estimate together with the analytic clustered  $p$ -value, two design-based alternatives (the restricted wild cluster bootstrap and an in-time placebo), and two placebo rankings: a cross-journal rank obtained by relabeling each control journal in turn as treated, and a synthetic-control rank that orders  $JAR$  and its donor journals by the ratio of post- to pre-policy fit. Because the policy was not randomly assigned and the number of journals is small, we report both exercises as ranks rather than as  $p$ -values; on the cross-journal rank  $JAR$  is first for every outcome, which is the most that exercise can show. The novelty effect is significant analytically ( $p = 0.01$ ) and ranks first among the synthetic-control placebos; the wild bootstrap (0.16) and in-time placebo (0.38) are weaker,

**Table OA11:** Identification designs by outcome

| Result                                               |       | Headline estimate                                                      | esti-        | Margin                 | Primary supporting design                                                                                                         | Where          |
|------------------------------------------------------|-------|------------------------------------------------------------------------|--------------|------------------------|-----------------------------------------------------------------------------------------------------------------------------------|----------------|
| Readability improves                                 | im-   | Abstract −1.25                                                         | Fog          | Behavioral             | Second treated journal ( <i>MS</i> 2019) within author; exempt-theory placebo; opacity and evidence-strength gradients            | §5.1.1, §5.2   |
| Hedging decline concentrates in weak-evidence papers |       | Contribution hedging −0.185 (weak papers)                              |              | Behavioral (mechanism) | Triple interaction with all lower-order terms; <i>JAR</i> first among all journal relabelings                                     | §5.1.4         |
| More supporting coefficients; fewer headline claims  |       | Multiplicity ratio +14.9; $\approx 15$ more coefficients within author |              | Behavioral             | The one channel that survives explicit author fixed effects; flat test-statistic distribution                                     | §5.1.3, §5.2.1 |
| Novelty rises                                        |       | Intro novelty +0.013 ( $\approx \frac{1}{3}$ SD)                       |              | Market (selection)     | Synthetic control; <i>JAR</i> first among all journal relabelings; no-text reference-network corroboration; stable core topic mix | §5.3.1         |
| Rare-topic rises                                     | share | +12.9pp                                                                |              | Market (selection)     | Step at 2015; <i>JAR</i> first among all journal relabelings; threshold, leave-one-year-out, and pre-only-definition robustness   | §5.3.1         |
| Citation counts fall; citing-journal status rises    |       | Count top-6 +2.5pp                                                     | −0.28; share | Market (reception)     | Tier shares and citing-SNIP robust; prestige gain concentrated in top accounting/business citers                                  | §5.3.3         |


*Notes.* Treatment is a single journal, so analytic clustered  $t$ -statistics are descriptive throughout and no single design supports every result; this table summarizes the design each result’s interpretation rests on. Where a result is described as *JAR* first among all journal relabelings, each control journal was relabeled as treated in turn; with six journals and no random assignment this is a placebo ranking, not a significance level. Full design-based inference for every estimate is collected in Appendix E.

reflecting the limited power of a single treated cluster, so we corroborate novelty with the convergent-validity evidence in Appendix H. The writing and citation effects keep their sign and magnitude but, in a single small journal, are estimated less precisely under the most conservative procedures. One caveat cuts the other way: with a single treated cluster the wild cluster bootstrap is known to under-reject (MacKinnon and Webb, 2017), so its large  $p$ -values understate rather than overturn the evidence. Figure OA2 shows the corresponding synthetic-control gaps: novelty is the series in which  $JAR$  separates from the placebo donor cloud after the policy, while the other outcomes lie within the placebo spread.

**Table OA12:** Design-based inference for one treated cluster

| Outcome                | Coef.  | Analytic<br>cluster | Wild<br>boot. | In-time<br>placebo | Placebo<br>rank | Synth.<br>ctrl rank |
|------------------------|--------|---------------------|---------------|--------------------|-----------------|---------------------|
| Novelty (intro-doc)    | +0.013 | 0.01                | 0.16          | 0.38               | 1/6             | 1/6                 |
| Abstract readability   | −1.250 | 0.04                | 0.30          | 0.50               | 1/6             | 8/9                 |
| Introduction hedging   | −0.028 | 0.02                | 0.43          | 0.31               | 1/6             | 6/9                 |
| Log citations          | −0.277 | 0.09                | 0.55          | 0.69               | 1/6             | 2/9                 |
| Cite share top-6 acct. | +0.028 | 0.00                | 0.09          | 0.19               | 1/6             | 4/8                 |

*Notes.* Each row reports the  $JAR \times Post$  point estimate, the analytic  $p$ -value, two design-based  $p$ -values, and two placebo ranks. *Analytic cluster:* cluster-robust on journal $\times$ year (the value behind Table 2). *Wild boot.:* restricted wild cluster bootstrap- $t$  on the journal (1,999 replications); conservative with one treated cluster (MacKinnon and Webb, 2017). *In-time placebo:* placebo policy years among controls. *Placebo rank:* each control journal relabeled treated in turn, ranked one-sided in the direction predicted by P1–P5;  $JAR$  ranks first of six for every outcome. With six journals and no random assignment this is a placebo ranking, not a significance level. *Synthetic control rank:*  $JAR$ 's position among itself and its donor journals when all are ordered by the ratio of post- to pre-policy RMSPE; the donor pool is the five accounting controls plus the finance journals with enough observations for the outcome, so the denominator varies by outcome. This too is a ranking, not a significance level. With a single treated cluster these procedures are deliberately conservative and no outcome clears every one: the writing and novelty effects clear the analytic cluster test, and novelty is the only outcome that ranks first among the synthetic-control placebos. Novelty's remaining soft spot is the in-time placebo (0.38); we therefore corroborate it with the convergent-validity evidence in Appendix H, where reference-network measures that use no text move the same way.



**Figure OA2: Synthetic-control gaps.** *JAR* minus synthetic *JAR* (solid) against placebo donor journals (grey), for covariate-adjusted journal×year outcomes. The vertical line marks the policy. Each panel reports *JAR*'s rank among itself and the donor journals when all are ordered by the ratio of post- to pre-policy RMSPE.

## F Supporting tables for the writing results

This appendix collects the coefficients and  $t$ -statistics that underlie the writing figures in Section 5.1.1. Table OA13 reports the section-level difference-in-differences for hedging and readability that Figure 3 plots, with a descriptive panel for the pooled sample. Table OA14 reports the same estimates within the introduction, move by move, formalizing the rhetorical-move asymmetry in the right panel of Figure 3: hedging falls in the gap move ( $-0.087$ ,  $t = -2.42$ ) and the contribution move ( $-0.091$ ,  $t = -3.46$ ) and is flat in the descriptive moves (motivation, results preview, present research). Both tables use the same specification and sample as the main estimates.

**Table OA13:** Paper-level sections: descriptives and difference-in-differences around the *JAR* January 2015 data- and code-sharing policy

**Panel A: Descriptive statistics by section**

|                                  | N     | Mean | SD   | Q1   | Median | Q3   |
|----------------------------------|-------|------|------|------|--------|------|
| <b><i>Whole paper</i></b>        |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,630 | 425  | 88   | 366  | 418    | 474  |
| <i>Hedging</i>                   | 4,630 | 0.35 | 0.06 | 0.30 | 0.34   | 0.39 |
| <i>Readability</i>               | 4,630 | 16.7 | 1.5  | 15.7 | 16.7   | 17.7 |
| <b><i>Abstract</i></b>           |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,630 | 6    | 2    | 5    | 6      | 7    |
| <i>Hedging</i>                   | 4,630 | 0.63 | 0.22 | 0.50 | 0.62   | 0.80 |
| <i>Readability</i>               | 4,630 | 19.1 | 3.0  | 17.1 | 19.0   | 21.0 |
| <b><i>Introduction</i></b>       |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,629 | 65   | 18   | 52   | 64     | 76   |
| <i>Hedging</i>                   | 4,629 | 0.50 | 0.10 | 0.44 | 0.51   | 0.57 |
| <i>Readability</i>               | 4,629 | 17.7 | 1.8  | 16.5 | 17.6   | 18.9 |
| <b><i>Theory/predictions</i></b> |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,282 | 69   | 31   | 47   | 64     | 85   |
| <i>Hedging</i>                   | 4,282 | 0.46 | 0.14 | 0.37 | 0.47   | 0.56 |
| <i>Readability</i>               | 4,282 | 17.3 | 1.9  | 16.1 | 17.3   | 18.5 |
| <b><i>Methodology/data</i></b>   |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,623 | 74   | 39   | 46   | 69     | 95   |
| <i>Hedging</i>                   | 4,623 | 0.18 | 0.08 | 0.12 | 0.18   | 0.24 |
| <i>Readability</i>               | 4,623 | 15.8 | 1.9  | 14.6 | 15.8   | 17.0 |
| <b><i>Results</i></b>            |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,629 | 195  | 78   | 139  | 187    | 243  |
| <i>Hedging</i>                   | 4,629 | 0.28 | 0.08 | 0.23 | 0.28   | 0.33 |
| <i>Readability</i>               | 4,629 | 16.4 | 1.7  | 15.3 | 16.4   | 17.5 |
| <b><i>Conclusions</i></b>        |       |      |      |      |        |      |
| <i>Length (sents)</i>            | 4,627 | 21   | 16   | 12   | 17     | 25   |
| <i>Hedging</i>                   | 4,627 | 0.62 | 0.17 | 0.50 | 0.62   | 0.74 |
| <i>Readability</i>               | 4,627 | 18.1 | 2.3  | 16.7 | 18.1   | 19.6 |



**Table OA13 (continued).** Difference-in-differences by section: sentence hedging and Gunning Fog readability as sub-rows under each section banner.

**Panel B: Hedging and Readability DiD by section**

|                                  | Pre-Policy<br>( $N = 84$ ) | Post-Policy<br>( $N = 78$ ) | Difference                  |           | Difference-in-Differences          |           |                                    |           |
|----------------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|------------------------------------|-----------|
|                                  |                            |                             | Within-JAR<br>( $N = 162$ ) |           | JAR vs. TAR + JAE<br>( $N = 788$ ) |           | JAR vs. All Top<br>( $N = 1,356$ ) |           |
|                                  |                            |                             | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                              | $t$ -stat |
| <hr/>                            |                            |                             |                             |           |                                    |           |                                    |           |
| <b><i>Whole paper</i></b>        |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.366                      | 0.362                       | -0.012-1.43                 |           | -0.013                             | -1.64     | -0.009                             | -1.09     |
| <i>Readability</i>               | 17.33                      | 16.87                       | -0.570** -2.56              |           | -0.768**                           | -2.35     | -0.533*                            | -1.82     |
| <b><i>Abstract</i></b>           |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.664                      | 0.610                       | -0.068* -1.73               |           | -0.028                             | -0.73     | -0.032                             | -0.93     |
| <i>Readability</i>               | 20.52                      | 19.36                       | -1.221** -2.32              |           | -1.352***                          | -2.60     | -1.250**                           | -2.03     |
| <b><i>Introduction</i></b>       |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.523                      | 0.503                       | -0.026* -1.81               |           | -0.027**                           | -2.02     | -0.028**                           | -2.41     |
| <i>Readability</i>               | 18.74                      | 17.98                       | -0.835** -2.29              |           | -1.079***                          | -2.74     | -0.959**                           | -2.52     |
| <b><i>Theory/predictions</i></b> |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.448                      | 0.417                       | -0.038* -1.88               |           | -0.018                             | -0.85     | -0.024                             | -1.07     |
| <i>Readability</i>               | 18.22                      | 17.40                       | -0.950** -2.51              |           | -1.017***                          | -2.90     | -0.669*                            | -1.96     |
| <b><i>Methodology/data</i></b>   |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.194                      | 0.186                       | -0.016-1.00                 |           | -0.001                             | -0.07     | -0.004                             | -0.29     |
| <i>Readability</i>               | 16.26                      | 16.06                       | -0.359-0.91                 |           | -0.539                             | -1.18     | -0.311                             | -0.73     |
| <b><i>Results</i></b>            |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.320                      | 0.323                       | -0.009-0.73                 |           | -0.012                             | -1.42     | -0.008                             | -0.84     |
| <i>Readability</i>               | 16.96                      | 16.73                       | -0.280-1.25                 |           | -0.547                             | -1.62     | -0.313                             | -1.05     |
| <b><i>Conclusions</i></b>        |                            |                             |                             |           |                                    |           |                                    |           |
| <i>Hedging</i>                   | 0.664                      | 0.605                       | -0.070*** -5.12             |           | -0.048**                           | -2.18     | -0.042**                           | -2.48     |
| <i>Readability</i>               | 18.53                      | 17.78                       | -1.038*** -4.45             |           | -1.006***                          | -2.83     | -0.823**                           | -2.31     |

*Notes.* Panel A pools all journals and years in the empirical sample of archival and experimental research articles ( $N = 4,630$ ; see Table 1). *Length* is the section sentence count (or the whole paper for *Whole paper*). *Hedging* is the fraction of section sentences flagged as uncertain by the DistilRoBERTa sentence classifier (binary label at  $t = 0.5$ ). *Readability* is Gunning Fog index. Each section row is conditional on the section being non-empty. Panel B reports difference-in-differences for the same sections, with two sub-rows per section: *Hedging* (sentence-level DistilRoBERTa rate at  $t = 0.5$ ) and *Readability* (Gunning Fog). Each row reports JAR's mean before/after January 1, 2015. *Within-JAR* is the OLS  $\beta$  on  $\text{Post}_t$  in a JAR-only specification with topic FE and submission-year-clustered SEs; *Difference-in-Differences* columns report  $\beta_3$  on  $\text{JAR} \times \text{Post}$  from a two-way-FE specification (topic + submission-year FE; SEs clustered at journal  $\times$  submission year). The DiD sample is empirical (archival + experimental) papers with policy assign-date in January 2012–December 2017, excluding those whose  $p_{10}/p_{90}$ -lag bounds straddle the cutoff. Controls  $X = \{\text{PAPERLENGTH}, \text{MEANHINDEX}, \text{MEANYEARSCAREERSCOPUS}, \text{NUMAUTHORS}, \text{TOP20UTD}\}$ . \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (two-sided tests).

**Table OA14:** Introduction paragraphs by CARS discourse-move category: descriptives and difference-in-differences around the *JAR* January 2015 data- and code-sharing policy

**Panel A: Descriptive statistics by discourse category**

|                                | N     | Mean | SD   | Q1   | Median | Q3   |
|--------------------------------|-------|------|------|------|--------|------|
| <b><i>Motivation</i></b>       |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 4,504 | 2.4  | 1.4  | 1.0  | 2.0    | 3.0  |
| <i>Hedging</i>                 | 4,504 | 0.46 | 0.22 | 0.33 | 0.46   | 0.60 |
| <i>Readability</i>             | 4,504 | 19.0 | 3.9  | 16.6 | 18.6   | 20.8 |
| <b><i>Gap</i></b>              |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 3,891 | 1.8  | 1.0  | 1.0  | 2.0    | 2.0  |
| <i>Hedging</i>                 | 3,891 | 0.62 | 0.21 | 0.50 | 0.62   | 0.75 |
| <i>Readability</i>             | 3,891 | 18.8 | 3.8  | 16.4 | 18.5   | 20.6 |
| <b><i>Present research</i></b> |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 4,612 | 3.2  | 1.7  | 2.0  | 3.0    | 4.0  |
| <i>Hedging</i>                 | 4,612 | 0.44 | 0.18 | 0.31 | 0.43   | 0.56 |
| <i>Readability</i>             | 4,612 | 19.1 | 3.1  | 17.2 | 18.9   | 20.8 |
| <b><i>Results preview</i></b>  |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 4,586 | 3.7  | 2.0  | 2.0  | 3.0    | 5.0  |
| <i>Hedging</i>                 | 4,586 | 0.59 | 0.16 | 0.50 | 0.60   | 0.70 |
| <i>Readability</i>             | 4,586 | 19.2 | 2.7  | 17.4 | 19.0   | 20.8 |
| <b><i>Contribution</i></b>     |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 4,488 | 2.6  | 1.3  | 2.0  | 3.0    | 3.0  |
| <i>Hedging</i>                 | 4,488 | 0.61 | 0.18 | 0.50 | 0.62   | 0.73 |
| <i>Readability</i>             | 4,488 | 18.6 | 3.1  | 16.6 | 18.4   | 20.1 |
| <b><i>Other</i></b>            |       |      |      |      |        |      |
| <i>Length (paragraphs)</i>     | 3,710 | 2.1  | 1.5  | 1.0  | 2.0    | 3.0  |
| <i>Hedging</i>                 | 3,670 | 0.17 | 0.24 | 0.00 | 0.00   | 0.29 |
| <i>Readability</i>             | 3,710 | 12.9 | 4.2  | 10.3 | 12.6   | 15.2 |

**Table OA14 (continued).** Difference-in-differences by CARS discourse category: sentence hedging and Gunning Fog readability as sub-rows under each category banner.

**Panel B: Hedging and Readability DiD by discourse category**


|                         | Pre-Policy<br>( $N = 82$ ) | Post-Policy<br>( $N = 76$ ) | Difference                  |           | Difference-in-Differences          |           |                                     |           |
|-------------------------|----------------------------|-----------------------------|-----------------------------|-----------|------------------------------------|-----------|-------------------------------------|-----------|
|                         |                            |                             | Within-JAR<br>( $N = 158$ ) |           | JAR vs. TAR + JAE<br>( $N = 765$ ) |           | JAR vs. All Top<br>( $N = 1, 313$ ) |           |
|                         |                            |                             | Coef.                       | $t$ -stat | Coef.                              | $t$ -stat | Coef.                               | $t$ -stat |
| <b>Motivation</b>       |                            |                             |                             |           |                                    |           |                                     |           |
| <i>Hedging</i>          | 0.502                      | 0.462                       | -0.043                      | -1.09     | -0.016                             | -0.61     | -0.005                              | -0.22     |
| <i>Readability</i>      | 20.16                      | 18.64                       | -1.645**                    | -2.50     | -2.046***                          | -2.64     | -1.135                              | -1.30     |
| <b>Gap</b>              |                            |                             |                             |           |                                    |           |                                     |           |
| <i>Hedging</i>          | 0.673                      | 0.582                       | -0.096***                   | -4.38     | -0.076*                            | -1.89     | -0.087**                            | -2.42     |
| <i>Readability</i>      | 20.53                      | 19.40                       | -1.327                      | -1.56     | -0.946                             | -1.19     | -0.491                              | -0.52     |
| <b>Present research</b> |                            |                             |                             |           |                                    |           |                                     |           |
| <i>Hedging</i>          | 0.439                      | 0.397                       | -0.054*                     | -1.74     | -0.022                             | -0.77     | -0.017                              | -0.59     |
| <i>Readability</i>      | 19.84                      | 19.09                       | -0.941                      | -1.40     | -1.264                             | -1.64     | -0.788                              | -1.06     |
| <b>Results preview</b>  |                            |                             |                             |           |                                    |           |                                     |           |
| <i>Hedging</i>          | 0.596                      | 0.575                       | -0.025                      | -0.92     | -0.017                             | -0.71     | -0.023                              | -0.98     |
| <i>Readability</i>      | 20.05                      | 18.87                       | -1.197**                    | -2.37     | -1.498***                          | -3.22     | -1.330***                           | -3.11     |
| <b>Contribution</b>     |                            |                             |                             |           |                                    |           |                                     |           |
| <i>Hedging</i>          | 0.664                      | 0.577                       | -0.085***                   | -2.71     | -0.094***                          | -3.76     | -0.091***                           | -3.46     |
| <i>Readability</i>      | 19.71                      | 18.72                       | -1.043**                    | -2.29     | -1.133***                          | -2.65     | -0.947**                            | -2.00     |

*Notes.* Introduction paragraphs are classified into one of six CARS discourse-move categories (Motivation, Gap, Present research, Results preview, Contribution, Other) using a CARS-style prompt with a Qwen LLM. Both panels use the canonical primary-or-secondary CARS convention (a paragraph counts toward a category if either its primary or secondary label matches). Panel A reports paragraph-count per paper, fraction of sentences flagged as uncertain (DistilRoBERTa,  $t = 0.5$ ), and Gunning Fog readability for the concatenated text of the category. Panel B reports difference-in-differences for the same five primary categories, with two sub-rows per category: *Hedging* (sentence-level DistilRoBERTa rate at  $t = 0.5$ ) and *Readability* (Gunning Fog). Each row reports JAR's mean before/after January 1, 2015. *Within-JAR* is the OLS  $\beta$  on  $\text{Post}_t$  in a JAR-only specification with topic FE and submission-year-clustered SEs; *Difference-in-Differences* columns report  $\beta_3$  on  $\text{JAR} \times \text{Post}$  from a two-way-FE specification (topic + submission-year FE; SEs clustered at journal  $\times$  submission year). The DiD sample is empirical (archival + experimental) papers with policy assign-date in January 2012–December 2017, excluding those whose  $p_{10}/p_{90}$ -lag bounds straddle the cutoff. Controls  $X = \{\text{PAPERLENGTH}, \text{MEANHINDEX}, \text{MEANYEARS CAREER SCOPUS}, \text{NUMAUTHORS}, \text{TOP20UTD}\}$ . \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$  (two-sided tests).

## G Threshold bunching: showcased versus incidental results

Section 5.1.2 reports that we find no robust evidence the mandate reduced threshold-driven inference. This appendix shows why. We split every extracted coefficient into the results authors showcase—the primary hypothesis tests, the results discussed in the introduction, and the single headline result—and the incidental coefficients that fill out the tables, and we estimate the difference-in-differences on the caliper share just above and just below each

conventional cutoff. Figure OA3 collects the picture. The just-above-one-percent share falls, and it falls most in the showcased results ( $-0.077$  for the primary results,  $-0.015$  for the incidental ones), which on its own looks like deterred specification search. But the displaced mass does not move into the non-significant region. It reappears just above the five-percent cutoff, where the showcased share rises by roughly 0.10, while the share of insignificant showcased results does not rise and the overall share of significant showcased results, if anything, increases. With a single treated journal the one-percent decline is also fragile: relabeling each control journal as treated places *JAR* first among the six journals, a ranking that is suggestive but cannot establish the effect on its own. The pattern is that of reported statistics sorting one threshold down, not of less specification search, so we treat the test-statistic channel as an unresolved null.



**Figure OA3: Just-above-threshold bunching: showcased versus incidental results.** *JAR*  $\times$  *Post* difference-in-differences estimates on the share of a paper's reported results in the just-above caliper of each conventional threshold, split by whether the result is one the author showcases (primary, introduction-discussed, headline) or an incidental coefficient (the full universe of extracted results). The one-percent just-above share falls most in showcased results, but the mass relocates just above the five-percent threshold rather than into the non-significant region. See Section 5.1.2 and the discussion in this appendix.

## H Novelty: robustness and convergent validity

The novelty result in Section 5.3.1 is the paper’s most inference-robust finding, so we document that it does not depend on a single measure or a single embedding choice. Table OA15 collects the evidence. Panel A re-estimates the  $JAR \times Post$  effect across nearest-neighbor distances at  $k \in \{1, 5, 10, 20, 50\}$ , a five-year reference window, an abstract-only variant, an “originality” variant, and a distance-to-field-centroid variant. The effect is stable between one-third and one-half of a standard deviation everywhere except the centroid measure, which is weak—the signature of a *local* shift, in which papers move away from their nearest predecessors while the field’s center holds still. Panel B turns to measures built from the reference list rather than the text: post-mandate *JAR* papers cite a wider range of fields (higher entropy, a larger non-accounting share, more distinct journals), and because these use no text they are an independent check on the embedding measure. Panel C confirms that the text and citation measures correlate positively at the paper level, while the rhetorical novelty-claim indicator is nearly orthogonal to all of them and is therefore reported as a distinct construct. The novelty we measure is one of breadth and textual distance rather than of unusual recombination.


Two descriptive figures make concrete what kind of novelty the mandate sorts in. Figure OA4 plots the density of introduction novelty in *JAR* papers before and after the policy: the distribution does not shift uniformly but grows a thicker upper tail, with more mass past the pre-mandate ninetieth percentile. Figure OA5 traces the mechanism on the citation side, showing that *JAR* papers come to draw a markedly larger share of their references from outside accounting after the mandate while the control journals hold flat.

Finally, two figures document the topic-composition evidence behind Section 5.3.1. Figure OA6 plots each coarse topic’s share of *JAR* papers before and after the mandate against the same movement in the controls: the shares track, and no per-topic difference-in-differences survives a Benjamini–Hochberg correction (Figure OA7; smallest  $q = 0.20$ ). The one exception, plotted separately in the main text, is the rare-topic margin (Figure 10), the only topic outcome that is BH-significant ( $q < 0.001$ ). Distribution-level tests agree with the per-topic view: the Jensen–Shannon divergence between *JAR*’s and the controls’ topic mixes moves by +0.007 bits at the mandate (*JAR* does not rank first when each control journal is relabeled as treated; within-*JAR* permutation  $p = 0.67$ ; exact  $\chi^2$   $p = 0.69$ ), and *JAR*’s granular-topic entropy rises by about half a nat relative to controls (from 11 to 20 distinct topics on a nearly constant paper count), which is the concentration-side expression of the same rare-topic entry, not an independent fact.


**Table OA15:** Novelty: robustness across measures and convergent validity

| Measure ( $JAR \times Post$ )                                          | SD units | $t$         | Placebo rank |
|------------------------------------------------------------------------|----------|-------------|--------------|
| <i>Panel A: Embedding-based measures</i>                               |          |             |              |
| Intro-document kNN-20 (headline)                                       | +0.36    | 2.52        | 1/6          |
| Abstract kNN-20                                                        | +0.43    | 2.83        | 1/6          |
| kNN-5 (intro)                                                          | +0.39    | 3.69        | 1/6          |
| kNN-50 (intro)                                                         | +0.41    | 2.75        | 1/6          |
| Originality ( $1 - \max$ similarity)                                   | +0.33    | 3.27        | 1/6          |
| kNN-20, five-year reference window                                     | +0.44    | 3.22        | 1/6          |
| Distance to field centroid                                             | +0.17    | 1.41        | 3/6          |
| <i>Panel B: Reference-network measures (no text input)</i>             |          |             |              |
| Non-accounting reference share                                         | +0.48    | 3.73        | 1/6          |
| Reference field entropy                                                | +0.40    | 3.12        | 1/6          |
| Distinct cited journals                                                | +0.40    | 2.49        | 1/6          |
| Mean reference age                                                     | +0.18    | 1.69        | 2/6          |
| Uzzi journal-pair atypicality                                          | −0.40    | −3.44       | 6/6          |
| <i>Panel C: Convergent validity (correlation with headline kNN-20)</i> |          |             |              |
| Abstract kNN-20                                                        |          | $r = 0.70$  |              |
| Non-accounting reference share                                         |          | $r = 0.35$  |              |
| Distinct cited journals                                                |          | $r = 0.29$  |              |
| Reference field entropy                                                |          | $r = 0.27$  |              |
| Uzzi journal-pair atypicality                                          |          | $r = -0.33$ |              |
| Rhetorical novelty claim (indicator)                                   |          | $r = 0.01$  |              |


*Notes.* Panels A and B report the  $JAR \times Post$  coefficient from (1) on each novelty measure, expressed in standard-deviation units of the measure, for the All-Accounting comparison set and the headline 2012–2017 window; the Top-3 set gives the same picture. “Placebo rank” is  $JAR$ ’s position when each control journal is relabeled treated in turn, ordered one-sided in the predicted direction (novelty *rises*); with six journals and no random assignment it is a placebo ranking, not a significance level.  $JAR$  ranks first in every embedding and reference-diversity measure; the two discriminant checks (distance to the field centroid and mean reference age, where no directional prediction is made) are ranked two-sided and, as expected, do not place  $JAR$  first. The embedding effect is stable across every  $k$ , reference window, and variant (SD magnitudes 0.33–0.44). It is *local*: distance to nearest predecessors moves but distance to the field centroid does not. Most tellingly, the reference-network measures use no text yet move the same way and also place  $JAR$  first, and they correlate positively with the text measure (Panel C)—an independent corroboration that does not depend on the embedding pipeline. Uzzi journal-pair atypicality is computed against a 200-shuffle degree-preserving null (seed 42). The rhetorical novelty-claim indicator is nearly orthogonal to every distance measure ( $r \approx 0.01$ ), so we treat it as a distinct self-report construct.




**Figure OA4: Distribution of introduction novelty, before versus after the mandate.** Kernel density of introduction novelty (kNN-20 embedding distance) in *JAR* papers before and after the mandate. Mass moves past the pre-mandate ninetieth percentile (dotted line), the distributional counterpart of the quantile estimates in Figure 9.



**Figure OA5: Non-accounting share of references, by submission year.** Mean non-accounting share of a paper's references by submission year, for *JAR* and the pooled accounting controls; the dashed line marks January 2015. The share jumps for *JAR* after the policy while the controls hold flat, a visual analogue of the reference-based difference-in-differences in Table OA15.



**Figure OA6: *JAR*'s core topic mix is stable at the mandate.** Coarse topic shares of *JAR* papers (navy) and the accounting controls (grey), pre (2012–14) versus post (2015–17) submission years. No share movement survives multiple-testing correction.




**Figure OA7: Per-topic difference-in-differences.**  $JAR \times Post$  estimates for each topic-share outcome with 95% confidence intervals (analytic clustering by journal×year), sorted by coefficient; filled markers survive Benjamini–Hochberg correction, hollow markers do not. Only the rare-topic margin survives.



# I Citations and selection: supporting figures


This appendix collects the supporting evidence for the market-margin results in Sections 5.3.2 and 5.3.3. Figure OA8 reports the  $JAR \times Post$  effect on every citing-journal prestige and tier measure in standard-deviation units, including the versions that drop the top-six accounting journals as citers; the effect is positive and significant across the board, so the rise in citing-journal status is not an artifact of any single measure or of the field’s own leaders. Figure OA9 shows the accompanying broadening in the fields that cite *JAR*: the business-and-accounting share of citations falls after the mandate while economics, finance, and the social and decision sciences pick up the difference.



**Figure OA8: Citing-journal status rises across the main prestige and tier measures.**  $JAR \times Post$  difference-in-differences on each citation-prestige and citation-tier measure, in standard-deviation units, with 95% confidence intervals (All-Accounting sample; analytic clustering by journal $\times$ year). The “excl. top-6 accounting” rows drop the six leading accounting journals from the citing pool; those estimates remain positive but are not statistically distinguishable against the accounting controls, so the prestige gain is concentrated in the field’s leading outlets.

# J Readability: robustness across indices

The readability result in Section 5.1.1 does not depend on the choice of the Gunning–Fog index. Figure OA10 re-estimates the  $JAR \times Post$  effect for the abstract and the introduction using both Fog and the SMOG grade, a separate readability index built entirely from the




**Figure OA9: After the mandate, *JAR* is cited from a wider set of fields.** Change in the share of *JAR*’s citations originating in each field, pre- versus post-mandate. Business and accounting fall; economics, finance, and the social and decision sciences rise.

count of polysyllabic words. The two indices give the same answer—each significant, each between roughly one-third and one-half of a standard deviation—and they track each other at the paper level ( $r = 0.98$  for the abstract,  $0.88$  for the introduction), so the estimated gain is a property of the prose, not an artifact of one formula. The two remaining components of Fog behave as Section 5.1.1 describes: the complex-word share moves, while average sentence length does not. Character-based indices (Flesch, Coleman–Liau) require the raw text, which our aggregated corpus does not retain, so we do not compute them. A quantile version of the readability effect shows broad-based simplification with a modest gradient—the hardest-to-read abstracts simplify somewhat more than the median—rather than a change confined to the tail.

## K Differential measurement error

Because our outcomes are produced by a text pipeline and our treatment is a calendar cutoff, a natural concern is that the measurement instrument itself behaves differently across treated and control journals or across the pre and post periods, which would load onto the  $JAR \times Post$  coefficient and masquerade as a treatment effect. We check this by regressing each pipeline primitive—text-extraction yield, the classifier’s predicted probabilities, the mass of near-threshold classifications, and the share of statistics whose  $p$ -value is imputed from significance



**Figure OA10: The readability gain is robust across indices.**  $JAR \times Post$  effect on abstract and introduction readability, in standard-deviation units, measured by Gunning–Fog (crimson) and the SMOG grade (blue), with 95% confidence intervals (All-Accounting sample; clustering by journal $\times$ year). SMOG, an index based only on polysyllabic-word counts, confirms the Fog result.

stars—on  $JAR \times Post$  with journal and year fixed effects (Table OA16). The reassuring news is on the extensive margin: the share of papers with a usable abstract, introduction, and at least one extracted result is near one in every cell and shows no differential change (all  $p > 0.15$ ), so the estimates are not driven by differential sample censoring, and the per-paper readability measures, which are deterministic transforms of the text rather than classifier outputs, are not exposed to classifier drift. The two rows that count LLM-extracted results discussed in each paper fall for  $JAR$  after the mandate because post-mandate papers make fewer headline claims (Section 5.1.3); they are outcomes rather than instrument checks, and the comprehensive coefficient census behind the total-coefficient count in that section is a separate extraction. The cautions are on the intensive margin and are concentrated in two measure families. First, the introduction yields about three more extracted sentences per  $JAR$  paper post-mandate and the hedging classifier labels those introductions slightly less decisively, so the introduction-hedging rate is the one writing outcome whose differential measurement we cannot fully separate from a true effect—which is one more reason we do not rest any behavioral claim on the *average* hedging decline. The weak-paper gradient of Section 5.1.4 is a within- $JAR$  contrast between papers scored by the same pipeline in the same period, so a level drift in the classifier differences out of it. Second,  $JAR$ ’s post-mandate share of star-imputed  $p$ -values is differentially higher, which compounds the caution on the reported-statistics distribution and is why we treat the test-statistic channel as an unresolved null (Section 5.1.2) and report it on the stars-excluded aggregates. The readability and novelty headlines are not implicated by either caution. We note one limitation that the authors will

close in revision: the frozen analysis files do not contain the held-out human-labeled validation set, so classifier  $F_1$ , precision, recall, and inter-rater agreement—ideally reported separately for the treated-post cell—belong in this appendix and are reported in the supplementary validation materials.

## **L Culture versus policy: supporting estimates**

This appendix supports the coincident-shift test in Section 5.4.4. Table OA17 reports the horse race between the mandate and the econ-shift proxies (Panel A) and the weak-paper concentration test conditioned on the shift (Panel B). The smoothness event study that justifies the conditioning is Figure 16 in the main text. The exercise has two limits. First, the smoothness result is an equivalence claim and is bounded by its power: the design excludes 2015 steps larger than roughly 0.13 standard deviations for the quasi-experimental share, which rules out breaks of the size the writing outcomes display but not arbitrarily small ones. Second, the proxies are contemporaneous, so the horse race is a stability argument rather than a causal decomposition; it is informative here only because the shift shows no 2015 break. A backward-citation version of the econ-shift measure, built from each paper’s reference list, is a refinement we have not run.

**Table OA16:** Differential measurement-error and coverage balance: *JAR* vs. controls, pre vs. post (empirical sample, January 2012–December 2017)

| Measure                                               | JAR    |        | Controls |        | DiD (JAR×Post) |          |
|-------------------------------------------------------|--------|--------|----------|--------|----------------|----------|
|                                                       | Pre    | Post   | Pre      | Post   | Estimate       | <i>p</i> |
| <i>Panel A: Text-extraction yield</i>                 |        |        |          |        |                |          |
| Abstract sentences extracted                          | 5.940  | 5.859  | 7.743    | 7.547  | -0.4955*       | 0.055    |
| Introduction paragraphs extracted                     | 13.917 | 15.487 | 15.245   | 16.767 | +0.1201        | 0.464    |
| Introduction sentences extracted                      | 60.286 | 68.846 | 65.488   | 72.036 | +2.1360***     | 0.000    |
| LLM-extracted results (all)                           | 11.619 | 11.026 | 11.482   | 11.954 | -1.0920***     | 0.000    |
| LLM-extracted results (primary)                       | 2.583  | 1.962  | 2.096    | 1.893  | -0.4452***     | 0.000    |
| Share w/ usable abstract ( $\geq 3$ sent)             | 1.000  | 1.000  | 0.996    | 1.000  | -0.0039        | 0.156    |
| Share w/ usable intro ( $\geq 5$ sent)                | 1.000  | 1.000  | 0.999    | 1.000  | -0.0016        | 0.340    |
| Share w/ $\geq 1$ extracted result                    | 1.000  | 1.000  | 0.983    | 0.985  | +0.0008        | 0.931    |
| <i>Panel B: Classifier confidence (decisiveness)</i>  |        |        |          |        |                |          |
| Abstract mean hedge prob.                             | 0.639  | 0.600  | 0.575    | 0.538  | -0.0049        | 0.607    |
| Intro mean hedge prob.                                | 0.517  | 0.500  | 0.494    | 0.495  | -0.0158***     | 0.000    |
| Abstract mean claim prob.                             | 0.361  | 0.400  | 0.425    | 0.462  | +0.0049        | 0.606    |
| Intro mean claim prob.                                | 0.483  | 0.500  | 0.506    | 0.505  | +0.0158***     | 0.000    |
| <i>Panel C: Near-threshold ambiguity (0.5 cutoff)</i> |        |        |          |        |                |          |
| Abstract near-0.5 ambiguity index                     | 0.591  | 0.680  | 0.639    | 0.653  | +0.0725***     | 0.000    |
| Intro near-0.5 ambiguity index                        | 0.867  | 0.863  | 0.859    | 0.858  | -0.0051        | 0.231    |
| Abstract hedged-sentence share                        | 0.664  | 0.610  | 0.586    | 0.545  | -0.0178        | 0.100    |
| Intro hedged-sentence share                           | 0.523  | 0.503  | 0.496    | 0.499  | -0.0204***     | 0.000    |
| <i>Panel D: Star-imputed statistics</i>               |        |        |          |        |                |          |
| Share primary <i>p</i> -values star-imputed           | 0.000  | 0.038  | 0.039    | 0.037  | +0.0346***     | 0.008    |
| Share all results lost if stars dropped               | 0.023  | 0.035  | 0.029    | 0.020  | +0.0175***     | 0.008    |

*Notes.* Empirical (archival/experimental) sample, JAR plus accounting and finance controls, treatment-certain papers, assignment years 2012–2017. Cells are group×period means. DiD is the JAR×Post coefficient from OLS with journal and assignment-year fixed effects, standard errors clustered by journal. Panel A measures whether extractable text/statistics differ across the cutoff; Panel B whether the classifier labels sentences more/less decisively (mean predicted probabilities); Panel C the mass of near-0.5 (ambiguous) classifications and the realized hedged-sentence share; Panel D the share of extracted coefficients whose *p*-value is imputed from significance stars. A significant DiD indicates the *instrument* (extraction or classifier) behaves differently for JAR post-2015, so the corresponding language outcome is not cleanly comparable across the cutoff. \* $p < 0.10$ , \*\* $p < 0.05$ , \*\*\* $p < 0.01$ . Held-out classifier F1/precision/recall and inter-rater agreement are author inputs and are not contained in the frozen data package.

**Table OA17:** Culture versus policy: the writing results conditioned on the shift toward economics-style work

| <i>Panel A: Horse race — JAR <math>\times</math> Post before and after conditioning on the econ-shift proxies</i> |                   |                     |          |                |
|-------------------------------------------------------------------------------------------------------------------|-------------------|---------------------|----------|----------------|
| Outcome                                                                                                           | Baseline $\beta$  | Conditioned $\beta$ | Retained | Oster $\delta$ |
| Abstract readability (Fog)                                                                                        | −1.250            | −1.086              | 87%      | 4.30           |
| Introduction hedging                                                                                              | −0.028            | −0.031              | 113%     | $\delta < 0$   |
| Novelty (intro-doc kNN-20)                                                                                        | +0.013            | +0.014              | 108%     | —              |
| Log total citations                                                                                               | −0.277            | −0.291              | 105%     | $\delta < 0$   |
| Citing share, top-6 accounting                                                                                    | +0.028            | +0.028              | 98%      | 10.15          |
| <i>Panel B: Concentration — net weak-paper hedging effect, conditioned on the econ shift</i>                      |                   |                     |          |                |
| Hedging outcome                                                                                                   | Baseline net-weak | Conditioned         | Retained | Placebo rank   |
| Contribution hedging (CARS)                                                                                       | −0.185            | −0.185              | 100%     | 1/6            |
| Introduction hedging                                                                                              | −0.053            | −0.056              | 105%     | 1/6            |
| Gap hedging (CARS)                                                                                                | −0.065            | −0.074              | 114%     | 1/6            |

*Notes.* Accounting-only sample, submission years 2012–2017, submission-year and topic fixed effects, paper-level controls, standard errors clustered by journal  $\times$  year, as in Table 2. The two econ-shift proxies are the paper-level quasi-experimental-design indicator (an extracted identification strategy of difference-in-differences, instrumental variables, regression discontinuity, event study, or matching) and the share of classified introduction paragraphs making an identification claim; both enter in levels and interacted with *Post*. Panel A reports the JAR  $\times$  Post coefficient from the baseline specification and from the specification that adds both proxies; “Retained” is the ratio of the two. Oster  $\delta$  is the degree of selection on unobservables, relative to the included controls, required to drive the conditioned coefficient to zero ( $\delta < 0$ : the coefficient strengthens under conditioning, so no positive  $\delta$  overturns it). Panel B reports the net weak-paper hedging effect of Table OA8 (headline  $|z|$  below the sample median) before and after adding the econ-shift proxies and their weak-paper interactions; “Placebo rank” is the one-sided position of the conditioned effect when each of the six accounting journals is relabeled as treated in turn (1/6 means *JAR* ranks first); it is a placebo ranking, not a significance level. The conditioning is valid only because the proxies show no 2015 break (Figure 16); see Section 5.4.4.